

Delegating SQL Parsing to PostgreSQL

Greg Potter

FOSDEM 2026

How do you build a tool that
understands a Postgres schema?

Option A: Write a SQL parser 😭

Option B: Ask Postgres 🤔

```
CREATE FUNCTION calculate_weighted_average(
    measurements jsonb,
    weights numeric[] DEFAULT ARRAY[1.0],
    ignore_nulls boolean DEFAULT true
) RETURNS TABLE (
    category text,
    weighted_avg numeric(10,4),
    sample_count bigint,
    quality text
) LANGUAGE sql STABLE PARALLEL SAFE
BEGIN ATOMIC
    SELECT
        m->>'category',
        sum((m->>'value')::numeric * w) / nullif(sum(w), 0),
        count(*),
        CASE
            WHEN count(*) >= 100 THEN 'high'
            WHEN count(*) >= 10 THEN 'medium'
            ELSE 'low'
        END
    FROM jsonb_array_elements(measurements) WITH ORDINALITY AS t(m, i)
    LEFT JOIN unnest(weights) WITH ORDINALITY AS w(w, wi) ON t.i = w.wi
    WHERE NOT ignore_nulls OR m->>'value' IS NOT NULL
    GROUP BY m->>'category';
END;
```

What Postgres Knows Now

```
postgres=# SELECT proname, proargtypes, proallargtypes, proargnames, prorettype, provolatile,  
proparallel, prosqlbody FROM pg_proc WHERE proname ='calculate_weighted_average';
```

```
-[ RECORD 1 ]-----  
proname      | calculate_weighted_average  
proargtypes   | 3802 1231 16          -- jsonb, numeric[], boolean  
proallargtypes| {3802,1231,16,25,1700,20,25} -- jsonb, numeric[], boolean, text, numeric, ...  
proargnames   | {measurements,weights,ignore_nulls,category,weighted_avg, sample_count,quality}  
prorettype    | 2249  
provolatile   | s                      -- STABLE  
proparallel   | s                      -- PARALLEL SAFE  
prosqlbody    | (parsed query tree)
```

Shadow Database



System Catalogs

| | |
|---------------|---|
| pg_class | tables, views, indexes, sequences, materialized views |
| pg_attribute | columns |
| pg_proc | functions, procedures, aggregates |
| pg_type | types (built-in, composite, enum, domain, range) |
| pg_constraint | primary keys, foreign keys, check, unique, exclusion |
| pg_index | index details (columns, expressions, predicates) |

Querying the catalogs

```
SELECT attname,
       format_type(atttypid, atttypmod) as type,
       NOT attnotnull as nullable,
       pg_get_expr(d.adbin, d.adrelid) as default
  FROM pg_attribute a
 LEFT JOIN pg_attrdef d ON a.attrelid = d.adrelid AND a.attnum = d.adnum
 WHERE a.attrelid = 'orders'::regclass
   AND a.attnum > 0;
```

| attname | type | nullable | default |
|----------|--------------|----------|-------------------------|
| id | integer | f | |
| status | order_status | f | 'pending'::order_status |
| customer | text | f | |

Order Matters

```
postgres=# DROP TABLE orders;
ERROR:  cannot drop table orders because other objects depend on it
DETAIL:  function pending_orders() depends on table orders
view pending_order_count depends on function pending_orders()
HINT:  Use DROP ... CASCADE to drop the dependent objects too.
```

```
postgres=# DROP TYPE order_status;
ERROR:  cannot drop type order_status because other objects depend on it
DETAIL:  column status of table orders depends on type order_status
function pending_orders() depends on type order_status
view pending_order_count depends on function pending_orders()
HINT:  Use DROP ... CASCADE to drop the dependent objects too.
```

```
pending_order_count (view)
    ↓ uses
pending_orders() (function)
    ↓ queries
orders (table)
    ↓ uses
order_status (type)
```

pg_depend stores the graph

| classid | objid | refobjid | deptype |
|------------|---------------------|----------------|---------|
| pg_rewrite | pending_order_count | pending_orders | n |
| pg_proc | pending_orders | orders | n |
| pg_class | orders | order_status | n |

From graph to order

pg_depend gives you edges:

pending_order_count → pending_orders

pending_orders → orders

orders → order_status

Topological sort gives you order:

CREATE: order_status → orders → pending_orders → pending_order_count

DROP: pending_order_count → pending_orders → orders → order_status

deptype

deptype = 'n' (normal)

view → function → table → type

Real dependencies. These are yours.

deptype = 'a' (automatic)

sequence → SERIAL column

index → table

Created together. Linked.

deptype = 'i' (internal)

TOAST table → parent

Postgres internals.

deptype = 'e' (extension)

st_distance → postgis

Belongs to an extension. Not yours.

What pg_depend can't see

```
-- pg_depend CAN'T see inside this (string-literal bodies)
CREATE FUNCTION pending_orders() RETURNS SETOF orders
LANGUAGE sql STABLE
AS $$

    SELECT * FROM orders WHERE status = 'pending';

$$;
```

pg_depend tracks:

function → return type (SETOF orders) ✓

pg_depend doesn't track:

function body → orders table ✗

function body → status column ✗

BEGIN ATOMIC fixes this

```
-- pg_depend CAN see inside this (BEGIN ATOMIC)
CREATE FUNCTION pending_orders() RETURNS SETOF orders
LANGUAGE sql STABLE
BEGIN ATOMIC
    SELECT * FROM orders WHERE status = 'pending';
END;
```

pg_depend tracks:

- function → return type (SETOF orders) ✓
- function body → orders table ✓
- function body → status column ✓

Implicit objects

```
CREATE TABLE my_table (
    id SERIAL PRIMARY KEY
);
          objid | refobjid | deptype
-----+-----+-----
my_table | my_table_id_seq | a
my_table | my_table_pkey   | a
```

Creates:

```
table: orders
sequence: orders_id_seq    (deptype 'a')
index: orders_pkey         (deptype 'a')
```

Array types

```
CREATE TYPE order_status AS ENUM (...);
```

Creates:

```
type: order_status
```

```
type: _order_status      (array type, depends on order_status, deptype 'i');
```

```
CREATE TABLE array_table (statuses order_status[]);
```

| objid | | refobjid | | deptype |
|-------------------|--|---------------|--|---------|
| -----+-----+----- | | | | |
| array_table | | _order_status | | n |
| _order_status | | order_status | | i |

Trade Offs

You need a running Postgres
Shadow database = real database

You have to build support explicitly
There are a lot of catalog tables

Postgres versions matter
Catalogs evolve
prosqlbody requires Postgres 14+

Review your output
Only as complete as your queries

The toolkit

Shadow database

→ Postgres parses, you query

System catalogs

→ Structure, types, columns

pg_depend

- Dependency graph
- Filter by deptype
- Know the blind spots
- Watch for implicit objects

Topological sort

→ Edges → correct order

pgmt

I built pgmt using all of this.

Schema-as-code for Postgres.

Edit .sql files, pgmt generates migrations.

pgmt.dev

github.com/gdpotter/pgmt

Not just for migrations

Drift detection

- Compare expected catalogs to live state

Schema linters

- Query catalogs for problems
(unused indexes, missing FKs, naming violations)

CI validation

- Apply to shadow database
- If it succeeds, the schema parses

Visualization

- `pg_depend` is the dependency graph
- Draw it

Postgres already knows your schema.

Just ask.