


Designing Fibers for systemd

Structured POSIX avoidance in PID 1


Background

- systemd has a lot of daemons
- These daemons have to handle requests concurrently




Threads + fork() == pain

- Signal handlers face similar problems



Event Loop + Callbacks == pain

Stack




Callback 1

```
char *s;  
  
s = strjoin("foo", "bar");  
...  
sd_event_add_io(..., callback2, s);  
...  
<stack unwinding starts>  
...
```

Callback 2

```
_cleanup_free_ char *q;  
  
q = strjoin(s, "quux");  
...  
<MEMORY LEAK>
```



Callback 1

```
char *s;  
  
s = strjoin("foo", "bar");  
...  
sd_event_add_io(e, callback2, s);  
...  
sd_event_source_set_destroy_callback(e, free);  
...  
<stack unwinding starts>  
...
```

Callback 2


```
_cleanup_free_ char *q;  
  
q = strjoin(s, "quux");  
...
```


Threads vs Callbacks

=> Stack Unwinding

Can we have the best of both worlds?

- Can we have separate stacks without separate threads?
- Can we pass around an entire stack between individual callbacks?






How to implement stack suspend/resume?

- ucontext.h
- getcontext()
- makecontext()
- swapcontext()

fiber_read()



Cancellation

- To ensure proper cleanup for fibers, stack unwinding is essential
- When shutting down a daemon, running fibers have to be interrupted
- Implemented by injecting the ECANCELED error into running fibers
 - And then waiting for the fibers to finish running

Considerations when spawning background fibers

- An existing fiber can spawn more fibers to run tasks in the background
- Thinking about ownership of data passed into background fibers is essential
 - Shared state!
- When in doubt, copy input arguments to each background fiber

Disadvantages?

- No threads, but still need locking, concurrency primitives, ...
- Less efficient than callbacks
- Not POSIX! (deprecated)
- swapcontext() not available on musl (need libucontext)
- Worse integration than threads (gdb, core dumps, ...)
- Performance?
 - POSIX mandates saving/restoring per thread signal mask for every swapcontext()
- Still no non-blocking disk I/O
 - But, Jens Axboe is working behind the scenes to make io-uring more widely available!

Current status

- PR open: github.com/systemd/systemd/pull/39771
- Introduces new libsystemd interface sd-future.h
 - Internal for now, but might become public in the future
 - Adds support for using fibers with sd-event.h
- Basic operations implemented
 - Socket IO
 - Child processes
 - Sleeping
 - Waiting for other fibers
 - ...
- Concurrency primitives still missing

Questions?