



eBPF Hookpoint Gotchas: Why Your Program Fires (or Fails) in Unexpected Ways

Chris Tarazi, Sr. Staff Software Engineer
Donia Chaiehoudj, Software Engineer and Community



Overview

- Tracing hookpoint overview
- Generic gotchas:
 - Kernel versions
 - Different architectures
 - Performance
- uprobe dynamic links gotchas
- kprobe/uprobe inlining gotchas
- Nested executions gotchas



Tracing hookpoint overview

Hookpoint	Scope	Stability	Kernel version	Mechanism	Use case
kprobe	*Any kernel function entry / exit, offsets	No guarantees	2.6.9	Various (interrupt-based, optimization exists)	Dynamic debugging, kernel internals
uprobe	Userspace function entry / exit, offsets	No guarantees	3.5	Same as kprobe (userspace traps into kernel)	Dynamic debugging, application internals
fprobe	Function entry / exit	Same as kprobe/uprobe	5.5	ftrace (NOP patching + trampoline)	More efficient kernel function tracing
tracepoint	Kernel static points	Designed to not break (exceptions possible)	2.6.28	NOP patching	Production tracing and performance accounting

*Inlined functions, blacklisted functions, or functions marked as notrace are excluded from probes.

Hookpoints performance overview

Empty probe attached:

Case	ns/op	Overhead ns/op	Overhead percent
no probe attached	117	0	0%
tracepoint empty	132	15	13%
fentry empty	141	24	21%
kprobe empty	254	137	117%

Probe with a simple map increment attached:

Case	ns/op	Overhead ns/op	Overhead percent
no probe attached	117	0	0%
tracepoint simple	152	35	30%
fentry simple	159	42	36%
kprobe simple	277	160	136%

Probe with a complex map increment attached:

Case	ns/op	Overhead ns/op	Overhead percent
no probe attached	117	0	0%
tracepoint complex	213	96	82%
fentry complex	220	103	88%
kprobe complex	346	229	196%

Hookpoint ranking

-  tracepoints
-  fprobes
-  kprobes (慢 slow due to interrupt-based implementation)


- note:
 - overhead depends on whether function is in the hot path or not
 - overhead may not be direct translation to system slowdown

source: <https://mastodon.ivan.computer/@mastodon/110737250286611183>



Gotcha #1 Kernel versions

“



Gotcha #1: Kernel versions



kprobes, fprobes have no stability guarantee:

Gotcha #1: Kernel versions



kprobes, fprobes have no stability guarantee:

- **kernel internal functions can change with any new release:**
 - function is renamed, removed, inlined: Failed to attach because the symbol not found
 - function arguments changed, struct field changed: Reads the wrong argument → garbage data

Gotcha #1: Kernel versions



kprobes, fprobes have no stability guarantee:

- **kernel internal functions can change with any new release:**
 - function is renamed, removed, inlined: Failed to attach because the symbol not found
 - function arguments changed, struct field changed: Reads the wrong argument → garbage data
- **Alternatives**
 - **Use tracepoints which are more reliable if possible**
 - Statically defined instrumentation points added by kernel developers
 - Guaranteed stable ABI: arguments are documented and maintained
 - `# cat /sys/kernel/debug/tracing/available_events`

Gotcha #1: Kernel versions



kprobes, fprobes have no stability guarantee:

- **kernel internal functions can change with any new release:**
 - function is renamed: "Failed to attach: symbol not found"
 - function is removed: "Attached: 0 events"
 - function inlined: "Failed to attach: symbol not found"
 - function arguments changed: "Reads wrong argument: garbage data"
- **Alternatives**
 - **Use tracepoints which are more reliable if possible**
 - Statically defined instrumentation points added by kernel developers
 - Guaranteed stable ABI: arguments are documented and maintained
 - `# cat /sys/kernel/debug/tracing/available_events`

Q: What if there is no tracepoint for the function you want to cover, what is your option?

- A. Hardcode the function name and hope it doesn't change
- B. Write separate BPF programs for each kernel version
- C. Use CO-RE with BTF for automatic adaptation
- D. Abandon eBPF and use kernel modules instead

Gotcha #1: Kernel versions



kprobes, fprobes have no stability guarantee:

- **kernel internal functions can change with any new release:**
 - function is renamed: "Failed to attach: symbol not found"
 - function is removed: "Attached: 0 events"
 - function inlined: "Failed to attach: symbol not found"
 - function arguments changed: "Reads wrong argument: garbage data"
- **Alternatives**
 - **Use tracepoints which are more reliable if possible**
 - Statically defined instrumentation points added by kernel developers
 - Guaranteed stable ABI: arguments are documented and maintained
 - `# cat /sys/kernel/debug/tracing/available_events`

Q: What if there is no tracepoint for the function you want to cover, what is your option?

- A. Hardcode the function name and hope it doesn't change
- B. Write separate BPF programs for each kernel version
- C. Use CO-RE with BTF for automatic adaptation
- D. Abandon eBPF and use kernel modules instead

Gotcha #1: Kernel versions




kprobes, fprobes have no stability guarantee:

- **kernel internal functions can change with any new release:**
 - function is renamed: "Failed to attach: symbol not found"
 - function is removed: "Attached: 0 events"
 - function inlined: "Failed to attach: symbol not found"
 - function arguments changed: "Reads wrong argument: garbage data"
- **Alternatives**
 - **Use tracepoints which are more reliable if possible**
 - Statically defined instrumentation points added by kernel developers
 - Guaranteed stable ABI: arguments are documented and maintained
 - `# cat /sys/kernel/debug/tracing/available_events`
 - **CO-RE (Compile Once Run Everywhere) with BTF**
 - CO-RE = BPF programs adapt to kernel at load time
 - BTF (BPF Type Format) = Type information embedded in kernel
 - e.g: `bpftool btf dump file /sys/kernel/btf/vmlinux > vmlinux.h`
 - Relocations = libbpf adjusts offsets automatically

Gotcha #2 Architectures

“



Gotcha #2: Different architectures



kprobes have an unfriendly interface: the context passed to a kprobe program is struct pt_regs which is a struct representing the registers. It stores the function parameters.

- **Same BPF code behaves differently on x86_64 vs arm64:**
 - kprobe
 - different registers (RDI vs X0)

Gotcha #2: Different architectures




kprobes have an unfriendly interface: the context passed to a kprobe program is struct pt_regs which is a struct representing the registers. It stores the function parameters.

- **Same BPF code behaves differently on x86_64 vs arm64:**
 - kprobe
 - different registers (RDI vs X0)
- **Alternatives**
 - Use other hook points: fprobes, tracepoints (no raw_tracepoints)
 - libbpf provides helpers for access ([tools/lib/bpf/bpf_tracing.h](#))

Gotcha #3 Dynamic Links

“



Gotcha #3: uprobe dynamic links



```
// libmath.c
int add(int a, int b) {
    return a + b;
}

// main.c: loads, attaches and runs add():
int result = add(2, 3);
printf("Result: %d\n", result);
```

Gotcha #3: uprobe dynamic links



```
// libmath.c
int add(int a, int b) {
    return a + b;
}

// main.c: loads, attaches and runs add():
int result = add(2, 3);
printf("Result: %d\n", result);

// Attach uprobe to add()
# sudo bpftrace -e 'uprobe:./libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
```

Symbol table:

add() → offset 0x1140	← Probe attaches HERE
-----------------------	-----------------------

Gotcha #3: uprobe dynamic links



```
// libmath.c
int add(int a, int b) {
    return a + b;
}

// main.c: loads, attaches and runs add():
int result = add(2, 3);
printf("Result: %d\n", result);

// Attach uprobe to add()
# sudo bpftrace -e 'uprobe:/libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
```

Symbol table:

add() → offset 0x1140	← Probe attaches HERE
-----------------------	-----------------------

Gotcha #3: uprobe dynamic links

```
// libmath.c
int add(int a, int b) {
    return a + b;
}

// main.c: loads, attaches and runs add():
int result = add(2, 3);
printf("Result: %d\n", result);

// Attach uprobe to add()
# sudo bpftrace -e 'uprobe:/libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
```

Symbol table:

add() → offset 0x1140

← Probe attaches HERE

```
==== Running main program ====
LD_LIBRARY_PATH=. ./main
add called: 2 + 3
Result: 5
```

```
// bpftrace output
sudo bpftrace -e 'uprobe:/libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
Attaching 1 probe...
uprobe: add(2, 3)
```

Gotcha #3: uprobe dynamic links



Update library

```
// libmath.c
// Insert new functions
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
// add function is now at the bottom
int add(int a, int b) { return a + b; }
```

Gotcha #3: uprobe dynamic links

Update library

```
// libmath.c
// Insert new functions
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
// add function is now at the bottom
int add(int a, int b) { return a + b; }
```

==== Running main program ===

LD_LIBRARY_PATH=. ./main

add called: 2 + 3

Result: 5

// bpftrace output

```
sudo bpftrace -e 'uprobe:/libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
```

Attaching 1 probe...

uprobe: add(2, 3)

Gotcha #3: uprobe dynamic links

Update library

```
// libmath.c
// Insert new functions
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
// add function is now at the bottom
int add(int a, int b) { return a + b; }
```

```
==== Running main program ====
LD_LIBRARY_PATH=. ./main
add called: 2 + 3
Result: 5

// bpftace output
sudo bpftace -e 'uprobe:/libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
Attaching 1 probe...
uprobe: add(2, 3)
```

Q: Your uprobe is tracing add(). A colleague updates libmath.c and adds new functions. You run your program again. What do you see?

- A. add(2, 3) - everything works fine
- B. subtract(2, 3) - probe fires on wrong function
- C. Nothing - probe silently stops working
- D. None of the above

Gotcha #3: uprobe broken dynamic links



Update library

```
// libmath.c  
  
// Insert new functions  
  
int subtract(int a, int b) { return a - b; }  
int multiply(int a, int b) { return a * b; }  
  
// add function is now at the bottom  
  
int add(int a, int b) { return a + b; }
```

```
==== Running main program ===
```

```
LD_LIBRARY_PATH=. ./main
```

```
add called: 2 + 3
```

```
Result: 5
```

```
// bpftrace output
```

```
sudo bpftrace -e 'uprobe:/libmath.so:add {  
    printf("uprobe: add(%d, %d)\n", arg0, arg1); }'
```

```
Attaching 1 probe...
```

```
uprobe: add(2, 3)
```

Q: Your uprobe is tracing add(). A colleague updates libmath.c and adds new functions. You run your program again. What do you see?

- A. add(2, 3) - everything works fine
- B. subtract(2, 3) - probe fires on wrong function
- C. Nothing - probe silently stops working**
- D. None of the above

Gotcha #3: uprobe broken dynamic links



Update library

```
// libmath.so
// Insert new functions
int subtract(int a, int b) { return a - b; }
int multiply(int a, int b) { return a * b; }
// add function is now at the bottom
int add(int a, int b) { return a + b; }
```

Symbol table AFTER update:

subtract	→ offset 0x1140
multiply	→ offset 0x1180
add	→ offset 0x11c0

← Probe still points here! ❌
← add() is now HERE

==== Running main program ===

LD_LIBRARY_PATH=. ./main

add called: 2 + 3

Result: 5

// bpftrace output

sudo bpftrace -e 'uprobe::./libmath.so:add {
printf("uprobe: add(%d, %d)\n", arg0, arg1); }'

Attaching 1 probe...

uprobe: add(2, 3)

Gotcha #3: uprobe broken dynamic links




Conclusion

- **Problem**
 - Library recompiled/updated, uprobes are attached to the old offset
- **In production:**
 - Library updates are rare... but they happen (security patches!)
 - your monitoring may silently break without warning
- **Solutions:**
 - Monitor for file changes (inotifywait)
 - Use a supervisor program that monitors updates to the relevant shared libraries to re-attach your uprobes after updates


Gotcha #4: uprobe/kprob e inlining

“



What is inlining?


“



What is inlining?

“

Inlining is an optimization where a function call is replaced with the actual code of the function itself at the point of call.



Gotcha #4: kprobe/uprobe inlining

```
// target.c

#include <stdio.h>
#include <stdlib.h>

int add(int a, int b) {
    return a + b;
}

int main(int argc, char *argv[]) {
    int x = atoi(argv[1]);

    if (x < 5) {
        printf("Usage: %d <number>\n", x);
        return 0;
    }

    int result = add(x, 10);
    printf("add(%d, 10) = %d\n", x, result);

    return 0;
}
```

Gotcha #4: kprobe/uprobe inlining

```
// target.c

#include <stdio.h>
#include <stdlib.h>

int add(int a, int b) {
    return a + b;
}

int main(int argc, char *argv[]) {
    int x = atoi(argv[1]);

    if (x < 5) {
        printf("Usage: %d <number>\n", x);
        return 0;
    }

    int result = add(x, 10);
    printf("add(%d, 10) = %d\n", x, result);

    return 0;
}
```

Compilation without any optimisation:

```
# gcc -g -O0 -o target_00 target.c
```

Disassembled program, symbol of the function is present:

```
# nm target_00 | grep add
00000000000007d8 T add
```

Gotcha #4: kprobe/uprobe inlining

```
// target.c

#include <stdio.h>
#include <stdlib.h>

int add(int a, int b) {
    return a + b;
}

int main(int argc, char *argv[]) {
    int x = atoi(argv[1]);

    if (x < 5) {
        printf("Usage: %d <number>\n", x);
        return 0;
    }

    int result = add(x, 10);
    printf("add(%d, 10) = %d\n", x, result);

    return 0;
}
```

Compilation without any optimisation:

```
# gcc -g -O0 -o target_00 target.c
```

Disassembled program, symbol of the function is present:

```
# nm target_00 | grep add
00000000000007d8 T add
```

Compilation with a level of optimisation O2:

```
# gcc -g -O2 -o target_02 target.c
```

Gotcha #4: kprobe/uprobe inlining

```
// target.c

#include <stdio.h>
#include <stdlib.h>

int add(int a, int b) {
    return a + b;
}

int main(int argc, char *argv[]) {
    int x = atoi(argv[1]);

    if (x < 5) {
        printf("Usage: %d <number>\n", x);
        return 0;
    }

    int result = add(x, 10);
    printf("add(%d, 10) = %d\n", x, result);

    return 0;
}
```

Compilation without any optimisation:

```
# gcc -g -O0 -o target_00 target.c
```

Disassembled program, symbol of the function is present:

```
# nm target_00 | grep add
00000000000007d8 T add
```

Compilation with a level of optimisation O2::

```
# gcc -g -O2 -o target_02 target.c
```


Disassembled program, symbol of the function is present:

```
# nm target_02 | grep add
0000000000000860 T add
```

Demo selective inlining

“

*I should be able to attach to
add() function in both cases:
without and with
optimisation?!*



Gotcha #4: kprobe/uprobe selective inlining

1. Build non-optimised and optimised versions

```
gcc -g -O0 -o target_00 target.c
gcc -g -O2 -o target_02 target.c
```

2. Check symbols - addf exists in BOTH binaries!

```
nm target_00 | grep addf
# 000000000000760 T addf

nm target_02 | grep addf
# 000000000000760 T addf      ← Symbol exists!
```

3. Trace 00 (no optimization).

```
# Terminal 1:
sudo bpftrace -e 'uprobe:./target_00:addf {
printf("addf called!\n"); }'

# Terminal 2:
./target_00 3      # x < 5: early return
./target_00 10     # x >= 5: calls addf

# Result: "addf called!" ✓
```

4. Trace 02 (optimized) - SAME probe, SAME program

```
# Terminal 1:
sudo bpftrace -e 'uprobe:./target_02:addf {
printf("addf called!\n"); }'

# Terminal 2:
./target_02 3      # nothing
./target_02 10     # nothing!
# Result: SILENCE ✗ (function inlined)
```

Gotcha #4: kprobe/uprobe selective inlining

5. Compare disassembly - WHY?

```
# 00: main CALLS addf
objdump -d target_00 | grep -A 15 "<main>:"
# ...
#   bl <addf>      ← BRANCH to addf function
```

02: main has addf INLINED

```
objdump -d target_02 | grep -A 15 "<main>:"
# ...
#   add w0, w0, #0xa   ← This IS addf(x,10), no
branch!
```

6. Solution: Attach at offset (uprobe/kprobe)

```
# Find where addf is inlined:
objdump -d target_02
# main at 0x6c0, "add w0,w0,#0xa" at 0x6e8
# offset = 0x6e8-0x6c0 = 0x28
```

```
sudo bpftrace -e 'uprobe:./target_02:main+0x28 {
printf("Hit inlined addf!\n"); }'
```

```
# Now ./target_02 10 triggers: "Hit inlined addf!" ✓
```

Gotcha #4: kprobe/uprobe selective inlining

Disassembled binary

No optimisation

```
# objdump -d -S --disassemble=main ./target_00
```

```
00000000000007f8 <main>:
// ...
int main(int argc, char *argv[]) {
 7f8: a9bd7bfd      stp    x29, x30, [sp, #-48]!
 7fc: 910003fd      mov    x29, sp
 800: b9001fe0      str    w0, [sp, #28]
 804: f9000be1      str    x1, [sp, #16]
  int x = atoi(argv[1]);
 808: f9400be0      ldr    x0, [sp, #16]
 80c: 91002000      add    x0, x0, #0x8
 810: f9400000      ldr    x0, [x0]
 814: 97ffff8f      bl    650 <atoi@plt>
 818: b9002be0      str    w0, [sp, #40]

  if (x < 5) {
 81c: b9402be0      ldr    w0, [sp, #40]
// ...
  }

  int result = addf(x, 10);
840: 52800141      mov    w1, #0xa
844: b9402be0      ldr    w0, [sp, #40]
848: 97ffffe4      bl    7d8 <add>
 84c: b9002fe0      str    w0, [sp, #44]
// ...
```

Gotcha #4: kprobe/uprobe selective inlining

Disassembled binary

No optimisation

```
# objdump -d -S --disassemble=main ./target_00
```

```
00000000000007f8 <main>:  
// ...  
int main(int argc, char *argv[]) {  
 7f8: a9bd7bfd    stp    x29, x30, [sp, #-48]!  
 7fc: 910003fd    mov    x29, sp  
 800: b9001fe0    str    w0, [sp, #28]  
 804: f9000be1    str    x1, [sp, #16]  
    int x = atoi(argv[1]);  
 808: f9400be0    ldr    x0, [sp, #16]  
 80c: 91002000    add    x0, x0, #0x8  
 810: f9400000    ldr    x0, [x0]  
 814: 97ffff8f    bl    650 <atoi@plt>  
 818: b9002be0    str    w0, [sp, #40]  
  
    if (x < 5) {  
 81c: b9402be0    ldr    w0, [sp, #40]  
// ...  
    }  
  
    int result = addf(x, 10);  
 840: 52800141    mov    w1, #0xa  
 844: b9402be0    ldr    w0, [sp, #40]  
 848: 97ffffe4    bl    7d8 <add>  
 84c: b9002fe0    str    w0, [sp, #44]  
// ...
```

Optimised

```
# objdump -d -S --disassemble=main ./target_02
```

```
00000000000006c0 <main>:  
// ...  
int main(int argc, char *argv[]) {  
 6c0: aa0103e0    mov    x0, x1  
 6c4: a9bf7bfd    stp    x29, x30, [sp, #-16]!  
// ...  
    int x = atoi(argv[1]);  
  
    if (x < 5) {  
 6e0: 7100101f    cmp    w0, #0x4  
 6e4: 5400012d    b.le   708 <main+0x48>  
}  
  
// ...  
  
 6e8: 11002803    add    w3, w0, #0xa  
 6ec: 90000001    adrp   x1, 0 <__abi_tag-0x278>  
 6f0: 52800040    mov    w0, #0x2  
 6f4: 91228021    add    x1, x1, #0x8a0  
 6f8: 97ffffd6    bl    650 <__printf_chk@plt>  
  
    int result = addf(x, 10);  
    printf("addf(%d, 10) = %d\n", x, result);  
// ...
```

Gotcha #4: kprobe/uprobe selective inlining

Conclusion

The Problem

- Symbol exists in binary (visible with nm)
- But compiler inlined the function into the caller
- uprobe is attached on symbol but it never fires,
→ executed instructions inlined in the caller

How to Detect Selective Inlining?

1. Show disassembly code

```
objdump -d s  
llvm-dwarfdump <your_binary>
```

2. Check if function is inlined

not inlined: bl <function> → real call

inlined: <instruction> → No bl, just the instructions directly

Gotcha #4: kprobe/uprobe selective inlining

Conclusion

The Problem

- Symbol exists in binary (visible with nm)
- But compiler inlined the function into the caller
- uprobe is attached on symbol but it never fires,
→ executed instructions inlined in the caller

How to Detect Selective Inlining?

1. Show disassembly code

```
objdump -d s  
llvm-dwarfdump <your_binary>
```

2. Check if function is inlined

not inlined: bl <function> → real call
inlined: <instruction> → No bl, just the instructions directly

Solution: Use uprobe or kprobe + offset


1. Look inside the caller function's disassembly
2. Find the instruction that does the work (e.g., add for addition)
3. Calculate offset: instr_address - fnc_start_address
4. Probe to the offset:
`uprobe:binary:<caller_function>+<offset>`
→ Fires when execution reaches that exact instruction

Caveats

- !! Multiple probes for all callsites of your function
- Offsets change with every recompilation
- Different compiler versions = different offsets
- Not practical for production, useful for debugging

Gotcha #5: kprobe/uprobe inlining

“



Gotcha #5: kprobe/uprobe inlining

```
// target.c
int allocate_resource(int size) {
    if (size <= 0 || size >= 1024) return -1;

    int resource_id = 0;
    for (int i = 0; i < size; i++) { ... } // Loop
    snprintf(log_buffer, ...);           // String formatting
    printf("%s\n", log_buffer);         // I/O
    return resource_id;
}
```

Gotcha #5: kprobe/uprobe inlining

```
// target.c
int allocate_resource(int size) {
    if (size <= 0 || size >= 1024) return -1;

    int resource_id = 0;
    for (int i = 0; i < size; i++) { ... }    // Loop
    snprintf(log_buffer, ...);                 // String formatting
    printf("%s\n", log_buffer);                // I/O
    return resource_id;
}
```

Q: How many symbols does this generate when compiling with the following command: `gcc -O2 -o target_gcc target.c` ?

- A. 1
- B. 2
- C. 10
- D. wth are symbols?

Gotcha #5: kprobe/uprobe inlining

A typical scenario...

```
// target.c
int allocate_resource(int size) {
    if (size <= 0 || size >= 1024) return -1;

    int resource_id = 0;
    for (int i = 0; i < size; i++) { ... } // Loop
    snprintf(log_buffer, ...);           // String formatting
    printf("%s\n", log_buffer);         // I/O
    return resource_id;
}
```

Q: How many symbols does this generate when compiling with the following command: `gcc -O2 -o target_gcc target.c` ?

- A. 1
- B. 2**
- C. 10
- D. wth are symbols?

```
$ nm target | grep -E "allocate_resource"
000000000000c40 T allocate_resource
000000000000ae0 t allocate_resource.part.0
```

Gotcha #5: kprobe/uprobe partial inlining

GCC Optimisation

gcc's -fpartial-inlining (enabled by -O2) splits functions based on execution patterns:

Fast path example:

- quick validation
- error path

Slow path example, extracted to .part:

- complex error handling
- complex processes

Before:

```
allocate_resource()  
  
    if (size <= 0 || size >= 1024)  
        return -1  
  
    for (i=0; i<size; i++)  
        resource_id += ...  
    sprintf(log_buffer, ...)  
    printf("%s", log_buffer)  
    return resource_id
```

After:

```
allocate_resource() 0xc20  
  
    if (size <= 0 || size >= 1024)  
        return -1  
    jump to .part.0  
  
allocate_resource.part.0 0xb90  
  
    for (i=0; i<size; i++)  
        resource_id += ...  
    sprintf(log_buffer, ...)  
    printf("%s", log_buffer)  
    return resource_id
```

Gotcha #5: kprobe/uprobe partial inlining



~~WTF is this optimization?~~

- Smaller hot/fast code → fits better in instruction cache
- Better branch prediction → CPU expects the fast path
- Reduced register pressure → cold/slow path variables don't pollute hot path

Gotcha #5: kprobe/uprobe partial inlining

1. Build with partial inlining
make target_partial

2. Check symbol table - see the split!
nm target_partial | grep allocate_resource
0xc20 T allocate_resource
0xb90 t allocate_resource.part.0

3. Compare sizes
readelf -s target_partial | grep allocate_resource

4. Probe the wrapper (catches valid calls)
sudo bpftrace -e 'uprobe:./target_partial:allocate_resource {
 printf("allocate_resource called\n");
'


5. Probe the error path
sudo bpftrace -e 'uprobe:./target_partial:allocate_resource.part.0 {
 printf("error path called!\n");
'

6. Run the target and observe
./target_partial

Demo partial inlining

“

**So what about
kprobe/uprobe
attachments and symbols?**



Gotcha #5: kprobe/uprobe partial inlining

Why probe misses events?

Program output:

```
./target

// Hot path
allocate_resource(-5) = -1

allocate_resource(9999) = -1
Allocated resource #2424 (size=100 bytes)
Total allocated so far: 100 bytes

// Cold path
allocate_resource(100) = 2424
Allocated resource #3024 (size=256 bytes)
Total allocated so far: 356 bytes

allocate_resource(256) = 3024
```

```
PROBE: allocate_resource
$ sudo bpftrace -e 'uprobe:...:allocate_resource {...}'
Attaching 1 probe...
(nothing)                                     ← SILENT! No events!
```

```
PROBE: allocate_resource.part.0
$ sudo bpftrace -e 'uprobe:...:allocate_resource.part.0 {...}'
Attaching 1 probe...
allocate_resource called with size=100
allocate_resource called with size=256
                                         ↑ ACTUAL ALLOCATIONS!
```

Gotcha #5: kprobe/uprobe partial inlining

Alternatives

- **generally:** look for alternative functions that has defined tracepoints (not 1:1 translation)
- **uprobe:** gcc -O2 -fnopartial-inlining -o <output> <your_program>
- **kprobe:**
 - kernel recompilation in production is not common → identify symbols using tools

1. Check for suffixed symbols

```
nm <your_binary> | grep -E "\.(part|cold|isra|constprop)"
```

2. Verify with DWARF


```
llvm-dwarfdump <your_binary> | grep -A5 "<your_function>"
```

3. See source mapping

```
objdump -S <your_binary> | grep -A30 "<your_function>:"
```

Gotcha #6: missed executions

“



Gotcha #6: missed executions



```
# bpftool p | grep misses
```

```
5845: tracing      name fentry_do_sys_openat2          tag xxx gpl recursion_misses 9482
5847: kprobe       name kprobe_do_sys_openat2         tag xxx gpl recursion_misses 245
11066: tracepoint  name tracepoint_lock_contention_begin_1 tag xxx gpl recursion_misses 25
```

Gotcha #6: missed executions



```
# bpftool p | grep misses

5845: tracing      name fentry_do_sys_openat2          tag xxx gpl recursion_misses 9482
5847: kprobe       name kprobe_do_sys_openat2         tag xxx gpl recursion_misses 245
11066: tracepoint  name tracepoint_lock_contention_begin_1 tag xxx gpl recursion_misses 25
```

Question: What does a "recursion_miss" counter represent?

- A) The number of times the BPF program failed to execute due to missing kernel functions
- B) How many times recursion protection prevented a program from running again while already executing
- C) A count of recursive function calls that the BPF verifier rejected
- D) The number of missed hardware events due to recursion in the perf subsystem

Gotcha #6: missed executions



```
# bpftool p | grep misses

5845: tracing      name fentry_do_sys_openat2          tag xxx gpl recursion_misses 9482
5847: kprobe       name kprobe_do_sys_openat2         tag xxx gpl recursion_misses 245
11066: tracepoint  name tracepoint_lock_contention_begin_1 tag xxx gpl recursion_misses 25
```

Question: What does a "recursion_miss" counter represent?

- A) The number of times the BPF program failed to execute due to missing kernel functions
- B) How many times recursion prevented the BPF program from running again while already executing**
- C) A count of recursive function calls that the BPF verifier rejected
- D) The number of missed hardware events due to recursion in the perf subsystem

Gotcha #6: missed executions



- be aware! your tracing programs may not always execute
- must assume that you will not always see all events in the kernel when using tracing programs (kprobes, fprobes, tracepoints)
- why you should care?
 - security monitoring
 - debugging
 - etc
- we will discuss
 - kprobes
 - fprobes
 - tracepoints

Gotcha #6: missed executions

- probes can miss executions in two different ways (**only if on the same CPU**)
 - recursion of the program
 - nested executions of any other eBPF programs
- how does a probe recurse?
 - a probe on a function which calls another function with another probe attached – don't do this!
 - within a probe, interrupt fires in which the interrupt handler calls the function that has a probe attached – random misses!

Example 1:

```
do_sys_openat2() [kprobe attached, in progress]
└ bpf_trace_printk()
  └ spin_lock(trace_printk_lock) [contention occurs]
    └ contention_begin [tracepoint BPF prog fires]
      └ bpf_trace_printk()
        └ spin_lock(trace_printk_lock)
          └ contention_begin [tracepoint executes again]
            └ SKIPPED (nmissed++)
```

Example 2:

```
do_sys_openat2() [probed, in progress]
└ IRQ/NMI fires
  └ IRQ handler calls do_sys_openat2()
    └ SKIPPED (nmissed++)
```

kprobes execution misses



- kprobe misses are handled at two different layers
 - handler / attach layer
 - kprobe-specific recursion check
 - BPF program execution (JIT)
 - per CPU recursion check

kprobe: handler / attach layer



- **3 handlers**
 - **int3 (breakpoint)**, fallback default behavior with:
 - CONFIG_KPROBES_ON_FTRACE=n
 - and no optimization

kprobe: handler / attach layer



- **3 handlers**
 - **int3 (breakpoint)**, fallback default behavior with:
 - CONFIG_KPROBES_ON_FTRACE=n and no optimization
 - **kprobe** ftrace handler relies on:
 - CONFIG_KPROBES_ON_FTRACE=y

kprobe: handler / attach layer



- **3 handlers**
 - **int3 (breakpoint)**, fallback default behavior with:
 - CONFIG_KPROBES_ON_FTRACE=n and no optimization
 - **kprobe** ftrace handler relies on:
 - CONFIG_KPROBES_ON_FTRACE=y
 - **"opt"**
 - CONFIG_KPROBES_ON_FTRACE=n
 - CONFIG_OPTPROBES=y (automatically enabled x86/x86-64 & non-preemptive kernel)
 - "debug.kprobes_optimization" sysctl = 1

Handler	Optimized
int3 (breakpoint)	No
kprobe ftrace, opt	Yes

kprobe: handler / attach layer

- **3 handlers**

- **int3 (breakpoint)**, fallback default behavior with:
 - CONFIG_KPROBES_ON_FTRACE=n and no optimization
- **kprobe** ftrace handler relies on:
 - CONFIG_KPROBES_ON_FTRACE=y
- **"opt"**
 - CONFIG_KPROBES_ON_FTRACE=n
 - CONFIG_OPTPROBES=y (automatically enabled x86/x86-64 & non-preemptive kernel)
 - "debug.kprobes_optimization" sysctl = 1

Handler	Optimized
int3 (breakpoint)	No
kprobe ftrace, opt	Yes

```
int handler(...) {  
    ...  
    if (kprobe_running()) {  
        kprobes_inc_nmissed_count(...);  
    }  
    ...  
}  
  
static inline struct kprobe *kprobe_running(void)  
{  
    return __this_cpu_read(current_kprobe);  
}
```

```
DECLARE_PER_CPU(struct kprobe *, current_kprobe);
```

kprobe: BPF program execution layer



- once the handlers proceed to invoke the BPF program, additional checks are made
- `trace_call_bpf()`
 - called for kprobes programs
 - checks if there is any BPF program running on the same CPU as the kprobe program

```
if (unlikely(__this_cpu_inc_return(bpf_prog_active) != 1)) {
    /*
     * since some bpf program is already running on this cpu,
     * don't call into another bpf program (same or different)
     * and don't send kprobe event into ring-buffer,
     * so return zero here
    */
    ...
    bpf_prog_inc_misses_counters(...);
    ...
}
```

fprobe: handler layer




- handler uses `ftrace_test_recursion_trylock()`
 - allows 1 level of nesting per CPU, per execution context

```
int handler(...) {  
    ...  
    if (ftrace_test_recursion_trylock(...) < 0) {  
        return;  
    }  
    ...  
}
```

fprobe: handler layer



- handler uses `ftrace_test_recursion_trylock()`
 - allows 1 level of nesting per CPU, per execution context




```
int handler(...) {  
    ...  
    if (ftrace_test_recursion_trylock(...) < 0) {  
        return;  
    }  
    ...  
}
```


fprobe: handler layer



- handler uses `ftrace_test_recursion_trylock()`
 - allows 1 level of nesting per CPU, per execution context



```
int handler(...) {  
    ...  
    if (ftrace_test_recursion_trylock(...) < 0) {  
        return;  
    }  
    ...  
}
```




fprobe: BPF program execution layer



- BPF trampoline-based
- two entrypoints for trampoline
 - `__bpf_prog_enter_recur()`
 - `__bpf_prog_enter_sleepable_recur()`
- per program check results in execution misses only for self-recursion

```
/* At BPF program execution layer */
if (unlikely(this_cpu_inc_return(*(prog->active)) != 1))
{
    bpf_prog_inc_misses_counter(prog);
}
```



kprobe & fprobe comparison



- notice the difference between the checks:
 - `prog->active` and `bpf_prog_active` (from previous slide)

Feature	Kprobes	Fprobes
Hook type basis	Primarily interrupt-based (int3) or optimized (jump ins.)	BPF trampoline-based
Handler layer check	Misses if any other kprobe is running	Allows 1 level of execution nesting/recursion, misses on further levels
BPF execution layer check	Misses if any other BPF program is running on the same CPU	Misses only on self-recursion (exact same fprobe program)
Nesting/recursion tolerance	No nesting or recursion of any BPF program on the same CPU	Allows nesting of different BPF programs on the same CPU. Only prevents self-recursion

- **tracepoints**: similar to kprobes (per CPU check) except without the kprobe-specific restrictions
- **raw tracepoints**: similar to fprobes (per program check) since they are both trampoline-based except without the fprobe-specific restrictions

BPF LSM hooks



```
SEC("kprobe/do_sys_openat2")
int BPF_KPROBE(kprobe_do_sys_openat2, ...) {
    ...
}
```

```
SEC("fentry/do_sys_openat2")
int BPF_PROG(fentry_do_sys_openat2, ...) {
    ...
}
```

```
SEC("lsm/file_open")
int BPF_PROG(lsm_file_open, struct file *file, ...)
{
    ...
}
```



BPF LSM hooks



```
SEC("kprobe/do_sys_openat2")
int BPF_KPROBE(kprobe_do_sys_openat2, ...) {
    ...
}
```

```
SEC("fentry/do_sys_openat2")
int BPF_PROG(fentry_do_sys_openat2, ...) {
    ...
}
```

```
SEC("lsm/file_open")
int BPF_PROG(lsm_file_open, struct file *file, ...)
{
```



-> otherwise, use tracepoints as they are the highest performance tracing hook type

hookpoint missed execution summary



Hookpoint	kprobe	fprobe / trampoline	tracepoint / perf event	raw tracepoint
When misses occur	Per CPU: <ul style="list-style-type: none">• Any nested kprobes• Nested BPF programs	Per program: <ul style="list-style-type: none">• Self nested fprobe	Per CPU: <ul style="list-style-type: none">• Nested BPF programs	Per program: <ul style="list-style-type: none">• Nested BPF programs

Key Takeaways



Generic

Performance

- Use tracepoints where possible

Kernel versions

- Use BTF and CORE

System Architecture

- use fprobe or tracepoint (no raw) instead of kprobe
- libbpf helpers
→ same code, ≠ architectures

kprobe/uprobe

Dynamic Links

- Library updated → offsets change
→ probe breaks silently
- Re-attach probes after library updates or use tracepoints

Selective Inlining

- Symbol exists but some call sites are inlined
- Solution: Trace the caller function or use offset

Partial Inlining

- Compiler splits function into fast/slow paths, missing probes
- Solution: kprobe: probe the .part suffix function, uprobe: compile with no partial inlining

Missed Executions

kprobes

- Cannot recurse
- Use optimized kprobes if necessary, otherwise prefer fprobes or tracepoints


fprobes, tracepoints

- Limited recursion

kprobes, fprobes, tracepoints

- If cannot tolerate any missed executions, consider BPF LSM.

Thanks, Credits & Resources



 **chris** 7 h
I'm curious if folks have run into any interesting eBPF hookpoint (kprobe, fprobe, tracepoint, cgroup, etc) gotchas or limitations that were surprising that you know of. Putting together a talk for FOSDEM (with [@Donia](#)) so wanted to see if folks had any interesting real world examples to share.
Happy to make the proper attribution / credit as well.
 **+5** **46 réponses** Dernière réponse aujourd'hui à 12 h 50

🙏 Thanks (in order of responses in the Slack thread) Martynas Pumputis, Mahé Tardy, Paul Chaignon, Daniel Borkman, Dylan Reimerink, Jiri Olsa, Korniliос Kourtis, Kev Sheldrake.

🙏 Thanks as well to Masami Hiramatsu.

LPC Talks:

- Kernel func tracing in the face of compiler optimization (<https://www.youtube.com/watch?v=kOYEsChbw-0>)
- Where have all the kprobes gone (<https://www.youtube.com/watch?v=Erqy3rxDp4g>)

Articles:

- Bouncing on trampolines to run eBPF programs (<https://bootlin.com/blog/bouncing-on-trampolines-to-run-ebpf-programs/>)
- eBPF Tracepoints, Kprobes, or Fprobes: Which One Should You Choose? (<https://labs.iximiuz.com/tutorials/ebpf-tracing-46a570d1>)
- An introduction to KProbes (<https://lwn.net/Articles/132196/>)
- Linux Tracing Technologies Guide (<https://docs.kernel.org/trace/>)

LKML:

- kprobe: Support nested kprobes (<https://lwn.net/ml/linux-kernel/158894789510.14896.13461271606820304664.stgit@devnote2/>)