

eBPF with Nix: laptop to testbed

Yifei Sun - PhD Student at

Inria, ENS de Lyon, Université Grenoble Alpes




- Website: <https://0.0.9.5.f.2.0.6.2.ip6.arpa>
- Email: fosdem@0.0.9.5.f.2.0.6.2.ip6.arpa

Goals

- DevOps
-

Goals

- DevOps
- Demonstrate how I use Nix to do BPF related work



Background

I started a project

- Multicast caching system for networked FS over XDP
-
-

Background

I started a project

- Multicast caching system for networked FS over XDP
- Its running late
-

Background

I started a project

- Multicast caching system for networked FS over XDP
- Its running late
- So here I am...

Background

Nix



- Declarative & functional
-
-
-

¹<https://zero-to-nix.com/concepts/derivations/>

²<https://zero-to-nix.com/concepts/closures/>

Background

Nix



- **Declarative** & functional
- Source code → Derivations¹ → Closure²
-
-

¹<https://zero-to-nix.com/concepts/derivations/>

²<https://zero-to-nix.com/concepts/closures/>

Background

Nix



- **Declarative** & functional
- Source code → Derivations¹ → Closure²
- NixOS: operating system as closure
-

¹<https://zero-to-nix.com/concepts/derivations/>

²<https://zero-to-nix.com/concepts/closures/>

Background

Nix



- **Declarative** & functional
- Source code → Derivations¹ → Closure²
- NixOS: operating system as closure
- Nix and NixOS devroom @ UA2.118 (Henriot)

¹<https://zero-to-nix.com/concepts/derivations/>

²<https://zero-to-nix.com/concepts/closures/>

Background


Testbed (Grid'5000 @ SLICES-FR)



slices RI

- Academic HPC cluster, reservation required
- **Ephemeral** bare metal machines

Background



Problem

- Environment setup and collaboration
 - Headers, compiler, editor...
 - KConfig, QEMU, ... (what if multiple machines are needed?)
-
-

Problem

- Environment setup and collaboration
 - Headers, compiler, editor...
 - KConfig, QEMU, ... (what if multiple machines are needed?)
- Development, deployment and benchmark
 - Cluster boot, data collection, ...
-

Problem

- Environment setup and collaboration
 - Headers, compiler, editor...
 - KConfig, QEMU, ... (what if multiple machines are needed?)
- Development, deployment and benchmark
 - Cluster boot, data collection, ...
- Peer review
 - Reproducing benchmark results

What worked for me

NixOS VM tests

- Basically Nix + Python + QEMU
- Multi-machines, different kernels, networking
- Binary cache, and benefits from using Nix

NixOS Compose

- Multi-flavor deployment tool for **ephemeral** experiments
 - ▶ systemd-nspawn
 - ▶ Docker
 - ▶ Bare metal
 - ▶ ...
- Substitute with your own stuff

Userspace tooling

Pull packages from pinned nixpkgs

```
devShells.x86_64-linux.default = pkgs.mkShell {  
    inputsFrom = [ <derivations> ];  
    packages = with pkgs.llvmPackages; [ clang-unwrapped libllvm ];  
};
```

- Compilers
- Libraries
- ...

Get a kernel

```
kernel = {  
    version = "6.19.0-rc5+multikernel";  
    modDirVersion = "6.19.0-rc5";  
    stdenv = pkgs.gcc13Stdenv;
```

Get a kernel

```
kernel = {
    version = "6.19.0-rc5+multikernel";
    modDirVersion = "6.19.0-rc5";
    stdenv = pkgs.gcc13Stdenv;

    # or ./ or fileset ...
    src = fetchFromGitHub {
        owner = "multikernel";
        repo = "linux";
        rev = "a3b4530cc04fe16ddef6b251baac488df3cae79";
        hash = "sha256-mum7rTLU5xUS2qex7br+EotjPyp0...";
    };
}
```

Get a kernel

```
kernel = {
    version = "6.19.0-rc5+multikernel";
    modDirVersion = "6.19.0-rc5";
    stdenv = pkgs.gcc13Stdenv;

    # or ./ or fileset ...
    src = fetchFromGitHub {
        owner = "multikernel";
        repo = "linux";
        rev = "a3b4530cc04fe16ddef6b251baac488df3cae79";
        hash = "sha256-mum7rTLU5xUS2qex7br+EotjPyp0...";
    };

    kernelPatches = [ ... ];
}
```

Get a kernel

```
kernel = {
    version = "6.19.0-rc5+multikernel";
    modDirVersion = "6.19.0-rc5";
    stdenv = pkgs.gcc13Stdenv;

    # or ./ or fileset ...
    src = fetchFromGitHub {
        owner = "multikernel";
        repo = "linux";
        rev = "a3b4530cc04fe16ddef6b251baac488df3cae79";
        hash = "sha256-mum7rTLU5xUS2qex7br+EotjPyp0...";
    };

    kernelPatches = [ ... ];

    structuredExtraConfig = {
        MULTIKERNEL = lib.kernel.yes;
    };
};
```

Get a kernel

```
kernel = {
    version = "6.19.0-rc5+multikernel";
    modDirVersion = "6.19.0-rc5";
    stdenv = pkgs.gcc13Stdenv;

    # or ./ or fileset ...
    src = fetchFromGitHub {
        owner = "multikernel";
        repo = "linux";
        rev = "a3b4530cc04fe16ddef6b251baac488df3cae79";
        hash = "sha256-mum7rTLU5xUS2qex7br+EotjPyp0...";
    };

    kernelPatches = [ ... ];

    structuredExtraConfig = {
        MULTIKERNEL = lib.kernel.yes;
    };
};

boot.kernelPackages = pkgs.linuxPackagesFor (
    pkgs.callPackage (
        { buildLinux, fetchFromGitHub, ... } @ args:
        buildLinux (
            args
            //
            kernel # <-
            //
            (args.argsOverride or { })
        )
    )
    { }
);
}
```

One machine

```
pkgs testers .runNixOSTest {  
    name = "one-machine-test";
```

- Boilerplate

One machine

```
pkgs testers.runNixOSTest {  
    name = "one-machine-test";  
  
    nodes.machine1 = {  
        imports = [ nixosModules.kernel ];  
        services.scx.enable = true;  
    };
```

- Boilerplate
- Declarative NixOS closure generation

One machine

```
pkgs testers.runNixOSTest {  
    name = "one-machine-test";  
  
    nodes.machine1 = {  
        imports = [ nixosModules.kernel ];  
        services.scx.enable = true;  
    };  
  
    testScript = ''  
        machine1.wait_for_unit("default.target")  
        machine1.succeed("")  
        machine1.fail("")  
    '';  
}
```

- Boilerplate
- Declarative NixOS closure generation
- Imperative Python stmts to invoke tests

More machines?

```
pkgs.testers.runNixOSTest {  
    name = "lots-of-machine-test";  
  
    nodes.machine1.imports = [ nixosModules.grafana ];  
    nodes.machine2.imports = with nixosModules; [  
        kernel exporter ebpf benchmark  
    ];  
  
    testScript = ''  
        start_all()  
        machine1...  
        machine2...  
    '';  
}
```

What's in there?

```
nix-repl> test =  
pkgs.testers.runNixOSTest { ... }  
nix-repl> :p test.  
test.config           test.name  
test.driver          test.nodes  
test.driverInteractive ...
```

What's in there?

```
nix-repl> test =  
pkgs.testers.runNixOSTest { ... }  
nix-repl> :p test.  
test.config           test.name  
test.driver          test.nodes  
test.driverInteractive ...
```

Node closure:

```
<test>.nodes.<name>.system.build.toplevel
```

What's in there?

```
nix-repl> test =  
pkgs.testers.runNixOSTest { ... }  
nix-repl> :p test.  
test.config           test.name  
test.driver          test.nodes  
test.driverInteractive ...
```

Node closure:

```
<test>.nodes.<name>.system.build.toplevel
```

Driver:

```
<test>.driver (run testScript)
```

```
<test>.driverInteractive (Python repl)
```

What's in there?

```
nix-repl> test =  
pkgs.testers.runNixOSTest { ... }  
nix-repl> :p test.  
test.config           test.name  
test.driver          test.nodes  
test.driverInteractive ...
```

Python test driver

- Nodes
 - qemu
- VLANs
 - vde_switch

Node closure:

```
<test>.nodes.<name>.system.build.toplevel
```

Driver:

```
<test>.driver (run testScript)
```

```
<test>.driverInteractive (Python repl)
```

Mock syscall

Say we want to troll ourselves:

```
SEC("ksyscall/statx")
int BPF_KSYSCALL(fsd_statx_entry, ... statx(2) args) {
    // generate a map entry to collect start timestamp
    // check path, if not match return
    // else override with a static statx content
    struct statx stx = { ... };
    bpf_probe_write_user(statxbuf, &stx, sizeof(stx));
    return bpf_override_return(ctx, 0);
}
```


And count how many times we can footgun ourselves

With a counter and a histogram

Declarative userspace program

- Auto-load the the program (feat. `ebpf_exporter`)
- Collect the metrics and plot them (feat. Prometheus & Grafana)
- Local testing (feat. NixOS VM test)
- Deployment (feat. NixOS-Compose)

How?



Local testing

For simplicity

- We will be using a readily available userspace tool
 - Loading the program
 - Read the map and re-expose the content over Prometheus

Complication is fast

- Build once and its immutable
- Push cache to server (or have a CI server build it)
- SBOM

Debugging is easy

- SSH backdoor enable with a knob

Demo

- Build interactive driver closure

```
nom build .#checks.x86_64-linux.default.driverInteractive
```

- Start the driver

```
$ ./result/bin/nixos-test-driver
start vlan
running vlan (pid 3859017; ctl /run/user/1000/vde1.ctl)
SSH backdoor enabled, the machines can be accessed like this:
  collector: ssh -o User=root vsock/3
  exporter:  ssh -o User=root vsock/4
```

Straight to prod

Bit-perfect reproducibility (*: for some store paths)

Everything is in closure

- Deployment harness is easy to write

Demo

- Build deployment closure (instrumented with NixOS test)

```
nxc build
```

- Schedule couple machines and deploy the closure to cluster

Conclusion

Less than 250 LoC (Nix)

- Portable modules
- Composable with other services
- Adding new programs to deployment only adds a couple characters

```
services.prometheus.exporters.ebpf = {  
    enable = true;  
    # bpf object file names  
    names = [  
        "oomkill"  
        "softirq-latency"  
        ...  
    ];  
};
```

Questions?

- Code: git.sr.ht/~stepbrobd/fosdem
- <https://team.inria.fr/datamove>
- <https://numpex.org>
- Website: <https://0.0.9.5.f.2.0.6.2.ip6.arpa>
- Email: fosdem@0.0.9.5.f.2.0.6.2.ip6.arpa

Our team is hiring!