Full disk encryption for Confidential Computing guests

FOSDEM 2026

Emanuele Giuseppe Esposito

Vitaly Kuznetsov

Red Hat

Confidential computing

Confidential VMs ares here to stay, you can easily get one from AWS/Google/Azure or deploy your own KVM + SEV-SNP/TDX

- TDX, SEV-SNP reliably protect the data in use.
- Protecting the storage in left to the guest operating system.

Confidential storage

Confidential storage must provide:

- Verity protection for readonly parts.
- Encryption + integrity protection for read-write parts.
- Attestability.
- Roll-back attack protection.

Verity protection

- Dm-verity is great!
  - integrity checking for block devices
- Rich support in systemd:
  - repart, veritysetup-generator, dissect,...
  - ***roothash=<hash>*** for full root protection in read-only or with ephemeral overlay (systemd.volatile=overlay)
  - ***usrhash=<hash>*** for /usr protection (aka hermetic /usr approach)
- Downside: verity target partition is RO!

Encryption

- LUKS is great! Partition is encrypted and RW!
- Each VM instance/volume needs to be individually encrypted in a safe environment:
  - Pre-encrypted by some ‘trusted’ infrastructure.
  - Self-encryption e.g. upon the first boot.

Encryption and systemd-repart

- Systemd-repart can create encrypted volumes and seal keys to the vTPM
  - Integrity protection for LUKS (experimental) is coming in v260
    - Full roll-back attack always remains possible.
  - A new feature to ensure that the created encrypted volume will not change before (first) usage.

Combining encryption with verity

Verity -> Encryption switch can give us ‘read-write experience’:

- Copy everything from verity-protected volume to the encrypted one.
- Use filesystem overlay (overlayfs) on the encrypted volume over verity-protected data.
- Use dm-clone and transfer verity-protected data to the encrypted volume.

Verity + encrypted overlay

- ***roothash=<hash>***
  - Root is found via roothash, verified and mounted
  - Currently, it can't be combined with a persistent overlay with standard systemd tooling to give read-write experience.
- ***usrhash=<hash>:*** a read-write solution is coming
  - /usr is separate, dm-verity protected partition
  - Root is created by systemd-repart and encrypted
  - /usr is mounted ro (verity) and an overlay makes it rw

Dm-clone approach

- *Almost* no influence on the boot time (depends on the hydration parameters).
- Can simplify things in case of complex storage configurations, e.g. LVM.
- Potentially allows to minimize storage use by dropping dm-verity protected data after full convergence.
- No support in the standard tooling but a solution is being work on.

Encryption and attestation

- Systemd already measures a derivation of LUKS volume key to PCR15.
  - *... alongside randomly generated machine-id*
- For self-encryption, a proof that the volume was created in a safe environment (e.g. on the first boot) is also needed.
  - See the proposal!

EFI system partition

- Cannot be verity protected and/or encrypted.
- SecureBoot keys (+ Measured boot) need to be trusted for ESP artifacts:
  - UKI
  - Cmdline extensions
  - Systemd sysext/confext

Using distro-shipped UKIs

- Expected verity hash (***roothash=.../ usrhash=...)*** can be supplied with a cmdline extension.
- Encryption logic can be a systemd sysext/confext:
  - Distros using dracut may need a new feature

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.


linkedin.com/company/red-hat


youtube.com/@redhat


facebook.com/RedHat


x.com/RedHat

Red Hat