

Generating Programmable NPUs from Linalg with MLIR and CIRCT

Josse Van Delm

LLVM dev room
FOSDEM 2026

whoami

Josse Van Delm

PhD Researcher at KU Leuven - MICAS research group



KU LEUVEN

Working on compilers for AI hardware accelerators since 2021

Fun fact! Looking for a new job at the end of this year 😊

What you'll learn in this talk

- *What are NPUs?*
- *How do people design NPUs now?*
- *How people should design NPUs in the future (spoiler: with compilers!)*

NPUs are processors to optimize dense linear algebra

```
void elementwise_mult(float* A, float* B, float* C) {
    for (int i = 0; i < 16; i++) {
        C[i] = A[i] * B[i];
    }
}
```

NPUs are processors to optimize dense linear algebra

```
void elementwise_mult(float* A, float* B, float* C) {  
    for (int i = 0; i < 16; i++) {  
        C[i] = A[i] * B[i];  
    }  
}
```

NPUs are processors to optimize dense linear algebra

```
void elementwise_mult(float* A, float* B, float* C) {  
    for (int i = 0; i < 16; i++) {  
        C[i] = A[i] * B[i];  
    }  
}  
  
void matmul_16×16(float A[16][16], float B[16][16], float C[16][16]) {  
    for (int i = 0; i < 16; i++) {  
        for (int j = 0; j < 16; j++) {  
            C[i][j] = 0;  
            for (int k = 0; k < 16; k++) {  
                C[i][j] += A[i][k] * B[k][j];  
            }  
        }  
    }  
}
```

NPUs are processors to optimize dense linear algebra


```
void softmax_16(float* input, float* output) {
    float sum = 0.0f;
    float max_val = input[0];
    for (int i = 1; i < 16; i++) {
        if (input[i] > max_val) max_val = input[i];
    }
    for (int i = 0; i < 16; i++) {
        output[i] = expf(input[i] - max_val);
        sum += output[i];
    }
    for (int i = 0; i < 16; i++) {
        output[i] /= sum;
    }
}
```

Aside: MLIR represents this in the “linalg” dialect

```
void elementwise_mult(float* A, float* B, float* C) {
    for (int i = 0; i < 16; i++) {
        C[i] = A[i] * B[i];
    }
}

%3 = linalg.mul ins(%arg0, %2 : tensor<16xf32>, tensor<16xf32>) outs(%0 : tensor<16xf32>) →
tensor<16xf32>


#map = affine_map<(d0) → (d0)>
%3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
ins(%arg0, %2 : tensor<16xf32>, tensor<16xf32>) outs(%0 : tensor<16xf32>) {
    ^bb0(%in: f32, %in_0: f32, %out: f32):
        %5 = arith.mulf %in, %in_0 : f32
        linalg.yield %5 : f32
} → tensor<16xf32>
```



[1]



[2]




[3]




[4]

Are they useful? According to research: yes!



There is more available silicon area on chip for specialized processors




TPU = much faster and energy efficient at linear algebra than CPUs or GPUs

But are they really useful?


- Yes. ... but mostly for marketing
 - In reality it is often too difficult to use them
- These units can be too **specific** for newer workloads
- Most compilers can't leverage them
- Hardware, Firmware, Libraries often **closed source!**

There are multiple interdependent problems




How did we get here?

How do people typically Create NPUs?



How do people typically Create NPUs?



This is too slow! Can we do this faster?

Yes! With compilers we can enable **programmable hardware synthesis**

But how?

- How do we quickly create an ML compiler?
- How do we quickly generate hardware?
- How do we adapt the compiler to the hardware?
- How do we quickly develop this?



MLIR



CIRCT



Dynamic Backend




Secret shortcuts


Fully Open source A blue GitHub logo icon.

**What are we trying to
create?**


Anatomy of an NPU platform




The programmable hardware synthesis flow




The programmable hardware synthesis flow




The programmable hardware synthesis flow




The programmable hardware synthesis flow




The programmable hardware synthesis flow



Today we'll talk about PE array conversion



Today we'll talk about PE array conversion




Create a PE Array for these two supported operations

Supported Operations (MLIR)




PE Array
(SystemVerilog)


Create a PE Array for these two supported operations




Create a PE Array for these two supported operations




Create a PE Array for these two supported operations



This NPU now supports `linalg.generic A and B`



If the source contains supported operations, they are guaranteed to work on this NPU!



Examples

Simple example no. 1

Create a 4-way parallel PE array that supports Elementwise operations:


- Addition
- Subtraction
- Multiplication
- Bitwise Exclusive OR

We want to support the following operations


```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
    %0 = tensor.empty() : tensor<16xi64>
    %1 = linalg.add ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) →
    tensor<16xi64>
    %2 = linalg.sub ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %3 = linalg.mul ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]} ins(%arg0, %3 :
    tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
    return %4 : tensor<16xi64>
  }
}
```

Create a single PE for each supported operation


```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
    %0 = tensor.empty() : tensor<16xi64>
    %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
      ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
        ^bb0(%in: i64, %in_0: i64, %out: i64):
          %5 = arith.addi %in, %in_0 : i64
          linalg.yield %5 : i64
      } → tensor<16xi64>
    %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
      ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
        ^bb0(%in: i64, %in_0: i64, %out: i64):
          %5 = arith.subi %in, %in_0 : i64
          linalg.yield %5 : i64
      } → tensor<16xi64>
    %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
      ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
        ^bb0(%in: i64, %in_0: i64, %out: i64):
          %5 = arith.muli %in, %in_0 : i64
          linalg.yield %5 : i64
      } → tensor<16xi64>
    %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
      ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
        ^bb0(%in: i64, %in_0: i64, %out: i64):
          %5 = arith.xori %in, %in_0 : i64
          linalg.yield %5 : i64
      } → tensor<16xi64>
    return %4 : tensor<16xi64>
  }
}
```



Aggregate PEs into one



Create a 4-way array



Now deploy the algorithm on simple accelerator no.1

Create a 4-way parallel PE array that supports Elementwise operations:

- Addition
- Subtraction
- Multiplication
- Bitwise Exclusive OR

Deploy this on the simple 4-way array

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
    %0 = tensor.empty() : tensor<16xi64>
    %1 = linalg.add ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) →
    tensor<16xi64>
    %2 = linalg.sub ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %3 = linalg.mul ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]} ins(%arg0, %3 :
    tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
    return %4 : tensor<16xi64>
  }
}
```

Deploy addition

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy addition

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy subtraction

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy subtraction

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy multiplication

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy multiplication

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```




Deploy exclusive or

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```



Deploy exclusive or

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64>
  %0 = tensor.empty() : tensor<16xi64>
  %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.addi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.subi %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %3 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.muli %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
    ins(%arg0, %3 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
  return %4 : tensor<16xi64>
}
```



Performance

Load data * 4

Store data * 4

Program accelerator * 4

Simple example no. 2 operator fusion

We can make this much faster if we support operator fusion

- Addition
- Subtraction
- Multiplication
- Bitwise Exclusive OR

$$(A \wedge (A * (A - (A + B))))$$

Deploy this on the simple 4-way array

```
#map = affine_map<(d0) → (d0)>
module {
  func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
    %0 = tensor.empty() : tensor<16xi64>
    %1 = linalg.add ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) →
    tensor<16xi64>
    %2 = linalg.sub ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %3 = linalg.mul ins(%arg0, %2 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) → tensor<16xi64>
    %4 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]} ins(%arg0, %3 :
    tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
      ^bb0(%in: i64, %in_0: i64, %out: i64):
        %5 = arith.xori %in, %in_0 : i64
        linalg.yield %5 : i64
    } → tensor<16xi64>
    return %4 : tensor<16xi64>
  }
}
```

Deploy this on the simple 4-way array

```
#map = affine_map<(d0) → (d0)>
module {
    func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
        %0 = tensor.empty() : tensor<16xi64>
        %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]} ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
            ^bb0(%in: i64, %in_0: i64, %out: i64):
                %2 = arith.addi %in, %in_0 : i64
                %3 = arith.subi %in, %2 : i64
                %4 = arith.muli %in, %3 : i64
                %5 = arith.xori %in, %4 : i64
                linalg.yield %5 : i64
        } → tensor<16xi64>
        return %1 : tensor<16xi64>
    }
}
```

Processing Element
($A^A * (A - (A + B))$)


Processing Element
($A^A * (A - (A + B))$)

Processing Element
($A^A * (A - (A + B))$)

Processing Element
($A^A * (A - (A + B))$)

Deploy this on the simple 4-way array

```
#map = affine_map<(d0) → (d0)>
module {
    func.func public @streamer_add(%arg0: tensor<16xi64>, %arg1: tensor<16xi64>) → tensor<16xi64> {
        %0 = tensor.empty() : tensor<16xi64>
        %1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]} ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
            ^bb0(%in: i64, %in_0: i64, %out: i64):
                %2 = arith.addi %in, %in_0 : i64
                %3 = arith.subi %in, %2 : i64
                %4 = arith.muli %in, %3 : i64
                %5 = arith.xori %in, %4 : i64
                linalg.yield %5 : i64
        } → tensor<16xi64>
        return %1 : tensor<16xi64>
    }
}
```



Performance

Store data * 1

Load data * 1

Program accelerator * 1

Discussion

Tradeoff - More specific accelerators or more performance?

How did we create this?


Shortcut no. 1 - xDSL for quick prototyping

xDSL, an open-source MLIR, CIRCT compatible framework in Python, to quickly prototype new MLIR dialects and combinations of compiler passes.



github.com/xdslproject/xdsl

Shortcut no. 2 - PE array + SNAX platform = NPU



Shortcut no. 3 - SNAX-MLIR sw compilation toolchain

- Compatible with RISC-V control core (rv32imc)
- `linalg.generic` → loads/stores + program accelerators with dynamic backends

github.com/KULeuven-MICAS/snax-mlir

In the future

- Overlay functionality on parts of the PE
- Support for multiple dataflows
- Multi-accelerator support
- Support for multi-cycle hardware (ready/valid)
- Floating-point support (not because I like FP, rather because I don't like quantization)

Thank you! Questions?

josse:matrix.org

[linkedin.com/in/jossevandelm/](https://www.linkedin.com/in/jossevandelm/)

github.com/JosseVanDelm

vandelm.com

Takeaways of this talk

- *NPUs* are dedicated processors to accelerate **dense linear algebra**
- Properly using and and designing NPUS them requires an **overhaul** of the **design process**
- We leverage *MLIR*, *CIRCT*, *xDSL* and *SNAX* to quickly create new NPUs and compilers together through **programmable hardware synthesis**

References:

- [1] “TPU v4” by Norman P. Jouppi, George Kurian, Sheng Li, Peter Ma, Rahul Nagarajan, Lifeng Nai, Nishant Patil, Suvinay Subramanian, Andy Swing, Brian Towles, Cliff Young, Xiang Zhou, Zongwei Zhou, and David Patterson licensed under [Creative Commons Attribution 4.0 International license](#)
- [2] “Apple M1” by Henriok licensed under the [Public Domain](#)
- [3] “AMD BC-160” by Мой Компьютер licensed under [Creative Commons Attribution 3.0 Unported license](#)
- [4] “NVIDIA H100” by Geekerwan licensed under [Creative Commons Attribution 3.0 Unported](#)
- [5] “Microprocessor Trend Data” by “Karl Rupp” licensed under [Creative Commons Attribution 4.0 International Public License](#)
- [6] Jouppi, Norman P., et al. "In-datacenter performance analysis of a tensor processing unit." Proceedings of the 44th annual international symposium on computer architecture. 2017.

Creating the hardware with the PHS dialect

A new dialect to create **programmable*** hardware

phs.pe:

An operation that represents what a PE of the NPU can do

phs.choose:

A *named* operation that represents a choice of operation in different MLIR operations

phs.mux:

“multiplexer” - an *unnamed* operation that represents a choice of dataflow for the PE

phs.yield:

a terminator operation similar to `scf.yield` and `linalg.yield`.

To go from `linalg` to a programmable PE

1. Convert each `linalg` to a `linalg.generic`
2. Convert each `linalg.generic` to a `phs.pe` op
3. Aggregate all `phs.pe` ops through their named `phs.choose_ops` and insert `phs.mux_ops` where necessary
4. Generate an array of `phs.pe` elements
5. Convert the aggregated `phs.pe` to systemverilog with CIRCT

2) Convert each `linalg.generic` to a `phs.pe` op


```
%1 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
ins(%arg0, %arg1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
^bb0(%in: i64, %in_0: i64, %out: i64):
  %5 = arith.addi %in, %in_0 : i64
  linalg.yield %5 : i64
} → tensor<16xi64>
```



```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {
%1 = phs.choose @i_i64_i64_o_i64_0 with %0 (%in : i64, %in_1 : i64) → i64
  0) {
  %2 = arith.addi %in, %in_1 : i64
  phs.yield %2 : i64
}
phs.yield %1 : i64
```

2) Convert each `linalg.generic` to a `phs.pe` op

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0 (%in : i64, %in_1 : i64) → i64
  0) {
    %2 = arith.addi %in, %in_1 : i64
    phs.yield %2 : i64
  }
  phs.yield %1 : i64
```



2) Convert each `linalg.generic` to a `phs.pe` op


```
%2 = linalg.generic {indexing_maps = [#map, #map, #map], iterator_types = ["parallel"]}
ins(%arg0, %1 : tensor<16xi64>, tensor<16xi64>) outs(%0 : tensor<16xi64>) {
^bb0(%in: i64, %in_0: i64, %out: i64):
  %5 = arith.subi %in, %in_0 : i64
  linalg.yield %5 : i64
} → tensor<16xi64>
```



```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0 (%in : i64, %in_1 : i64) → i64
  0) {
    %2 = arith.subi %in, %in_1 : i64
    phs.yield %2 : i64
  }
  phs.yield %1 : i64
```

2) Convert each `linalg.generic` to a `phs.pe` op

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0 (%in : i64, %in_1 : i64) → i64
  0) {
  %2 = arith.subi %in, %in_1 : i64
  phs.yield %2 : i64
}
phs.yield %1 : i64
```



3) Aggregate all **phs.pe** ops through their named **phs.choose_ops** and insert **phs.mux_ops** where necessary


```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {  
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0  
    (%in : i64, %in_1 : i64) → i64  
  0) {  
    %2 = arith.addi %in, %in_1 : i64  
    phs.yield %2 : i64  
  }  
  phs.yield %1 : i64
```

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {  
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0  
    (%in : i64, %in_1 : i64) → i64  
  0) {  
    %2 = arith.subi %in, %in_1 : i64  
    phs.yield %2 : i64  
  }  
  phs.yield %1 : i64
```

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {  
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0  
    (%in : i64, %in_1 : i64) → i64  
  0) {  
    %2 = arith.addi %in, %in_1 : i64  
    phs.yield %2 : i64  
  }  
  1) {  
    %3 = arith.subi %in, %in_1 : i64  
    phs.yield %3 : i64  
  }  
  phs.yield %1 : i64  
}
```


3) Aggregate all **phs.pe** ops through their named **phs.choose_ops** and insert **phs.mux_ops** where necessary

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {  
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0  
  (%in : i64, %in_1 : i64) → i64  
  0) {  
    %2 = arith.addi %in, %in_1 : i64  
    phs.yield %2 : i64  
  }  
  1) {  
    %3 = arith.subi %in, %in_1 : i64  
    phs.yield %3 : i64  
  }  
  phs.yield %1 : i64  
}
```



3) Aggregate all **phs.pe** ops through their named **phs.choose_ops** and insert **phs.mux_ops** where necessary

```
phs.pe @acc1 with %0 (%in : i64, %in_1 : i64) {  
  %1 = phs.choose @i_i64_i64_o_i64_0 with %0 (%in : i64, %in_1 : i64) → i64  
  0) {  
    %2 = arith.addi %in, %in_1 : i64  
    phs.yield %2 : i64  
  }  
  1) {  
    %3 = arith.subi %in, %in_1 : i64  
    phs.yield %3 : i64  
  }  
  2) {  
    %4 = arith.muli %in, %in_1 : i64  
    phs.yield %4 : i64  
  }  
  3) {  
    %5 = arith.xori %in, %in_1 : i64  
    phs.yield %5 : i64  
  }  
  phs.yield %1 : i64  
}
```



```

hw.module private @acc1(in %in data_0: i64, in %in_1 data_1: i64,
                      in %0 switch_0: i2, out out_0: i64) {
    %1 = arith.addi %in, %in_1 : i64
    %2 = arith.subi %in, %in_1 : i64
    %3 = arith.muli %in, %in_1 : i64
    %4 = arith.xori %in, %in_1 : i64
    %5 = hw.array_create %4, %3, %2, %1 : i64
    %6 = hw.array_get %5[%0] : !hw.array<4xi64>, i2
    hw.output %6 : i64
}

hw.module @acc1_array(in %0 data_0: !hw.array<4xi64>,
                      in %1 data_1: !hw.array<4xi64>,
                      in %2 switch_0: i2, out out_0:
                      !hw.array<4xi64>) {
    %3 = arith.constant 0 : i2
    %4 = hw.array_get %0[%3] : !hw.array<4xi64>, i2
    %5 = arith.constant 0 : i2
    %6 = hw.array_get %1[%5] : !hw.array<4xi64>, i2
    %7 = hw.instance "acc1_pe_0" @acc1(data_0: %4: i64,
                                         data_1: %6: i64, switch_0: %2: i2) → (out_0: i64)
    %8 = arith.constant 1 : i2
    %9 = hw.array_get %0[%8] : !hw.array<4xi64>, i2
    %10 = arith.constant 1 : i2
    %11 = hw.array_get %1[%10] : !hw.array<4xi64>, i2
    %12 = hw.instance "acc1_pe_1" @acc1(data_0: %9: i64,
                                         data_1: %11: i64, switch_0: %2: i2) → (out_0: i64)
    ...
    %23 = hw.array_create %22, %17, %12, %7 : i64
    hw.output %23 : !hw.array<4xi64>
}

```

4) Generate an array of phs.pe elements

