

# Using OpenMP's interop for calling GPU-vendor libs with GCC

Tobias Burnus



# About Me

&



Contributor to the **OpenMP specification**

Contributor to **GCC** related to  
GPU offloading, OpenMP, Fortran, ...  
(+ co-maintainer/reviewer)

Background in numerical/computational  
**physics** (time-dependent density-functional  
theory, then related to transition-metal  
oxides). Contribution to GCC's gfortran  
started back then.

[tburnus@baylibre.com](mailto:tburnus@baylibre.com)

Doing contract work related to

- open-source compilers, linker, debugger, ...
- Linux kernel, Zephyr, Android, Yocto, ...

*with staff and customers on three continents*

[baylibre.com/blog/](http://baylibre.com/blog/)





# GCC — OpenMP and Offloading

- Support for **OpenMP** and **OpenACC**
- Most OpenMP **5.2** features  
(but no OMPT and OMPD, yet)  
[gcc.gnu.org/projects/gomp/](http://gcc.gnu.org/projects/gomp/)  
[gcc.gnu.org/onlinedocs/libgomp/](http://gcc.gnu.org/onlinedocs/libgomp/)
- Offloading to **Nvidia** and **AMD GPUs**  
[gcc.gnu.org/wiki/Offloading](http://gcc.gnu.org/wiki/Offloading)  
For build + install infos; Linux distro ship  
Offload support in optional packages.
- Use **GCC 15** (released April 2025),  
[gcc.gnu.org/gcc-15/](http://gcc.gnu.org/gcc-15/)
- **GCC 16** (April 2026?)  
[gcc.gnu.org/gcc-16/changes.html](http://gcc.gnu.org/gcc-16/changes.html)

## OpenMP **interop**

- Since GCC 15 [CUDA, HIP, HSA] with most OpenMP 6.0 additions (e.g. Fortran support)

WIP/near-term plan (a bit for 16, most GCC 17)

- Reduce missing bits for OpenMP 5.x + bug fixes
- Offload-performance improvements
- OMPT

And minor 6.x additions

Contributors welcome – code, documentation, bug reports  
GSoC 2026 ?



## interop – Starting Point

- Existing code – possibly with OpenMP for host and/or GPU parallelization
- System with a GPU
- Hot code using FFT, (sparse)BLAS, LAPACK, random-number generation
  - To be calculated by vendor lib on the GPU

OpenMP's interop + variant functions:

- Make it easier to hide some complexity while coding
- Improve portability a bit
- Permit dependency handling of OpenMP code vs. vendor-lib code



# Vendor Lib Call Example

Example: CUDA call for  $Y = \alpha X + Y$  [scalar times vector + vector; BLAS' DaXpY]

```
// Alloc device mem + copy data
3  cudaMalloc ((void**) &d_X, N * sizeof(*X));
4  cudaMemcpy (d_X, X, N * sizeof(*X), cudaMemcpyHostToDevice);
...
// Create streaming object and handle
5  cudaStreamCreate (&stream);
6  cublasCreate (&handle);
7  cublasSetStream (handle, stream);

// Y = alpha * X + Y
8  cublasDaxpy (handle, N, &alpha, d_X, 1, d_Y, 1);

9  cublasDestroy (handle);
10 cudaStreamDestroy (stream);

// Copy Y - free X as no longer needed
11 cudaFree (d_X);
12 cudaMemcpy (Y, d_Y, N * sizeof(*X), cudaMemcpyDeviceToHost);
13 cudaFree (d_Y);
```



# Vendor Lib Call Example

Example: CUDA call for  $Y = \alpha X + Y$  [scalar times vector + vector; BLAS' DaXpY]

```
// Alloc device mem + copy data
3 cudaMalloc ((void**) &d_X, N * sizeof(*X));
4 cudaMemcpy (d_X, X, N * sizeof(*X), cudaMemcpyHostToDevice);
...
// Create streaming object and handle
5 cudaStreamCreate (&stream);
6 cublasCreate (&handle);
7 cublasSetStream (handle, stream);

// Y = alpha * X + Y
8 cublasDaxpy (handle, N, &alpha, d_X, 1, d_Y, 1);

9 cublasDestroy (handle);
10 cudaStreamDestroy (stream);

// Copy Y - free X as no longer needed
11 cudaFree (d_X);
12 cudaMemcpy (Y, d_Y, N * sizeof(*X), cudaMemcpyDeviceToHost);
13 cudaFree (d_Y);
```

**First step:** Use OpenMP for  
memory allocation, host-device transfer,  
and on-device initialization

→ reduce/localize vendor specific code.



# Vendor Lib Call Example

OpenMP – memory handling (+ on-device var initialization)

```
// Initialize X + copy Y to the device
#pragma omp target enter data map(alloc: X[:N]) map(to:Y[:N])
#pragma omp target
    for (size_t i = 0; i < N; i++)
        X[i] = 1.0;
// Get device ptr:
d_X = (double*) omp_get_mapped_ptr (X, omp_get_default_device());
...
cublasDaxpy (handle, N, &alpha, d_X, 1, d_Y, 1);
...

// Copy Y - free X as no longer needed
#pragma omp target exit data map(release: X) map(from: Y[:N])
```



# Vendor Lib Call Example

## OpenMP – memory handling (+ on-device var initialization)

```
// Initialize X + copy Y to the device
#pragma omp target enter data map(alloc: X[:N]) map(to:Y[:N])
#pragma omp target
    for (size_t i = 0; i < N; i++)
        X[i] = 1.0;
// Get device ptr:
d_X = (double*) omp_get_mapped_ptr (X, omp_get_default_device());
...
cublasDaxpy (handle, N, &alpha, d_X, 1, d_Y, 1);
...

// Copy Y - free X as no longer needed
#pragma omp target exit data map(release: X) map(from: Y[:N])
```

### Variants:

- Use unified shared memory
- `omp_target_alloc`  
use directly or associate with host variable
- OpenMP allocators  
pinned, device-accessible, managed memory, ...



# OpenMP→CUDA device number via interop

Previous slides: assumption they are the same. Proper way + *interop* intro

```
1  omp_interop_t obj;  
2  #pragma omp interop init(target, prefer_type("cuda") : obj) \  
   device(omp_get_default_device())  
  
  // Check result: got an interop object - and for CUDA:  
5  if (obj != omp_interop_none  
6      && omp_ifr_cuda == omp_get_interop_int (obj, omp_ipr_fr_id, nullptr))  
{  
7    int dev_num = omp_get_interop_int (obj, omp_ipr_device_num, nullptr);  
    // e.g. → 0  
9    char* str = omp_get_interop_str (obj, omp_ipr_fr_name, nullptr));  
    // → "cuda"  
}  
10 #pragma omp interop destroy(obj)
```



# OpenMP→CUDA device number via interop

Previous code assume they are the same. Proper way + *interop intro*

```
1  omp_interop_t obj;  
2  #pragma omp interop init(target, prefer_type("cuda") : obj) \  
   device(omp_get_default_device())  
  
  // Check result: got an interop object - and for CUDA:  
5  if (obj != omp_interop_none  
   && omp_ifr_cuda == omp_get_interop_int (obj, omp_ipr_fr_id, nullptr))  
{  
  7   int dev_num = omp_get_interop_int (obj, omp_ipr_device_num, nullptr);  
    // e.g. → 0  
  9   char* str = omp_get_interop_str (obj, omp_ipr_fr_name, nullptr));  
    // → "cuda"  
}  
10 #pragma omp interop destroy(obj)
```

Available foreign runtimes:  
→ depends on device/GPU + compiler

For GCC:

- Nvidia GPUs: **cuda**, cuda\_driver, hip
- AMD GPUs: **hip**, hsa

→ [OpenMP.org](http://OpenMP.org) for list + assoc data types



# OpenMP → Streaming Object

cudaStreamCreate + cudaStreamDestroy → OpenMP interop (targetsync)

```
omp_interop_t obj;
#pragma omp interop init(targetsync, prefer_type("cuda")): obj
cudaStream_t stream = (cudaStream_t)
    omp_get_interop_ptr (obj, omp_ipr_targetsync, nullptr);
...
cublasSetStream (handle, stream);
...
13 #pragma omp interop destroy(obj)
```



# OpenMP → Streaming Object

cudaStreamCreate + cudaStreamDestroy → OpenMP interop (targetsync)

```
omp_interop_t obj;
#pragma omp interop init(targetsync, prefer_type("cuda")): obj
cudaStream_t stream = (cudaStream_t)
    omp_get_interop_ptr (obj, omp_ipr_targetsync, nullptr);
...
cublasSetStream (handle, stream);
...
13 #pragma omp interop destroy(obj)
```

foreign-runtime-ids		data types
id	name	
1	cuda	<b>cudaStream_t</b>
2	cuda_driver	<b>CUstream</b>
3	opencl	<b>cl_queue</b>
4	sycl	<b>cl::sycl::queue</b>
5	hip	<b>hipStream_t</b>
6	level_zero	<b>ze_command_queue_handle_t</b>
7	hsa	<b>hsa_queue_t *</b>

Less library/compiler specific code – but also:

- Same device a OpenMP
- Dependency handling with asynchronous execution supported

→ next page



# Interop with ‘depend’ and ‘nowait’ | HIP + Fortran

```
#pragma omp target nowait depend(out: x[0:N]) map(from: x[0:N])
myVectorSet(N, 1.0, x);

#pragma omp task depend(out: y[0:N])
myVectorSet(N, -1.0, y);

#pragma omp interop init(targetsync: obj) \
depend(in: x[0:N]) \
depend(inout: y[0:N])
...
#pragma omp interop destroy(obj) nowait \
depend(out: y[0:N])

#pragma omp target depend(inout: x[0:N])
myDscal(N, scalar, x);

#pragma omp taskwait
```

```
integer(omp_interop_kind) :: obj
integer(omp_interop_fr_kind) :: fr
type(c_ptr) :: hip_sm

!$omp interop init(target, targetsync, prefer_type("hip")): obj

fr = omp_get_interop_int (obj, omp_ipr_fr_id, res)
hip_sm = omp_get_interop_ptr (obj, omp_ipr_targetsync, res)
...
!$omp interop destroy(obj)
```

Full example with explanations: OpenMP Examples  
document  
[www.openmp.org/specifications](http://www.openmp.org/specifications)



# Variant Functions

```
bool use_variant = false;

void variant_fn (int n, int *A) { /* ... */ }

#pragma omp declare variant(variant_fn) match( user={condition(use_variant)} )
// #pragma omp declare variant(variant_fn) match( construct={parallel} )
// #pragma omp declare variant(variant_fn) match( target_device={kind(gpu)} )

void base_fn (int n, int *A) { /* ... */ }

void test() {
...
base_fn (n, arr); // Calls base function
use_variant = true;
base_fn (n, arr); // Calls variant function 'variant_fn'
}
```



# Variant Function for Device Variant

Assume base function:

```
void daxpy (int n, double *da, double *dx, int incx, double *dx, int incy);
```

And device wrapper function – issues: device hard coded, device-ptr assumptions

```
void device_daxpy (int n, double *da, double *dx, int incx, double *dx, int incy) {  
    cudaStreamCreate (&stream)  
    cublasCreate (&handle)  
    cublasSetStream (handle, stream);  
    #pragma omp target data use_device_ptr(dx, dy)  
        cublasDaxpy (handle, n, da, dx, incx, dy, incy);  
    cublasDestroy (handle); cudaStreamDestroy (stream);  
}  
  
#pragma omp declare variant(device_daxpy) match( ... )  
void daxpy (int n, double *da, double *dx, int incx, double *dx, int incy);
```



# Variant Function for Device Variant

Autoconvert host→device pointer + create streaming object

```
void device_daxpy(int n, double *da, double *dx, int incx,
                  double *dx, int incy, omp_interop_t obj)
{
    cudaStream_t stream = omp_get_interop_ptr (obj, omp_ipr_targetsync, nullptr);
    cublasSetStream (handle, stream);
    cublasDaxpy (handle, n, da, dx, incx, dy, incy);
    cublasDestroy (handle);
}

#pragma omp declare variant(device_daxpy) adjust_args(need_device_ptr: dx, dy) \
    append_args(interp(targetsync)) match( construct={dispatch} )
void daxpy (int n, double *da, double *dx, int incx, double *dx, int incy);

void test() { /* ... */
    daxpy (N, &alpha, d_X, 1, d_Y, 1); // Base function
    #pragma omp dispatch device(my_dev) // Specify a device to use
        // Variant function called with device ptr for dx/dy + 'targetsync' interop object
    daxpy (N, &alpha, d_X, 1, d_Y, 1);
```



# Sneak preview: Low-Overhead OpenMP



## GPU kernel

Simple & automatically — vs. — explicit but faster

- CUDA + HIP (as language): Low-overhead + fast code
- OpenMP: Flexible, multiple design patterns, e.g. nested ‘parallel’  
no need to specify number of teams, threads,  
stateful (internal control variables, ICV),

For hot code, faster code would be better → OpenMP 6.x ... →



# OpenMP's Kernel-Language-Like Features

## CUDA/HIP

- Team-private variables: `__SHARED__`
  - for static vars
  - with size specified at kernel-launch time
- ```
dim3 nblocks(N/256, N/256, N/256);
dim3 nthreads(256, 256, 256);
kernel<<<nblocks, nthreads>>>(...)
```

  - and —  

```
int x = blockIdx.x * blockDim.x +
threadIdx.x;
```
- By construction: low launch overhead

## OpenMP

- `omp groupprivate(x)` [OpenMP 6.0]
- `omp target dyn_groupprivate(size)`  
+ `omp_get_dyn_groupprivate_ptr()`  
[TR14 → OpenMP 6.1]
- [WIP, 6.1?] `omp target teams`  
`num_teams(dims(3):N/256,N/256,N/256)`
  - and —  
`int x = omp_get_team_num_dim(1) *
omp_get_num_teams_dim(1) + ...`
- [6.1/7?] optionally, disable features to reduce launch/on-device overhead



# Covered Topics



**OpenMP®**

- GCC's OpenMP / offload support
- interop – easier access to vendor libs
- OpenMP CUDA/HIP-like features ('kernel language')

More to the last two → SC25 booth/tech talk → [link.openmp.org/sc25talks](https://link.openmp.org/sc25talks)

Thanks for listening! — Questions, comments?

Tobias Burnus <[tburnus@baylibre.com](mailto:tburnus@baylibre.com)>



CFP



Home



Matrix Space



<https://embedded-recipes.org/>

# Venture capital for opensource projects <https://commit.fund/>

We believe in open-source, so does our partner >commit



>commit is a VC fund investing in early-stage Commercial Open-Source Startups with European roots.

Companies backed by >commit & its team



Twenty



UMA



Mastrra



White Circle



Keep



Better-Auth



Graphcore



Pangolin



Pyannote



Pandas AI



Stealth project

Disclaimer: BayLibre SAS is a Limited Partner in the >commit fund

Learn more at [commit.fund](https://commit.fund/)