

Innovations with YAML/CABAC in H.264/AVC software decoding


Thibault Raffaillac, Ph.D.

Software engineer at Quantum Surgical (Montpellier, France)


edge264

H.264/AVC decoder for Progressive High & MVC 3D profiles


- BSD-3-Clause license
- x86/x64, **ARM32/64, WASM (in progress)**
- Linux, Windows, Mac
- **Prefetching**
- **CI/CD**
- **Custom bitstream encoder**
- **Stress testing (in progress)**
- **Netflix support**



H.264/AVC decoder structure



1. YAML logging output




1-1. Debugging

```
- nal_ref_idc: 1
  nal_unit_type: 7 # Sequence parameter set
  profile_idc: 66 # Baseline
  constraint_set_flags: [1,1,1,0,0,0]
  level_idc: 1.2
  chroma_format_idc: 1 # 4:2:0 # inferred
  bit_depth: {luma: 8, chroma: 8} # inferred
  log2_max_frame_num: 16
  pic_order_cnt_type: 0
  log2_max_pic_order_cnt_lsb: 16
  max_num_ref_frames: 1
  gaps_in_frame_num_value_allowed_flag: 0
  pic_size_in_mbs: {width: 11, height: 9}
  frame_mbs_only_flag: 1
  direct_8x8_inference_flag: 1
  max_num_reorder_frames: 16 # inferred
  max_dec_frame_buffering: 16 # inferred
  decode_NAL_result: 0
```

- **Readability** by slim syntax (vs. XML/JSON), coloring, comments (vs. JSON)
- **Compactness** by recursive format (vs. CSV), inline structs
- **Scriptability** by design (vs. TXT, HTML)

1-2. Reencoding



Use a high-level language!

Include context from past parameter sets → low overhead, easier to reencode


frame_num: 0



frame_num: {bits: 4, value: 0}

- Versus lib h264bitstream: **9400** loc with C → **500** loc with Python


1-2. Reencoding



Python **integers** = infinite big-endian **bit buffers** (with leading set bit)

- Initialize empty bit buffer `buf = 1`
- Insert n-bit value `buf = buf<<n | value`
- Get buffer size `nbits = buf.bit_length() - 1`
- Remove leading bit & pad to bytes `buf = (buf ^ 1<<nbits) << (-nbits%8)`
- Write to file `f.write(buf.to_bytes((nbits+7)//8, byteorder="big"))`

1-3. Custom test streams




YAML encoder → create bitstreams for **stress testing**


- Catch up after ffmpeg/openh264 on 10+ years of field testing
- Search for *shall* clauses in spec! (e.g. «A sequence parameter set *RBSP*, with that particular value of *seq_parameter_set_id*, shall be available to the decoding process prior to its activation. »)
- Can be included in public repository & CI/CD since I own them
- Much smaller than conformance bitstreams → much faster to run test suite

1-4. Data analysis



Current thread scheduling



Frame dependencies



Decoding time per size




2. CABAC decoding


Screenshot of Firefox Profiler (collected with samply)

Total (samples)		Self	
41 %	2 387	2 387	► get_ae libedge264.1.0.0.dylib
29 %	1 689	1 689	► decode_inter libedge264.1.0.0.dylib
8,3 %	484	484	► deblock_mb libedge264.1.0.0.dylib
7,3 %	429	429	► parse_residual_coeffs_cabac libedge264.1.0.0.dylib
3,0 %	176	176	► add_idct4x4 libedge264.1.0.0.dylib
2,8 %	166	166	► parse_slice_data_cabac libedge264.1.0.0.dylib
2,3 %	134	134	► parse_NxN_residual_cabac libedge264.1.0.0.dylib
2,0 %	116	116	► decode_direct_mv_pred libedge264.1.0.0.dylib
1,4 %	80	80	► parse_mvd_pair_cabac libedge264.1.0.0.dylib
0,9 %	55	55	► add_idct8x8 libedge264.1.0.0.dylib
0,8 %	45	45	► parse_ref_idx_cabac libedge264.1.0.0.dylib
0,6 %	36	36	► parse_chroma_residual_cabac libedge264.1.0.0.dylib
0,4 %	22	22	► parse_slice_layer_without_partitioning libedge264.1.0.0.dylib
0,3 %	18	18	► parse_inter_residual_cabac libedge264.1.0.0.dylib
0,1 %	8	8	► madvise libsystem_kernel.dylib
0,1 %	6	6	► transform_dc2x2 libedge264.1.0.0.dylib
0,1 %	5	5	► decode_intra4x4 libedge264.1.0.0.dylib
0,0 %	2	2	► decode_intra8x8 libedge264.1.0.0.dylib
0,0 %	2	2	► pthread_cond_broadcast libsystem_pthread.dylib
0,0 %	1	1	► worker_loop libedge264.1.0.0.dylib
0,0 %	1	1	► decode_intraChroma libedge264.1.0.0.dylib
0,0 %	1	1	► _platform_memmove\$VARIANT\$Haswell libsystem_platform.dylib
0,0 %	1	1	► _kernelrpc_mach_vm_deallocate_trap libsystem_kernel.dylib


2. CABAC decoding



2. CABAC decoding



2. CABAC decoding



2-1. Extending CABAC state to size_t

Initially **load 64 bits in offset**, and **shift range up 55 bits** ($64 - 9$)


At each renormalization, offset & range are \leq 8-bit, refill 56 bits (7 bytes) into offset & shift range up the same

- Trades less frequent renormalizations with new counting of extra bits
- Wide loads work well with on-the-fly unescaping (c.f. last year)


range $\in [256;511] \rightarrow$ range's 9th bit is set \rightarrow **bit count = 64 - clz(range)**

- Allows keeping bit count without an extra variable or extra set bit (ffmpeg)

2-2. Batch-decoding CABAC bypass



2-2. Batch-decoding CABAC bypass



2-2. Batch-decoding CABAC bypass

To get N bypass bits (*with size_t format*):

- Ensure offset & range have $\geq N$ extra bits (renorm otherwise)
- Shift range down N bits
- **offset / range** \rightarrow N bypass bits
- $\text{offset \% range} \rightarrow$ new offset

To return M unconsumed bits:

- $\text{offset + range * unconsumed} \rightarrow$ new offset
- Shift range up M bits

Thank you for your attention!

<https://github.com/tvlabs/edge264>

traf@ik.me

- 1. YAML logging output
 - 1. Decoding
 - 2. Reencoding (*Python 🤖*)
 - 3. Custom test streams (*shall 🔎*)
 - 4. Data analysis
- 2. CABAC decoding
 - 1. Extending state to `size_t`
 - 2. Batch-decoding bypass (`div %`)

1. x86-64 microarchitecture runtime variants

Compile entire lib for x86-64-v1, compile lib minus top-level functions for v2 and v3, then branch at runtime

- Easy `make` selection of optional variants
- Allows SIMD everywhere (except top-level functions)
- Used for logging support too

