

Performance and reliability pitfalls of eBPF

Usama Saqib @ FOSDEM/ebpf

January 2026



DATADOG

About Me


- SWE @ Datadog
- Working with eBPF as part of the ebpf-platform team
- Enjoy exploring and working with kernel internals

Contents

- Kretprobe scaling pains
- A deadlock in fentry
- A production incident due to uprobe

Kretprobe scaling pains - Spinlocks


- Customer impact
- Running amazon linux 2 with old 4.14 kernel
- Large instance with 96 cores
- Large impact to network throughput



Distro	Kernel	CNM?	Throughput
AL2	4.14	✗	✓ 4.238M
AL2	4.14	✓	✗ 119.1K


Kretprobe scaling pains - Internals (kernel 4.14)

- The trampoline saves the return address for each kretprobe invocation in a unique struct kretprobe_instance
- The list has max_active slots configured during kprobe registration
- Each invocation requires it's own instance.



Kretprobe scaling pains - Bottleneck (kernel 4.14)

- The free list of instances is guarded by a spinlock
- Parallel executions of the same kretprobe bottleneck here.



- Perf top reflects this observation.
- CPU time spent in the pre-handler for high contention case
- With massive drop in throughput

```
--38.63%--inet_recvmsg
    tcp_recvmsg
    __key.28673
    ftrace_ops_assist_func
    kprobe_ftrace_handler
    aggr_pre_handler
    |
    --38.63%--pre_handler_kretprobe
    |
    --38.60%--_raw_spin_lock_irqsave
    |
    --38.60%--queued_spin_lock_slowpath
                native_queued_spin_lock_slowpath
```

Kretprobe scaling pains - Cache contention (kernel 6.6)

- Customer impact
- Running newer kernel version 6.6
- Very large instance type

Distro	Kernel	CNM?	Throughput
COS	6.6	✗	✓ 8.5M
COS	6.6	✓	✗ 85K

Kretprobe scaling pains - New mechanism (kernel 6.6)

- Kernel patch removes the spinlock bottleneck in favor of a lockless algorithm.

```
tree f4e6ad94a04db211bfacc7cc6a931a96ecd897fd
parent e563604a5f5a891283b6a8db4001cee833a7c6b8
author Peter Zijlstra <peterz@infradead.org> Sat Aug 29 22:03:56 2020 +0900
committer Ingo Molnar <mingo@kernel.org> Mon Oct 12 18:27:28 2020 +0200

kprobes: Replace rp->free_instance with freelist

Gets rid of rp->lock, and as a result kretprobes are now fully
lockless.

Signed-off-by: Peter Zijlstra (Intel) <peterz@infradead.org>
Signed-off-by: Masami Hiramatsu <mhiramat@kernel.org>
Signed-off-by: Ingo Molnar <mingo@kernel.org>
Link: https://lore.kernel.org/r/159870623583.1229682.17472357584134058687.stgit@devnote2
```

- Hint of the issue can be found in the kernel comments.

```
/*
 * Copyright: cameron@moodycamel.com
 *
 * A simple CAS-based lock-free free list. Not the fastest thing in the world
 * under heavy contention, but simple and correct (assuming nodes are never
 * freed until after the free list is destroyed), and fairly speedy under low
 * contention.
 *
 * Adapted from: https://moodycamel.com/blog/2014/solving-the-aba-problem-for-lock-free-lists
 */
```

Kretprobe scaling pains - Bottleneck (kernel 6.6)

- New lock less design
 - No lock so no lock contention!
- On a big enough machine memory becomes the bottleneck.
 - The freelist manipulation causes cache bouncing
 - The problem gets worse with more CPUs
- Shown by perf stat where we see we are bottlenecked on memory

Kretprobe scaling pains - Objpool (kernel 6.6+)

```
8 tree db31d9a59c0ac9b21a7aa243bb8a2d02f0525c4f
7 parent ef8a257b4e499a979364b1f9caf25a325f6ee8b8
6 author Masami Hiramatsu <mhiramat@kernel.org> Sat Mar 26 11:27:05 2022 +0900
5 committer Alexei Starovoitov <ast@kernel.org> Mon Mar 28 19:38:09 2022 -0700
4
3 kprobes: Use rethook for kretprobe if possible
2
1 Use rethook for kretprobe function return hooking if the arch sets
  CONFIG_HAVE_RETHOOK=y. In this case, CONFIG_KRETPROBE_ON_RETHOOK is
1 set to 'y' automatically, and the kretprobe internal data fields
2 switches to use rethook. If not, it continues to use kretprobe
3 specific function return hooks.
4
5 Suggested-by: Peter Zijlstra <peterz@infradead.org>
6 Signed-off-by: Masami Hiramatsu <mhiramat@kernel.org>
7 Signed-off-by: Alexei Starovoitov <ast@kernel.org>
8 Link: https://lore.kernel.org/bpf/164826162556.2455864.12255833167233452047.stgit@devnote2
9
```

```
/*
 * objpool: ring-array based lockless MPMC queue
 *
 * Copyright: wuqiang.matt@bytedance.com,mhiramat@kernel.org
 *
 * objpool is a scalable implementation of high performance queue for
 * object allocation and reclamation, such as kretprobe instances.
 *
 * With leveraging percpu ring-array to mitigate hot spots of memory
 * contention, it delivers near-linear scalability for high parallel
 * scenarios. The objpool is best suited for the following cases:
 * 1) Memory allocation or reclamation are prohibited or too expensive
 * 2) Consumers are of different priorities, such as irqs and threads
 *
```

An Fentry and RCU bug

- Lets move away from kretprobes and into the modern world of fentry
- Reports of failures when exiting the agent container

```
root@ddvm:~# docker container stop dd-agent-foo
Error response from daemon: cannot stop container: dd-agent-foo: tried to kill container, but did not receive an exit event
from Docker."}
```

An Fentry and RCU bug

```
echo w > /proc/sysrq-trigger
```

- Two tasks in
uninterruptible sleep
1. System-probe
(agent)
 2. rcu_tasks_kthread

```
[ 183.865554] task:system-probe state:D stack:0
[ 183.865559] Call Trace:
[ 183.865562] <TASK>
[ 183.865567] __schedule+0x2cb/0x760
[ 183.865578] schedule+0x63/0x110
[ 183.865583] schedule_timeout+0x157/0x170
[ 183.865593] wait_for_completion+0x88/0x150
[ 183.865603] __wait_rcu_gp+0x17e/0x190
[ 183.865611] synchronize_rcu_tasks_generic+0x64/0xe0
[ 183.865617] ? __pfx_call_rcu_tasks+0x10/0x10
[ 183.865623] ? __pfx_wakeme_after_rcu+0x10/0x10
[ 183.865632] synchronize_rcu_tasks+0x15/0x20
[ 183.865638] ftrace_shutdown.part.0+0xd1/0x1f0
[ 183.865651] ? 0xfffffffffc11f0000
[ 183.865686] unregister_ftrace_function+0x47/0x170
[ 183.865694] ? 0xfffffffffc11f0000
[ 183.865702] unregister_ftrace_direct+0x60/0x1e0
[ 183.865710] ? bpf_trampoline_6442514787+0xba/0x1000
[ 183.865723] ? 0xfffffffffc11f0000
[ 183.865729] ? __pfx__ia32_sys_setregid16+0x10/0x10
[ 183.865736] bpf_trampoline_update+0x505/0x5f0
[ 183.865745] ? __radix_tree_delete+0x9e/0x150
[ 183.865757] bpf_trampoline_unlink_prog+0x9d/0x130
[ 183.865765] bpf_tracing_link_release+0x16/0x50
[ 183.865772] bpf_link_free+0x52/0x80
[ 183.865778] bpf_link_release+0x26/0x40
[ 183.865782] __fput+0xf9/0x2c0
[ 183.865790] __fput+0xe/0x20
[ 183.865794] task_work_run+0x5e/0xa0
[ 183.865799] exit_to_user_mode_loop+0x105/0x130
[ 183.865805] exit_to_user_mode_prepare+0xa5/0xb0
[ 183.865808] syscall_exit_to_user_mode+0x29/0x60
[ 183.865814] do_syscall_64+0x62/0x90
[ 183.865821] ? __rseq_handle_notify_resume+0x37/0x70
[ 183.865828] ? exit_to_user_mode_loop+0xe5/0x130
[ 183.865832] ? exit_to_user_mode_prepare+0x30/0xb0
[ 183.865836] ? syscall_exit_to_user_mode+0x37/0x60
[ 183.865840] ? do_syscall_64+0x62/0x90
[ 183.865843] ? do_syscall_64+0x62/0x90
[ 183.865846] ? do_syscall_64+0x62/0x90
[ 183.865850] entry_SYSCALL_64_after_hwframe+0x73/0xdd
[ 183.865851] RIP: 0033:0x1ce792e
```

```
[ 183.865273] task:rcu_tasks_kthre state:D stack:0
[ 183.865284] Call Trace:
[ 183.865288] <TASK>
[ 183.865322] __schedule+0x2cb/0x760
[ 183.865340] schedule+0x63/0x110
[ 183.865345] schedule_timeout+0x157/0x170
[ 183.865357] wait_for_completion+0x88/0x150
[ 183.865367] __synchronize_srcu+0x91/0xd0
[ 183.865374] ? __pfx_wakeme_after_rcu+0x10/0x10
[ 183.865383] ? ktime_get_mono_fast_ns+0x3c/0xa0
[ 183.865392] synchronize_srcu+0xb9/0x110
[ 183.865399] rcu_tasks_postscan+0x39/0x60
[ 183.865405] rcu_tasks_wait_gp+0x116/0x310
[ 183.865411] ? __raw_spin_unlock_irqrestore+0x11/0x60
[ 183.865418] ? rcu_tasks_need_gpcb+0x199/0x2f0
[ 183.865427] ? __pfx_rcu_tasks_kthread+0x10/0x10
[ 183.865433] rcu_tasks_one_gp+0x108/0x140
[ 183.865440] rcu_tasks_kthread+0x33/0x40
[ 183.865447] kthread+0xef/0x120
[ 183.865456] ? __pfx_kthread+0x10/0x10
[ 183.865464] ret_from_fork+0x44/0x70
[ 183.865472] ? __pfx_kthread+0x10/0x10
[ 183.865479] ret_from_fork_asm+0x1b/0x30
[ 183.865495] </TASK>
```

An Fentry and RCU bug - echo t > /proc/sysrq-trigger

s6-svscan in interruptible sleep.


S6 is used by docker for process management.

Parent of system-probe

Pid 1 in the pid namespace of the container.

```
[ 322.722345] task:svscan      state:S stack:0 pi
[ 322.722346] task:s6-svscan    state:S stack:0 pi
[ 322.722347] Call Trace:
[ 322.722348] <TASK>
[ 322.722349] __schedule+0x2cb/0x760
[ 322.722352] schedule+0x63/0x110
[ 322.722353] do_wait+0x173/0x320
[ 322.722356] kernel_wait4+0xbd/0x170
[ 322.722359] ? __pxf_child_wait_callback+0x10/0x10
[ 322.722362] zap_pid_ns_processes+0x115/0x1b0
[ 322.722365] forget_original_parent+0x356/0x370
[ 322.722367] ? cgroup_exit+0xf3/0x1b0
[ 322.722370] do_exit+0x5a9/0x6f0
[ 322.722371] ? do_wait+0x1b1/0x320
[ 322.722374] do_group_exit+0x35/0x90
[ 322.722376] get_signal+0x8d8/0x940
[ 322.722377] ? __pxf_child_wait_callback+0x10/0x10
[ 322.722381] arch_do_signal_or_restart+0x39/0x120
[ 322.722385] exit_to_user_mode_loop+0x9a/0x130
[ 322.722386] exit_to_user_mode_prepare+0xa5/0xb0
[ 322.722388] syscall_exit_to_user_mode+0x29/0x60
[ 322.722390] do_syscall_64+0x62/0x90
[ 322.722391] ? exit_to_user_mode_prepare+0x9b/0xb0
[ 322.722392] ? syscall_exit_to_user_mode+0x37/0x60
[ 322.722394] ? do_syscall_64+0x62/0x90
[ 322.722398] ? switch_fpu_return+0x55/0xf0
[ 322.722401] ? exit_to_user_mode_prepare+0x9b/0xb0
[ 322.722403] ? syscall_exit_to_user_mode+0x37/0x60
[ 322.722404] ? do_syscall_64+0x62/0x90
[ 322.722406] entry_SYSCALL_64_after_hwframe+0x73/0xdd
```

An Fentry and RCU bug - Fentry Internals




An Fentry and RCU bug - RCU-tasks

The RCU tasks subsystem is an RCU flavor designed to figure out when no **process** is holding a reference to a structure.

- Normal RCU allows to figure out when no CPU is holding a reference to a structure.
 - Only one process on a CPU can hold a reference
 - References can only be held in an “atomic context”
 - Meaning of “atomic context” is dependent on preemption modes
 - In all cases except PREEMPT_RT explicit blocking is always illegal
 - With BPF we are dealing with cases where blocking+preemption is not allowed.
 - Other preemption modes is a whole other can of worms as far as BPF is concerned.
- In contrast RCU-tasks
 - Multiple processes on a CPU can hold a reference
 - The quiescent state here is that all tasks have performed a voluntary context switch

An Fentry and RCU bug - RCU-tasks



An Fentry and RCU bug - RCU-tasks

1. Scan the tasklist for all tasks on the runqueue, and add them to the holdout list. The holdout list is all tasks which have not yet performed a voluntary context switch.
2. Wait for all tasks which are in late stage do_exit to complete. Tasks in late stage of do_exit are not on the tasklist, so therefore, are not visible to the scan done above. The wait is performed by calling synchronize_srcu(&tasks_rcu_exit_srcu). This blocks the current task until all tasks_rcu_exit_srcu read-side critical sections have completed.
3. Loop over the holdout list until all tasks in the list have been seen to perform a voluntary context switch. This is tracked by placing a hook in the scheduler.

An Fentry and RCU bug - RCU-tasks

```
tree 5b6d2c42aaaf8b20397bd09c0ac31738618f57046
parent 53c6d4edf874d3cbc031a53738c6cba9277faea5
author Paul E. McKenney <paulmck@linux.vnet.ibm.com> Mon Aug 4 06:10:23 2014 -0700
committer Paul E. McKenney <paulmck@linux.vnet.ibm.com> Sun Sep 7 16:27:22 2014 -0700
```


rcu: Make TASKS_RCU handle tasks that are almost done exiting

Once a task has passed `exit_notify()` in the `do_exit()` code path, it is no longer on the task lists, and is therefore no longer visible to `rcu_tasks_kthread()`. This means that an almost-exited task might be preempted while within a trampoline, and this task won't be waited on by `rcu_tasks_kthread()`. This commit fixes this bug by adding an `srcu_struct`. An exiting task does `srcu_read_lock()` just before calling `exit_notify()`, and does the corresponding `srcu_read_unlock()` after doing the final `preempt_disable()`. This means that `rcu_tasks_kthread()` can do `synchronize_srcu()` to wait for all mostly-exited tasks to reach their final `preempt_disable()` region, and then use `synchronize_sched()` to wait for those tasks to finish exiting.


An Fentry and RCU bug - Zapper

- Sends kill signal to all children
- Waits until they are dead
- Does this inside a tasks_rcu_exit_srcu read-side critical section !!!!


An Fentry and RCU bug - Full picture



An Fentry and RCU bug - Full picture



An Fentry and RCU bug - Full picture



An Fentry and RCU bug - Conclusion

- At that time we were blocked from using fentry since the fix for this was not backported

```
tree 3e97b5789d9c52bf95808d07be34c86d2eb27020
parent 46faf9d8e1d52e4a91c382c6c72da6bd8e68297b
author Paul E. McKenney <paulmck@kernel.org> Fri Feb 2 11:28:45 2024 -0800
committer Boqun Feng <boqun.feng@gmail.com> Sun Feb 25 14:21:43 2024 -0800

rcu-tasks: Maintain lists to eliminate RCU-tasks/do_exit() deadlocks

This commit continues the elimination of deadlocks involving do_exit()
and RCU tasks by causing exit_tasks_rcu_start() to add the current
task to a per-CPU list and causing exit_tasks_rcu_stop() to remove the
current task from whatever list it is on. These lists will be used to
track tasks that are exiting, while still accounting for any RCU-tasks
quiescent states that these tasks pass through.

[ paulmck: Apply Frederic Weisbecker feedback. ]

Link: https://lore.kernel.org/all/20240118021842.290665-1-chenzhongjin@huawei.com/
```

Production impact from uprobe

- Service unavailability on a few nodes in production
- Khungtaskd reporting a large number of tasks hung for more than 120 seconds!
 - Some cases only tasks associated with the production application
 - Some cases tasks from bash, systemd, runc, and many more
- Among these were tasks
 - Registering uprobes
 - Firing uprobes

Uprobe Registration	Uprobe Firing
<code>uprobe_register -> register_for_each_vma</code>	<code>handle_swbp -> find_active_uprobe</code>

Production impact from uprobe - Investigation

Is it a deadlock again?

- We found an *almost* circular lock dependency but not complete.
- So cannot conclude deadlock.

Memory corruption?

- We found a patch missing in our kernel which discussed a possible memory corruption effect struct `mm_struct`.
- We built a custom kernel to deterministically trigger the condition for that corruption.
- Did not result in the symptoms we were seeing.

Livelock due to memory pressure?

- Hypothesis that the kernel was looping trying to allocate some memory but failing due to memory pressure.
- Evaluated all page allocation paths in the uprobe register and uprobe firing path and this did not seem possible.

Production impact from uprobe - Investigation

Monitoring kernel locks

- Began monitoring kernel locks to establish deadlock and failed.


trace_contention_begin

trace_contention_end

Production impact from uprobe - Uprobe Registration

Input to uprobe registration

1. Inode of the file to be probed
2. Address to be probed




Production impact from uprobe - Uprobe Registration


```
__uprobe_register
    down_write(&uprobe->register-rwsem) // write lock on semaphore
    register_for_each_vma
        percpu_down_write(&dup_mmap_sem) // write lock on semaphore
        for each interesting MM:
            mmap_write_lock(mm) // write lock on mmap_lock semaphore
            do_work();
            mm_write_unlock(mm) // unlock mmap_lock semaphore
```

Mmap lock is required to synchronize with uprobe firing, incase the uprobe is being removed.

Production impact from uprobe - Uprobe Registration



Production impact from uprobe - Uprobe Firing



Production impact from uprobe - Uprobe Firing


- Larger number of tasks blocked on `find_active_uprobe`
- Blocked on acquire read lock on `mmap_lock`
- Lots of application threads piling up waiting on this lock

Production impact from uprobe - mmap_lock & rw_semaphore

- Allows concurrent readers but exclusive writers
- Use contention tracepoints to build model
 - If reader is blocked, why?
 - Can we get the owner?
 - If write is blocked, why?
 - Reader owned or writer owned?
 - Can we get the owner?
 - Depth of wait queue
 - Contention duration
 - When is forward progress made

Production impact from uprobe - Root Cause


- Deep queue of mmap read locks interspersed with writes
 - Write from mmap, madvise, uprobe register



- `Rw_semaphore` wakes up only max 250 reader threads
 - To avoid thundering herd
 - But this means even batches of reads become serialized
 - No concurrent reader if writer at head of wait queue
 - Avoid writer starvation

Production impact from uprobe - Root Cause

- If a uprobe_register write_lock is blocked then whole system bricks



Production impact from uprobe - Root Cause

- If a uprobe firing `read_lock` is blocked then applicate suffer
 - Due to piling threads stuck in the kernel

Production impact from uprobe - Kernel Fix

Speculative mmap_lock acquisition

```
parent 83e3dc9a5d4d7402adb24090a77327245d593129
author Andrii Nakryiko <andrii@kernel.org> Thu Nov 21 19:59:22 2024 -0800
committer Peter Zijlstra <peterz@infradead.org> Mon Dec 2 12:01:38 2024 +0100
```

uprobes: add speculative lockless VMA-to-inode-to-uprobe resolution

Given filp_cachep is marked SLAB_TYPESAFE_BY_RCU (and FMODE_BACKING files, a special case, now goes through RCU-delated freeing), we can safely access vma->vm_file->f_inode field locklessly under just rCU_read_lock() protection, which enables looking up uprobe from uprobes_tree completely locklessly and speculatively without the need to acquire mmap_lock for reads. In most cases, anyway, assuming that there are no parallel mm and/or VMA modifications. The underlying struct file's memory won't go away from under us (even if struct file can be reused in the meantime).

We rely on newly added mmap_lock_speculate_{try_begin, retry}() helpers to validate that mm_struct stays intact for entire duration of this speculation. If not, we fall back to mmap_lock-protected lookup. The speculative logic is written in such a way that it will safely handle any garbage values that might be read from vma or file structs.

Thank you

Questions?



DATADOG