

String kfuncs

Simplifying string handling in eBPF programs

Viktor Malík

Principal Software Engineer @ Red Hat

January 31, 2026

Motivation: Why string handling?

- Many hooks (functions, LSMs, skb data) allow to access strings from eBPF programs
- Examples where string handling is useful:
 - Matching **filesystem paths** (e.g., is the file inside /home/\$USER/ . . .?)
 - Matching **process names** (e.g., does comm contain systemd?)
 - Searching **environment variables**
 - Parsing **protocol** (e.g. HTTP) headers

Motivation: String processing, the old way

- Copy string from kernel memory using `bpf_probe_read()`
- Manually implement all string operations

Motivation: String processing, the old way

- Copy string from kernel memory using `bpf_probe_read_str()`
- Manually implement all string operations
- Example:

```
SEC("lsm/bprm_check_security")
int BPF_PROG(path_check, struct linux_binprm *bprm) {
    char path[32];
    int path_len, last_slash = -1;

    path_len = bpf_probe_read_str(path, sizeof(path), bprm->filename);
    if (path_len < 0 || path_len >= 32) return 0;

    for (int i = 0; i < path_len; i++) {
        if (path[i] == '/')
            last_slash = i;
    }
    [...]
```

Motivation: String processing, the old way

- Copy string from kernel memory using `bpf_probe_read_str()`
- Manually implement all string operations
- Example:

```
SEC("lsm/bprm_check_security")
int BPF_PROG(path_check, struct linux_binprm *bprm) {
    char path[32];
    int path_len, last_slash = -1;

    path_len = bpf_probe_read_str(path, sizeof(path), bprm->filename);
    if (path_len < 0 || path_len >= 32) return 0;

    for (int i = 0; i < path_len; i++) {
        if (path[i] == '/')
            last_slash = i;
    }
    [...]
```

Motivation: String processing, the old way

- Copy string from kernel memory using `bpf_probe_read_str()`
- Manually implement all string operations
- Example:

```
SEC("lsm/bprm_check_security")
int BPF_PROG(path_check, struct linux_binprm *bprm) {
    char path[32];
    int path_len, last_slash = -1;

    path_len = bpf_probe_read_str(path, sizeof(path), bprm->filename);
    if (path_len < 0 || path_len >= 32) return 0;

    for (int i = 0; i < path_len; i++) {
        if (path[i] == '/')
            last_slash = i;
    }
    [...]
```

String kfuncs

- BPF Kernel Functions (kfuncs) are the modern way of exposing kernel functionality to eBPF programs
- String kfuncs provide implementations of the most common string operations

String kfuncs

- BPF Kernel Functions (kfuncs) are the modern way of exposing kernel functionality to eBPF programs
- String kfuncs provide implementations of the most common string operations
- Main advantages:
 - No need to reimplement the functions manually
 - No need to copy the strings onto the stack (which is limited) or to maps (which bring overhead)
 - Small performance benefit

Naive solution

- First idea: just call in-kernel implementations:

```
__bpf_kfunc int bpf_strlen(const char *s)
{
    return strlen(s);
}

size_t strlen(const char *s)
{
    const char *sc;
    for (sc = s; *sc != '\0'; ++sc)
        /* nothing */;
    return sc - s;
}
```

Naive solution

- First idea: just call in-kernel implementations:

```
__bp�_kfunc int bpf_strlen(const char *s)
{
    return strlen(s);
}
```

```
size_t strlen(const char *s)
{
    const char *sc;
    for (sc = s; *sc != '\0'; ++sc)
        /* nothing */;
    return sc - s;
}
```

- Not possible since we need to ensure:

- Safety

- Verifier is only able to check the first byte of the strings
 - Cannot use naked dereference as the memory doesn't have to be paged in

Naive solution

- First idea: just call in-kernel implementations:

```
__bpf_kfunc int bpf_strlen(const char *s)
{
    return strlen(s);
}

size_t strlen(const char *s)
{
    const char *sc;
    for (sc = s; *sc != '\0'; ++sc)
        /* nothing */;
    return sc - s;
}
```

- Not possible since we need to ensure:

- Safety

- Verifier is only able to check the first byte of the strings
 - Cannot use naked dereference as the memory doesn't have to be paged in

- Termination

- Strings are not necessarily null-terminated so the common algorithms can loop forever

Better solution

- Open code the functions:

```
__bpf_kfunc int bpf_strlen(const char *s__ign)
{
    guard(pagefault)();
    for (i = 0; i < XATTR_SIZE_MAX; i++) {
        __get_kernel_nofault(&c, s__ign, char, err_out);
        if (c == '\0') { return i; }
        s__ign++;
    }
    return i == XATTR_SIZE_MAX ? -E2BIG : i;
err_out:
    return -EFAULT;
}
```

Better solution

- Open code the functions:

```
__bpf_kfunc int bpf_strlen(const char *s__ign)
{
    guard(pagefault) ();
    for (i = 0; i < XATTR_SIZE_MAX; i++) {
        __get_kernel_nofault(&c, s__ign, char, err_out);
        if (c == '\0') { return i; }
        s__ign++;
    }
    return i == XATTR_SIZE_MAX ? -E2BIG : i;
err_out:
    return -EFAULT;
}
```

- Features:
 - page faults disabled

Better solution

- Open code the functions:

```
__bpf_kfunc int bpf_strlen(const char *s__ign)
{
    guard(pagefault) ();
    for (i = 0; i < XATTR_SIZE_MAX; i++) {
        __get_kernel_nofault(&c, s__ign, char, err_out);
        if (c == '\0') { return i; }
        s__ign++;
    }
    return i == XATTR_SIZE_MAX ? -E2BIG : i;
err_out:
    return -EFAULT;
}
```

- Features:

- page faults disabled
- __get_kernel_nofault instead of plain dereference

Better solution

- Open code the functions:

```
__bpf_kfunc int bpf_strlen(const char *s__ign)
{
    guard(pagefault) ();
    for (i = 0; i < XATTR_SIZE_MAX; i++) {
        __get_kernel_nofault(&c, s__ign, char, err_out);
        if (c == '\0') { return i; }
        s__ign++;
    }
    return i == XATTR_SIZE_MAX ? -E2BIG : i;
err_out:
    return -EFAULT;
}
```

- Features:

- page faults disabled
- __get_kernel_nofault instead of plain dereference
- upper limit on the string length

Verifier support

- String kfunc arguments:
 - Null-terminated (not necessarily) strings located in eBPF or kernel memory → **unsafe pointers**
 - Size arguments for bounded functions → **general integers**

Verifier support

- String kfunc arguments:
 - Null-terminated (not necessarily) strings located in eBPF or kernel memory → **unsafe pointers**
 - Size arguments for bounded functions → **general integers**
- Verifier supports "special suffices" for kfunc arguments
 - `__str` for string literals

Verifier support

- String kfunc arguments:
 - Null-terminated (not necessarily) strings located in eBPF or kernel memory → **unsafe pointers**
 - Size arguments for bounded functions → **general integers**
- Verifier supports "special suffices" for kfunc arguments
 - ~~__str for string literals~~
 - ~~__sz for buffer size literals~~

Verifier support

- String kfunc arguments:
 - Null-terminated (not necessarily) strings located in eBPF or kernel memory → **unsafe pointers**
 - Size arguments for bounded functions → **general integers**
- Verifier supports "special suffices" for kfunc arguments
 - ~~__str for string literals~~
 - ~~__sz for buffer size literals~~
 - **__ign** for ignoring argument verification

Verifier support

- String kfunc arguments:
 - Null-terminated (not necessarily) strings located in eBPF or kernel memory → **unsafe pointers**
 - Size arguments for bounded functions → **general integers**
- Verifier supports "special suffices" for kfunc arguments
 - ~~__str for string literals~~
 - ~~__sz for buffer size literals~~
 - **__ign for ignoring argument verification**
- We can afford ignoring verification since safety is ensured dynamically

API details

- Read-only functions for now
- Initially added functions:
 - `bpf_strcmp`
 - `bpf_strlen`, `bpf_strnlen`
 - `bpf_strstr`, `bpf_strnstr`
 - `bpf_strchr`, `bpf_strnchr`, `bpf_strchrnul`, `bpf_strrchr`
 - `bpf_strspn`, `bpf_strcspn`
- Several new functions added after the initial patchset:
 - `bpf_strcasecmp`, `bpf_strncasecmp`
 - `bpf_strcasestr`, `bpf_strncasestr`

API details

Comparison with stdlib functions

- Return **indices** instead of pointers

```
__bpf_kfunc int bpf_strchr(const char *s__ign, char c)
```

- Input pointers are unsafe → output pointers would also be unsafe
- Memory may not be there on successive reads → would still need explicit bounds check

API details

Comparison with stdlib functions

- Return **indices** instead of pointers

```
__bpf_kfunc int bpf_strchr(const char *s__ign, char c)
```

- Input pointers are unsafe → output pointers would also be unsafe
- Memory may not be there on successive reads → would still need explicit bounds check

- Return negative **error codes** on errors

- -EFAULT when reading paged out memory
- -E2BIG when the strings are too big
- -ERANGE when the strings are outside of the kernel address space
- -ENOENT as “item not found”

Practical usage

Basic example

Manual implementation

```
SEC("lsm/bprm_check_security")
int BPF_PROG(path_check, struct linux_binprm *bprm) {
    char path[32];
    int path_len, last_slash = -1;

    path_len = bpf_probe_read_str(path, sizeof(path),
                                  bprm->filename);
    if (path_len < 0 || path_len >= 32)
        return 0;

    for (int i = 0; i < path_len; i++) {
        if (path[i] == '/')
            last_slash = i;
    }

    // Basename is at path + last_slash + 1
}
```

Using string kfuncs

```
SEC("lsm/bprm_check_security")
int BPF_PROG(path_check, struct linux_binprm *bprm) {
    int last_slash;

    last_slash = bpf strrchr(bprm->filename, '/');
    if (last_slash < 0)
        return 0;

    // Basename is at bprm->filename + last_slash + 1
}
```

Practical usage

bpftace extensions

- Thanks to string kfuncs, bpftace now has more built-in functions for string manipulation:
 - bool **strcontains**(string haystack, string needle)
 - int64 **strstr**(string haystack, string needle)
 - uint64 **strlen**(string exp)
 - And many more are coming!
- This allows complex in-kernel filtering and string processing directly in bpftace scripts.

Future directions

- Adding **more read-only functions** – already happening
- Adding functions which **manipulate memory** (e.g. `bpf_strncpy`)
 - Problem: source and destination cannot overlap (could be enforced by the verifier?)
 - Is there a use-case for such functions?
- From **bpftrace** perspective:
 - Avoid `bpf_probe_read_str` when not necessary
 - Add wrappers for more functions

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



linkedin.com/company/red-hat



youtube.com/user/RedHatVideos



facebook.com/redhatinc



twitter.com/RedHat