

The GIL and API Performance

The Past, Present, and Free-Threaded Future

whoami

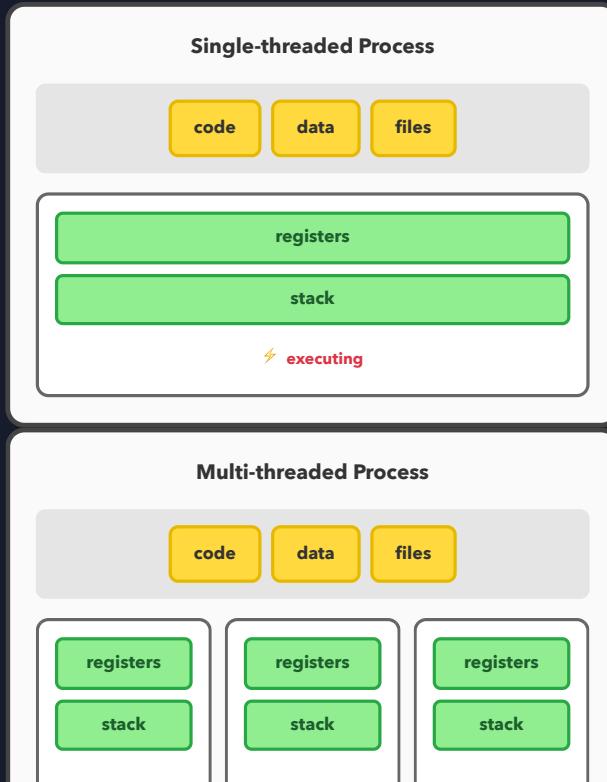
Ruben Hias

Software Engineer at TechWolf

floww.dev: Workflow automation for developers

Refresher: Threads vs Processes

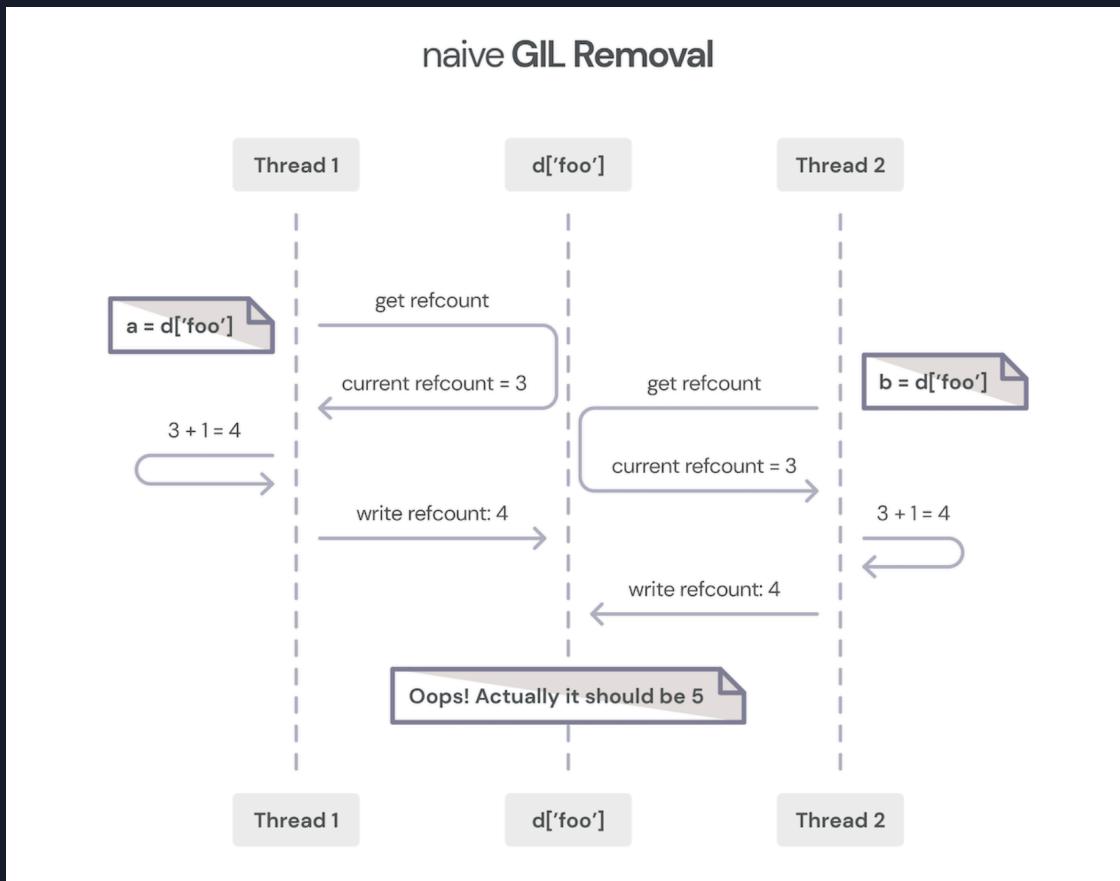
- **Process:** Independent program with its own memory space
 - Isolated memory → safer but higher overhead
 - Communication via IPC
- **Thread:** Lightweight execution unit within a process
 - Shared memory space → efficient but requires synchronization
 - Direct communication via shared data



GIL: Global Interpreter Lock

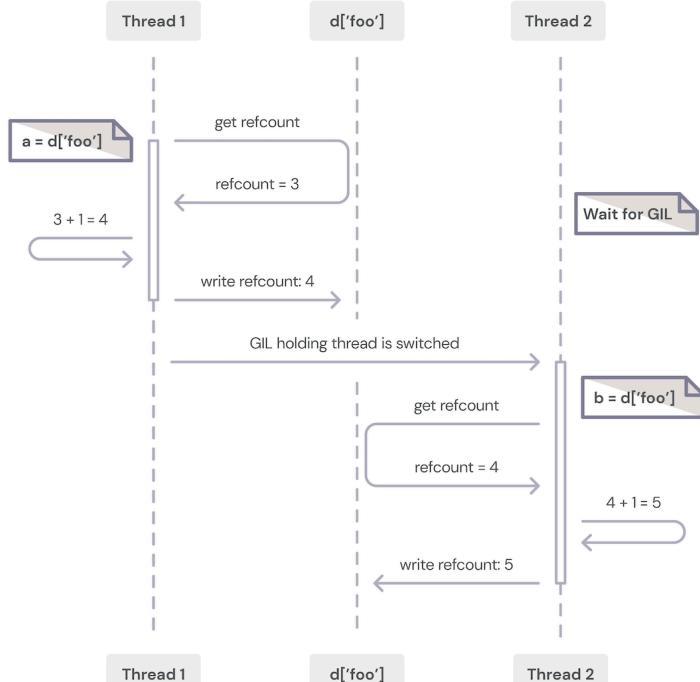
- *A lock that allows only **one thread** to execute Python bytecode at a time*
- CPython Implementation detail
- Protects internal data structures

naive GIL Removal



¹ The Python GIL - Backblaze

GIL Serializing Reference Count Increment



¹ The Python GIL - Backblaze

Commit 1984f1e



gvanrossum committed on Aug 4, 1992

- * Makefile adapted to changes below.
- * split pythonmain.c in two: most stuff goes to pythonrun.c, in the library.
- * new optional built-in threadmodule.c, build upon Sjoerd's thread.{c,h}.
- * new module from Sjoerd: mmodule.c (dynamically loaded).
- * new module from Sjoerd: sv (svgen.py, svmodule.c.proto).
- * new files thread.{c,h} (from Sjoerd).
- * new xxmodule.c (example only).
- * myselect.h: bzero -> memset
- * select.c: bzero -> memset; removed global variable

▸ **main** · ↗ v3.15.0a1 ⋯ 2.0

✓ Python/ceval.c

```
127 +
128 + void *
129 + save_thread()
130 + {
131 + #ifdef USE_THREAD
132 +     if (interpreter_lock) {
133 +         void *res;
134 +         res = (void *)current_frame;
135 +         current_frame = NULL;
136 +         release_lock(interpreter_lock);
137 +         return res;
138 +     }
139 +     else
140 +         return NULL;
141 + #endif
142 + }
143 +
144 + void
145 + restore_thread(x)
146 +     void *x;
147 +
148 + #ifdef USE_THREAD
149 +     if (interpreter_lock) {
150 +         int err;
151 +         err = errno;
152 +         acquire_lock(interpreter_lock, 1);
153 +         errno = err;
154 +         current_frame = (frameobject *)x;
155 +     }
156 + #endif
157 + }
```

Why even have threads then?

Why even have threads then?

1. I/O

- HDD Read speed: 0.5-2 MB/s
- Network Latency: 100's of ms
- Network Bandwidth: 14.4 KB/s

Why even have threads then?

Why even have threads then?

3. Simple to Extend



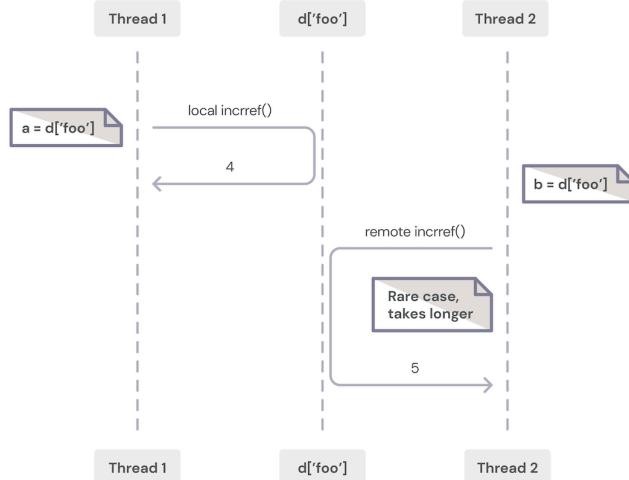
GIL removal attempts

- 1996
- 2007
- 2015 (x2)
- 2021

Core innovations

- `nogil` ~ Sam Gross
- Biased reference counting & immortal objects
- Thread-safe memory allocator
- Many more small tweaks
- **Officially accepted as PEP**

3.9-nogil



Recap: How we got here

- GIL is a CPython implementation detail to protect internal state
- Multiple removal attempts failed (1996, 2007, 2015) due to performance penalties
- Python 3.13 finally achieves free-threading via biased reference counting
- Trade-off: small single-threaded slowdown for true parallelism
- Now the question is: How do we use this?

```
uv run -p 3.14t main.py
```

Basic multi-threading

```
from concurrent.futures import ThreadPoolExecutor

N_THREADS = 8

def do_work(n: int = 10_000_000) → int:
    return sum(i * i for i in range(n))

with ThreadPoolExecutor(max_workers=N_THREADS) as ex:
    futures = [
        ex.submit(do_work)
        for _ in range(N_THREADS)
    ]
```

3.14: 2.352 s ± 0.014 s

3.14t: 768.9 ms ± 11.7 ms

~3x speedup

Single-threaded

```
def do_work(n: int = 10_000_000) → int:  
    return sum(i * i for i in range(n))  
  
for i in range(8):  
    do_work()
```

3.14: $2.341 \text{ s} \pm 0.022 \text{ s}$

3.14t: $2.370 \text{ s} \pm 0.043 \text{ s}$

1.2% overhead

Freethreading for CPU-bound applications

- True parallelism possible: 3x+ speedup for parallel CPU-bound tasks
- Best for embarrassingly parallel workloads (independent calculations)
- Single-threaded overhead now minimal (~3% vs Python 3.14)
- Thread overhead matters: use thread pools, avoid creating too many threads

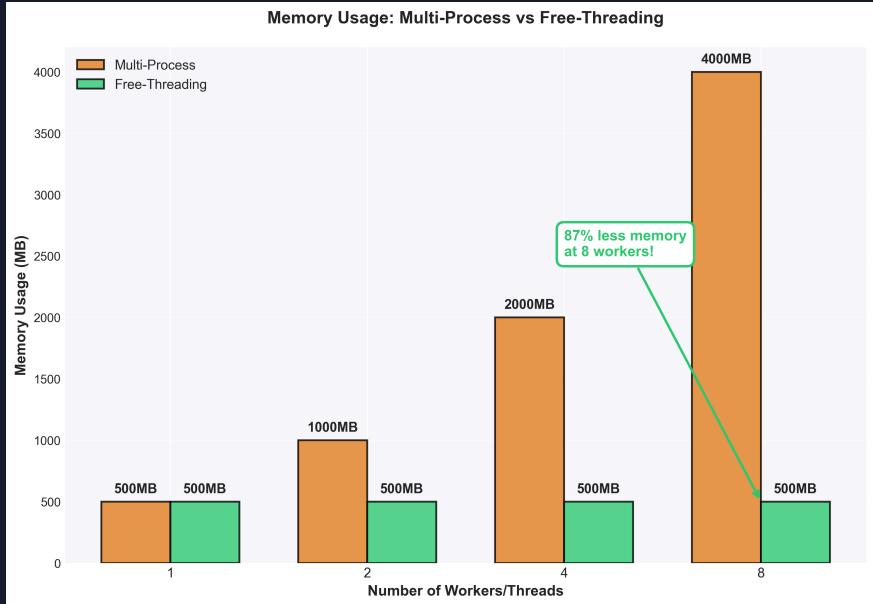
For the API developers

How Python APIs Scale Today

- Concurrent requests require parallel execution
- GIL blocks thread-based parallelism
- Solution: spawn multiple worker processes (Gunicorn, uWSGI)
- But processes have trade-offs...

Issues with processes

- Processes don't share memory space
- Data sharing is hard
- Communication between processes is slow



BUT doesn't have GIL contention

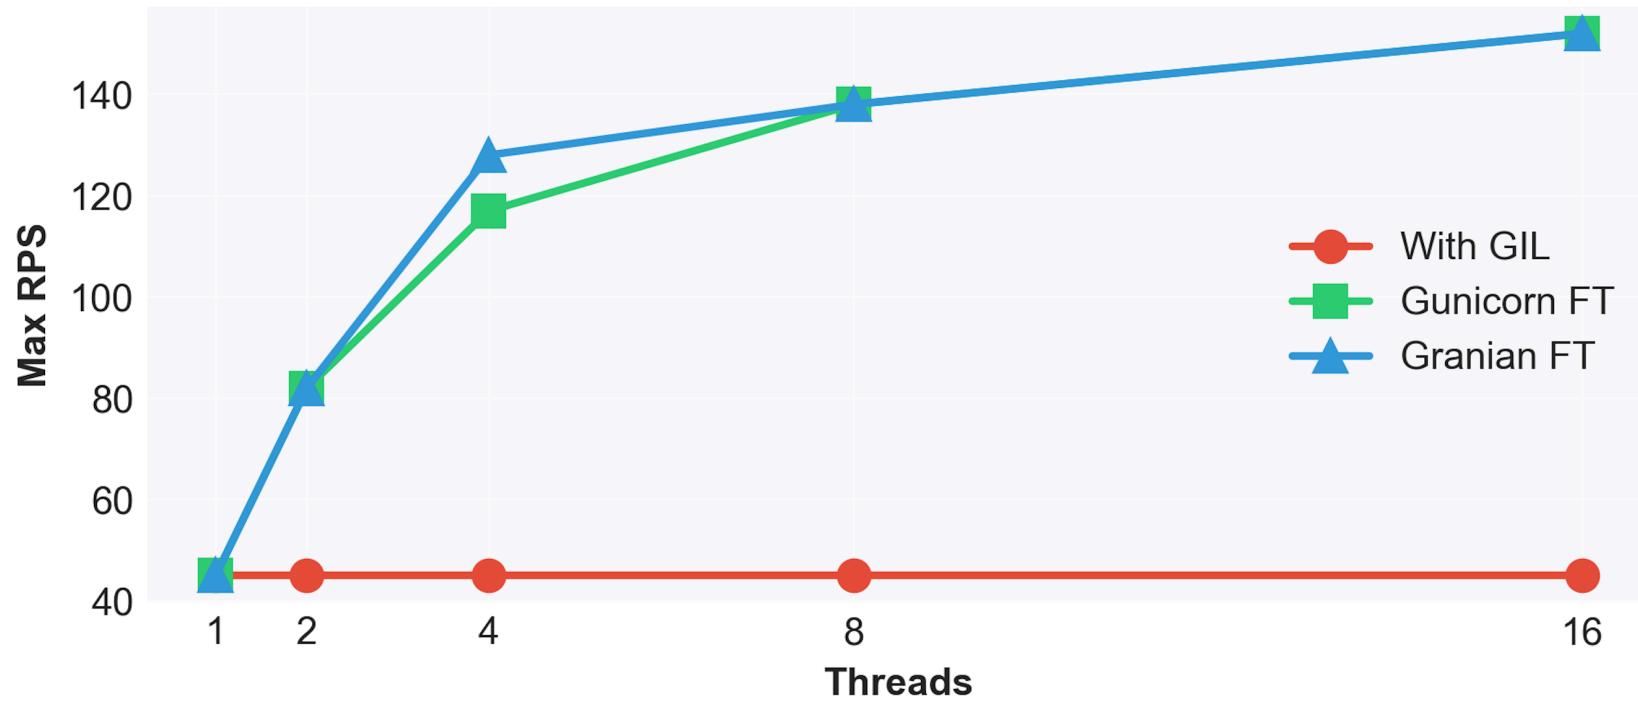
Benchmark Time

Configurations tested:

- Traditional: Gunicorn threads, Granian threads
- Async: Uvicorn/Granian ASGI, Gevent
- Free-threading: Python 3.14t with threads
- Baseline: Multi-process (current best practice)

Question: Can free-threading match multi-process performance?

CPU-bound API Throughput Comparison

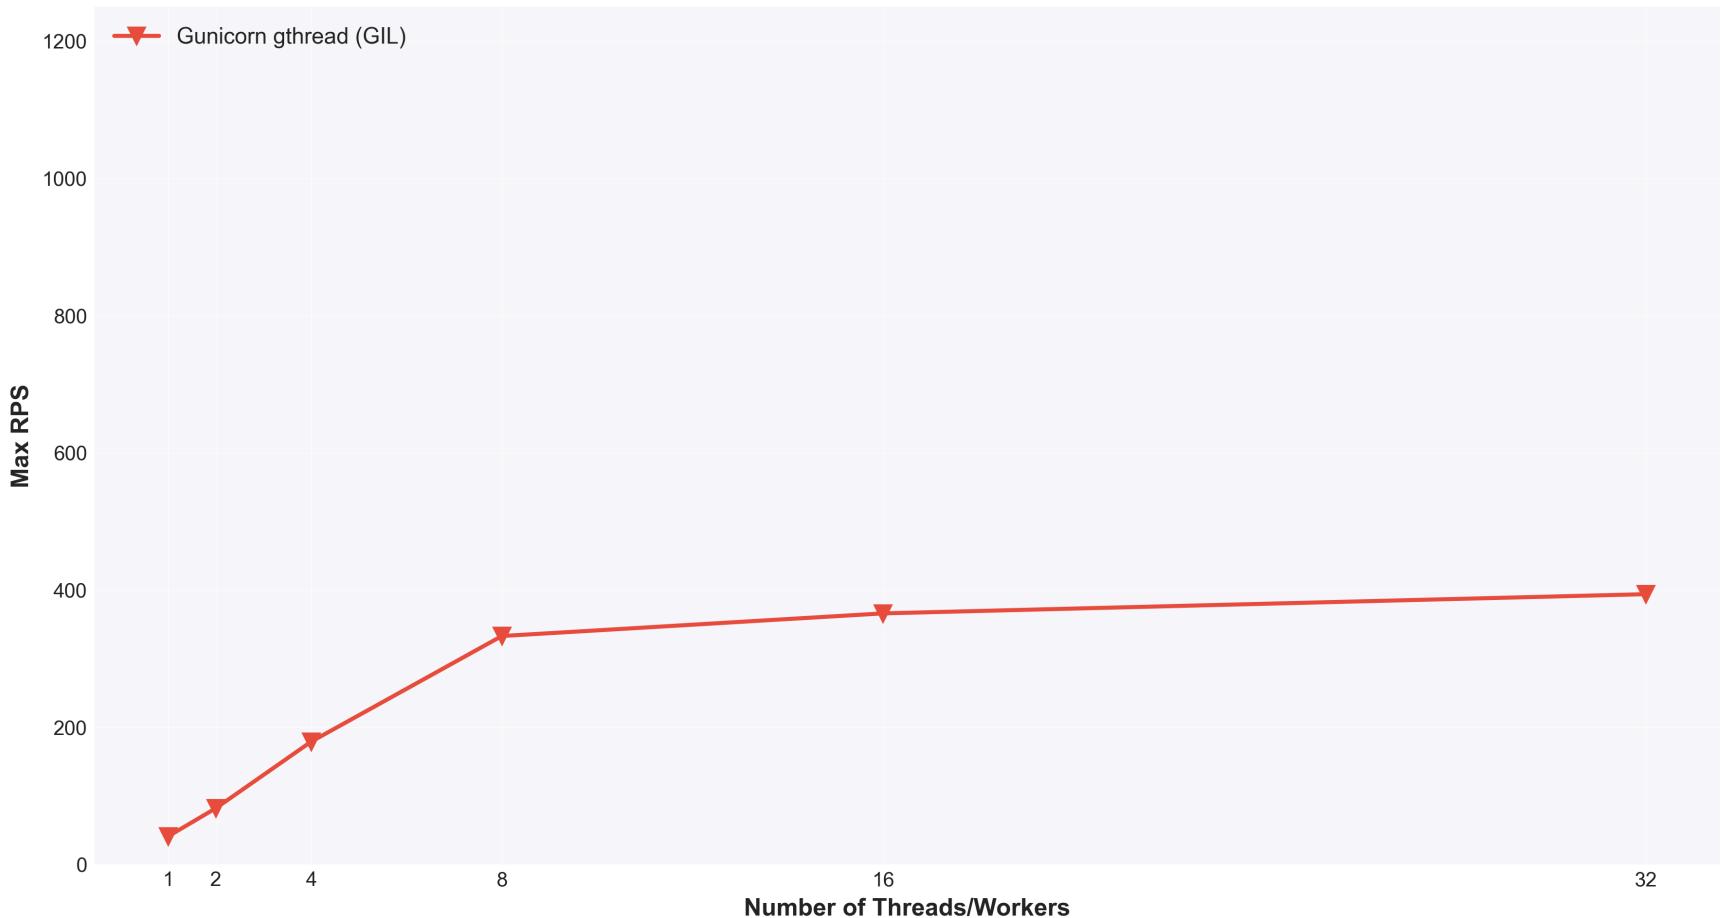


Benchmark: Mixed Workload

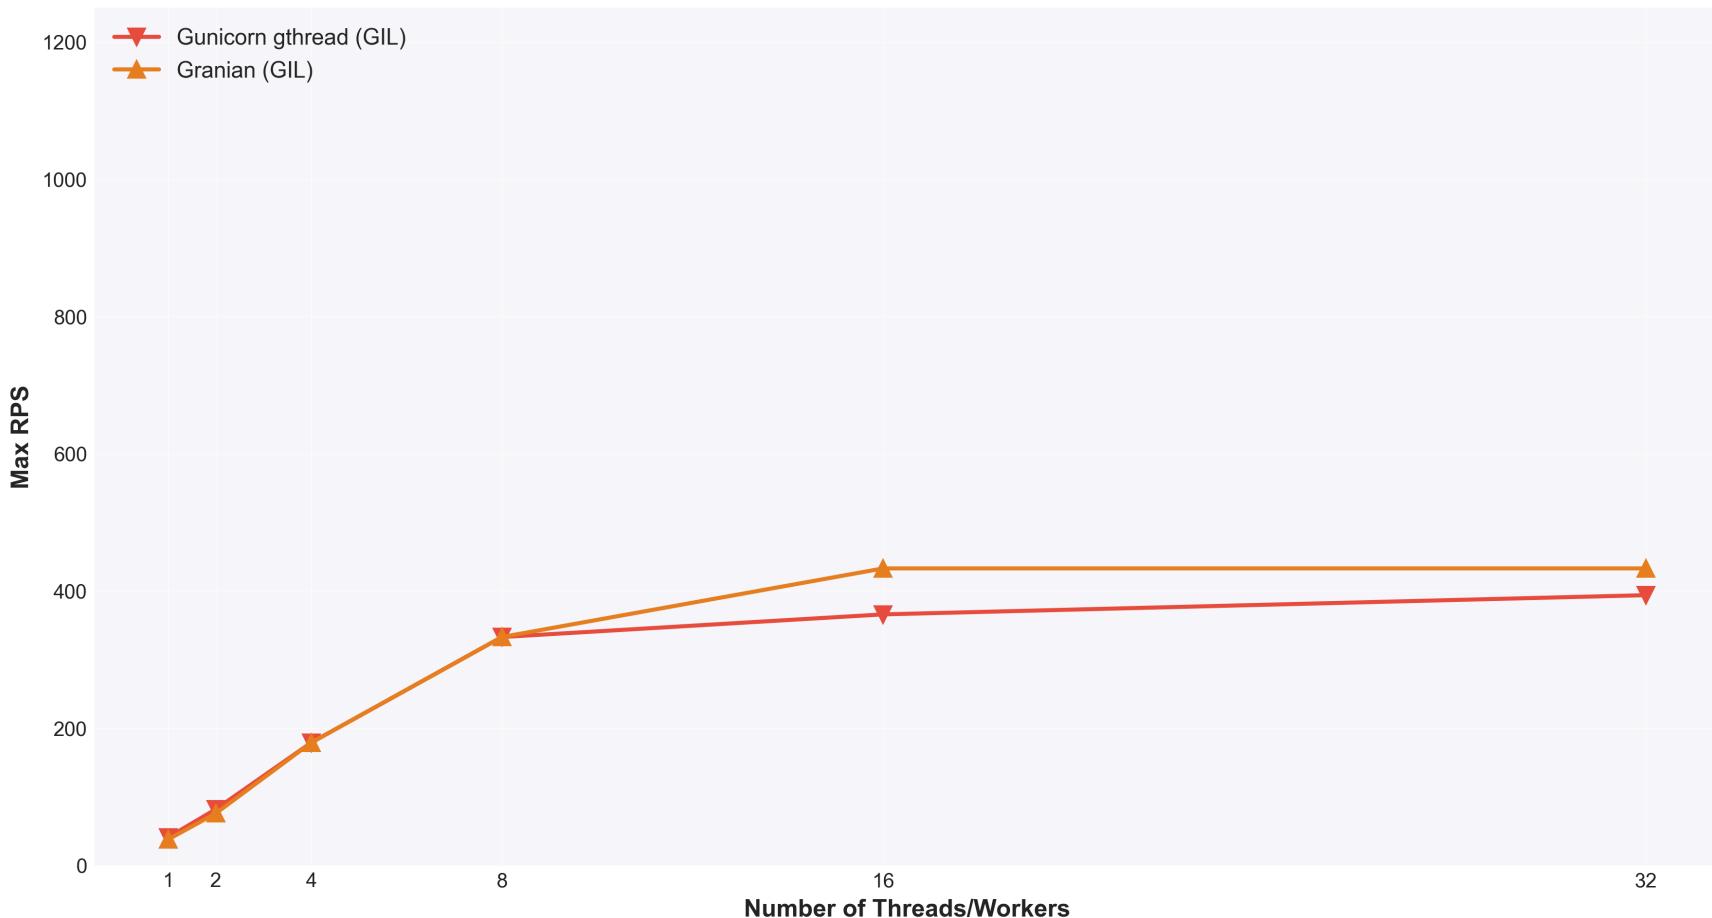
Testing a realistic API: 80% I/O operations, 20% CPU work

- Comparing different scaling strategies
- Varies thread/worker count from 1 to 32
- Measures maximum requests per second (RPS)

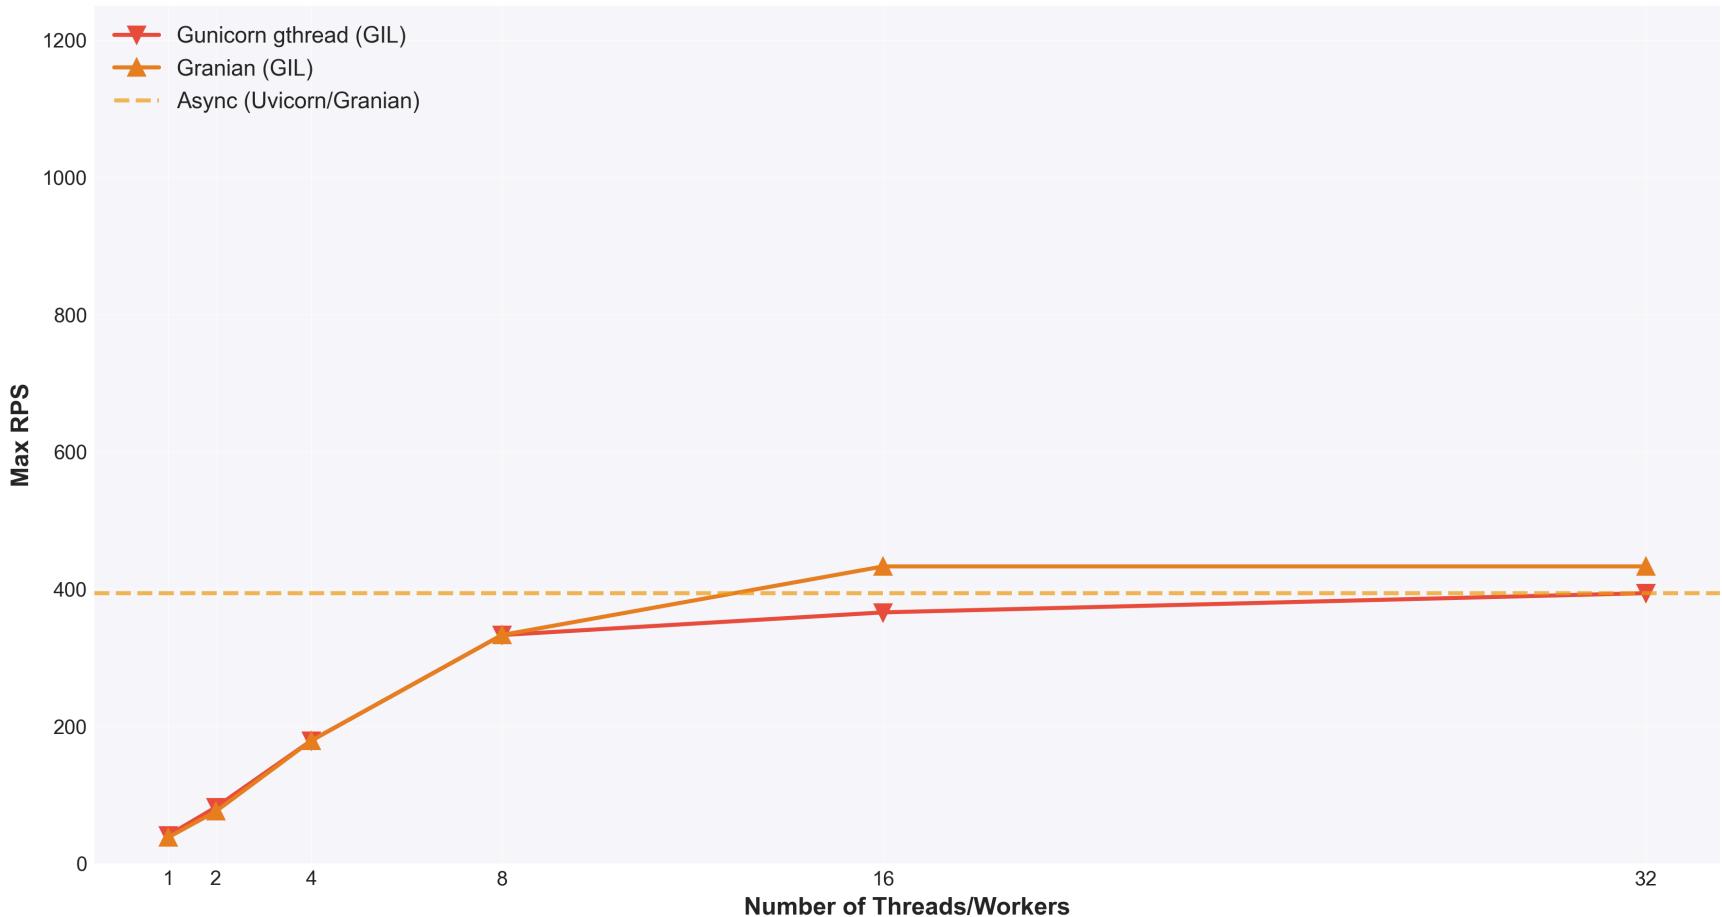
Mixed Workload Scaling (80% I/O, 20% CPU)



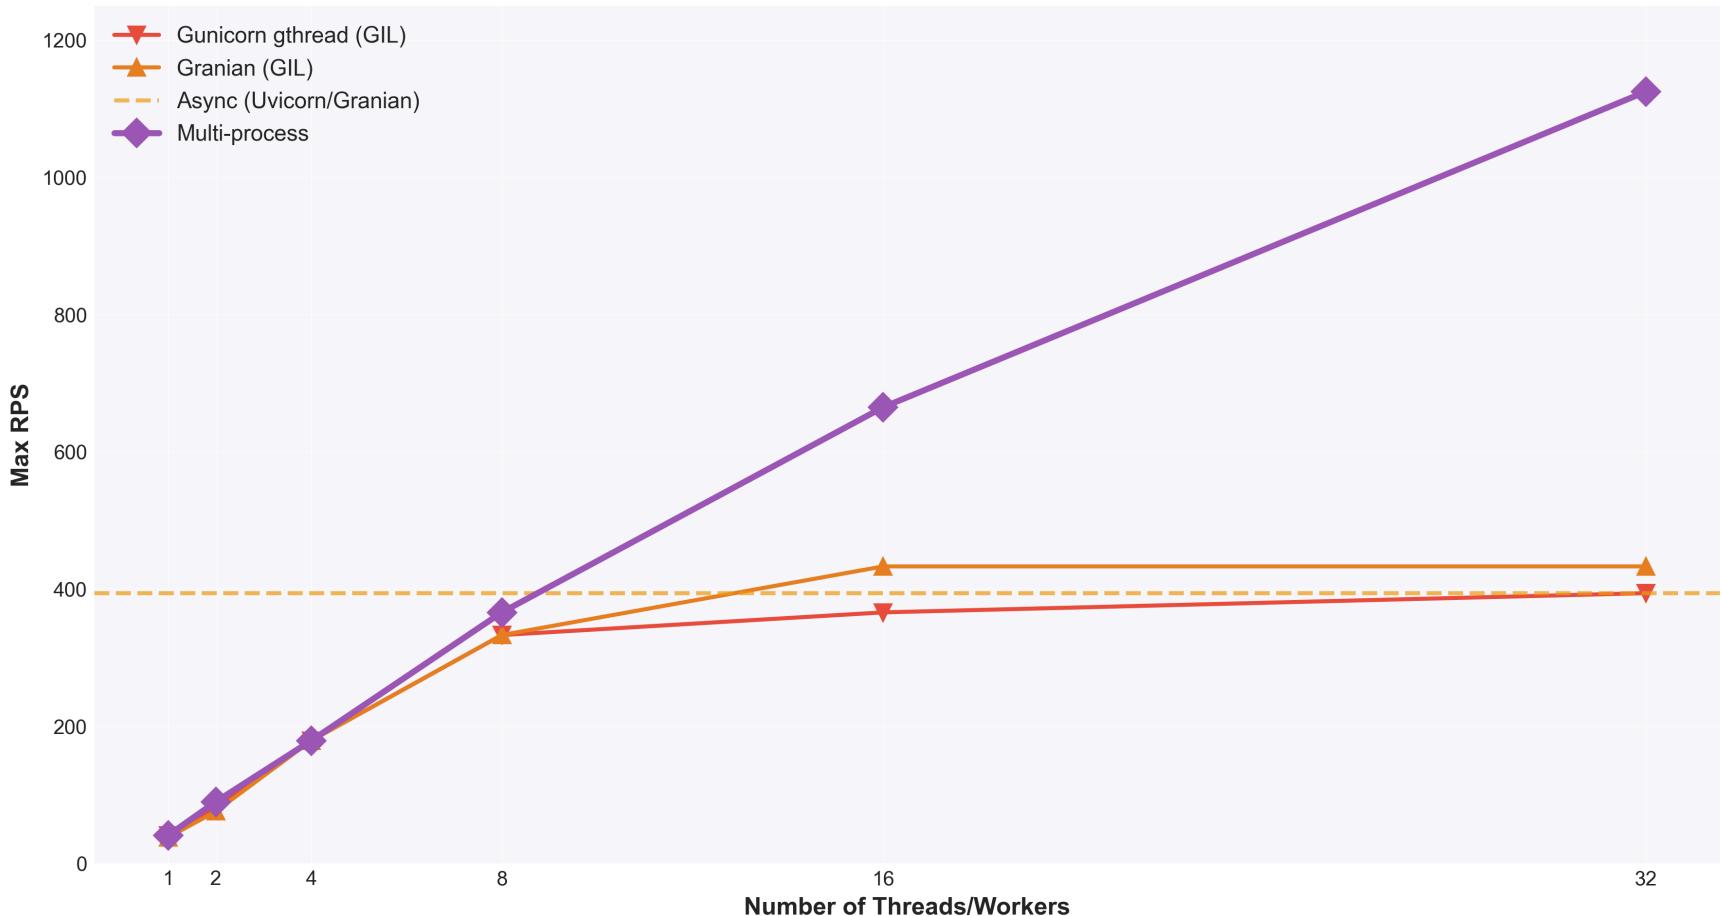
Mixed Workload Scaling (80% I/O, 20% CPU)



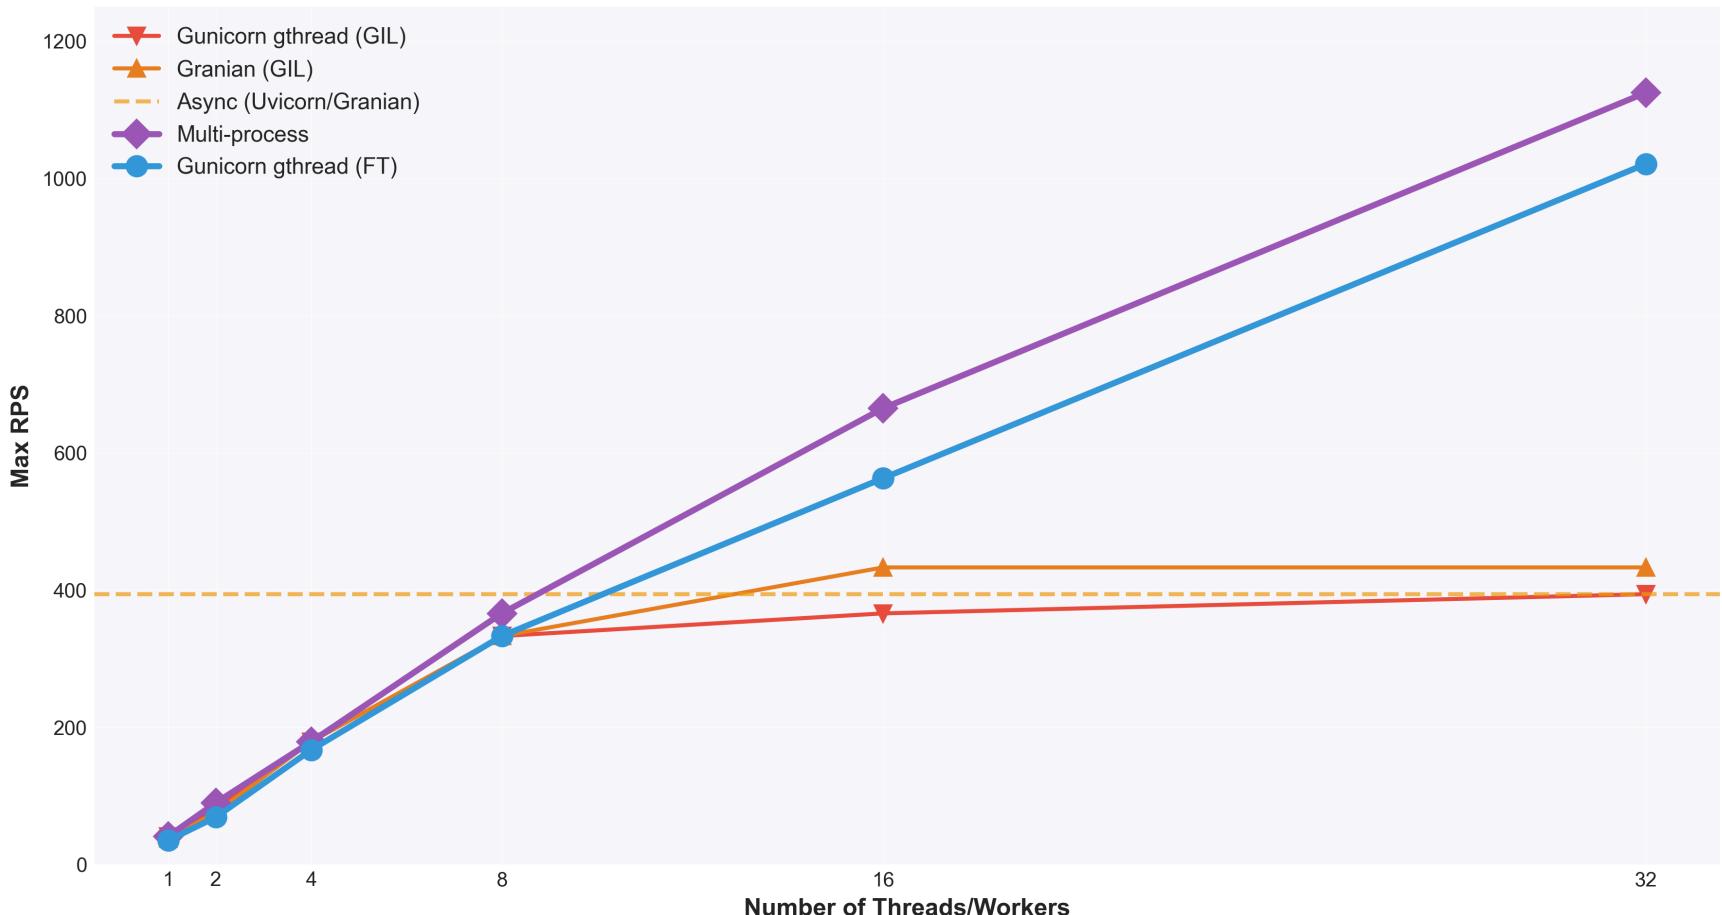
Mixed Workload Scaling (80% I/O, 20% CPU)



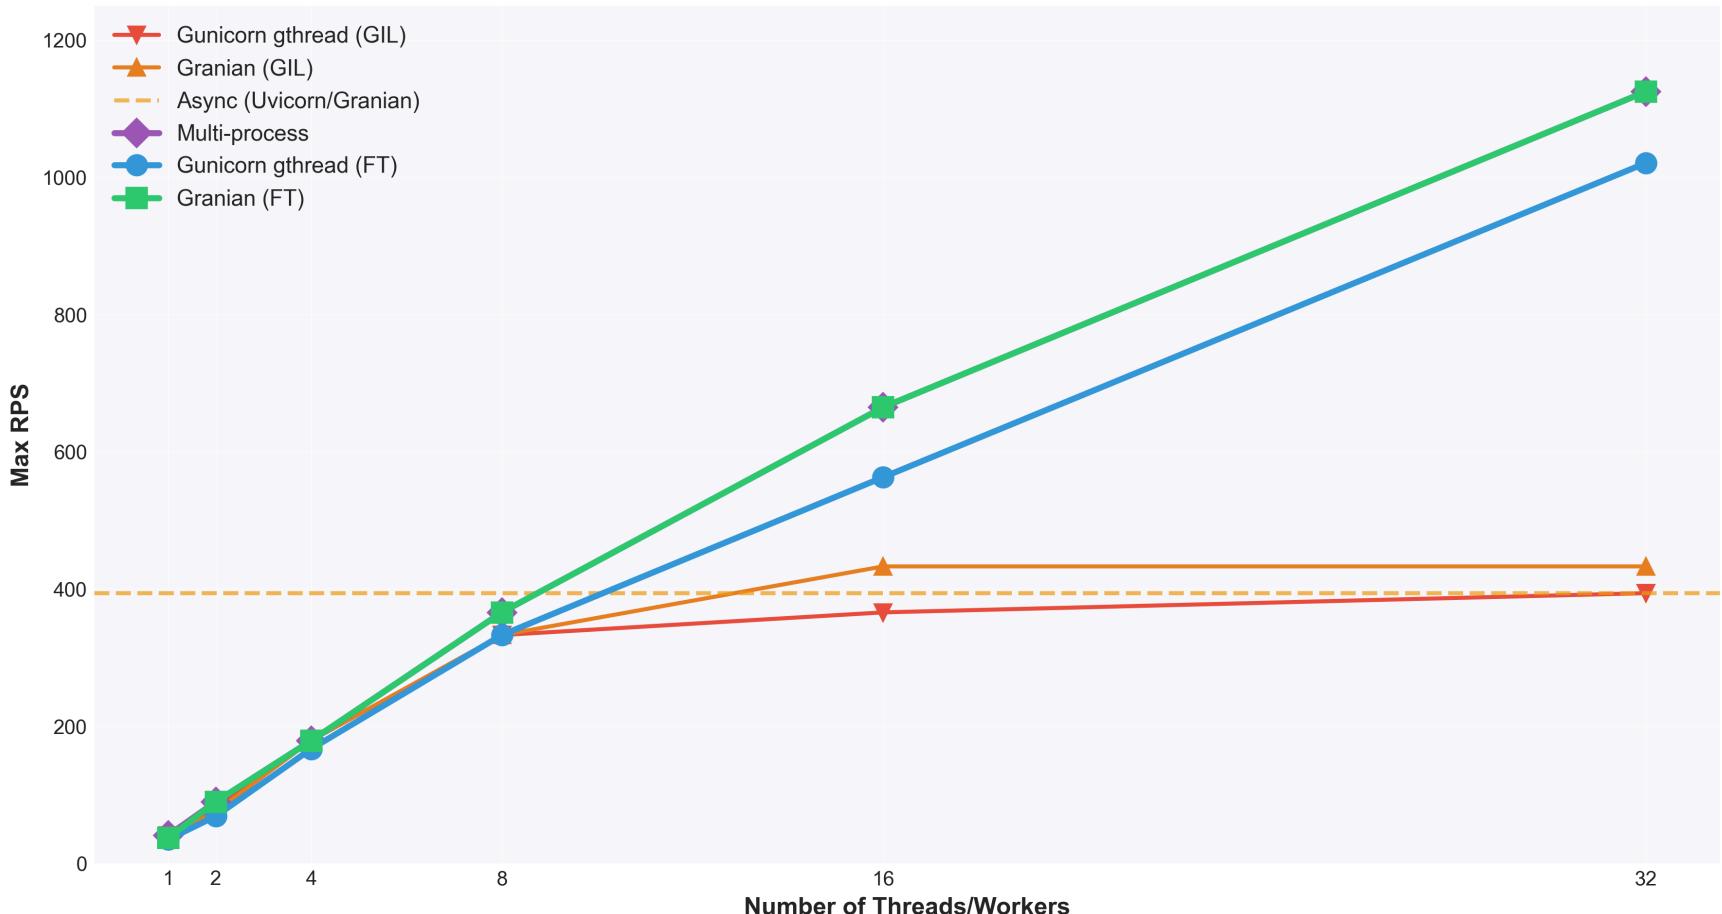
Mixed Workload Scaling (80% I/O, 20% CPU)



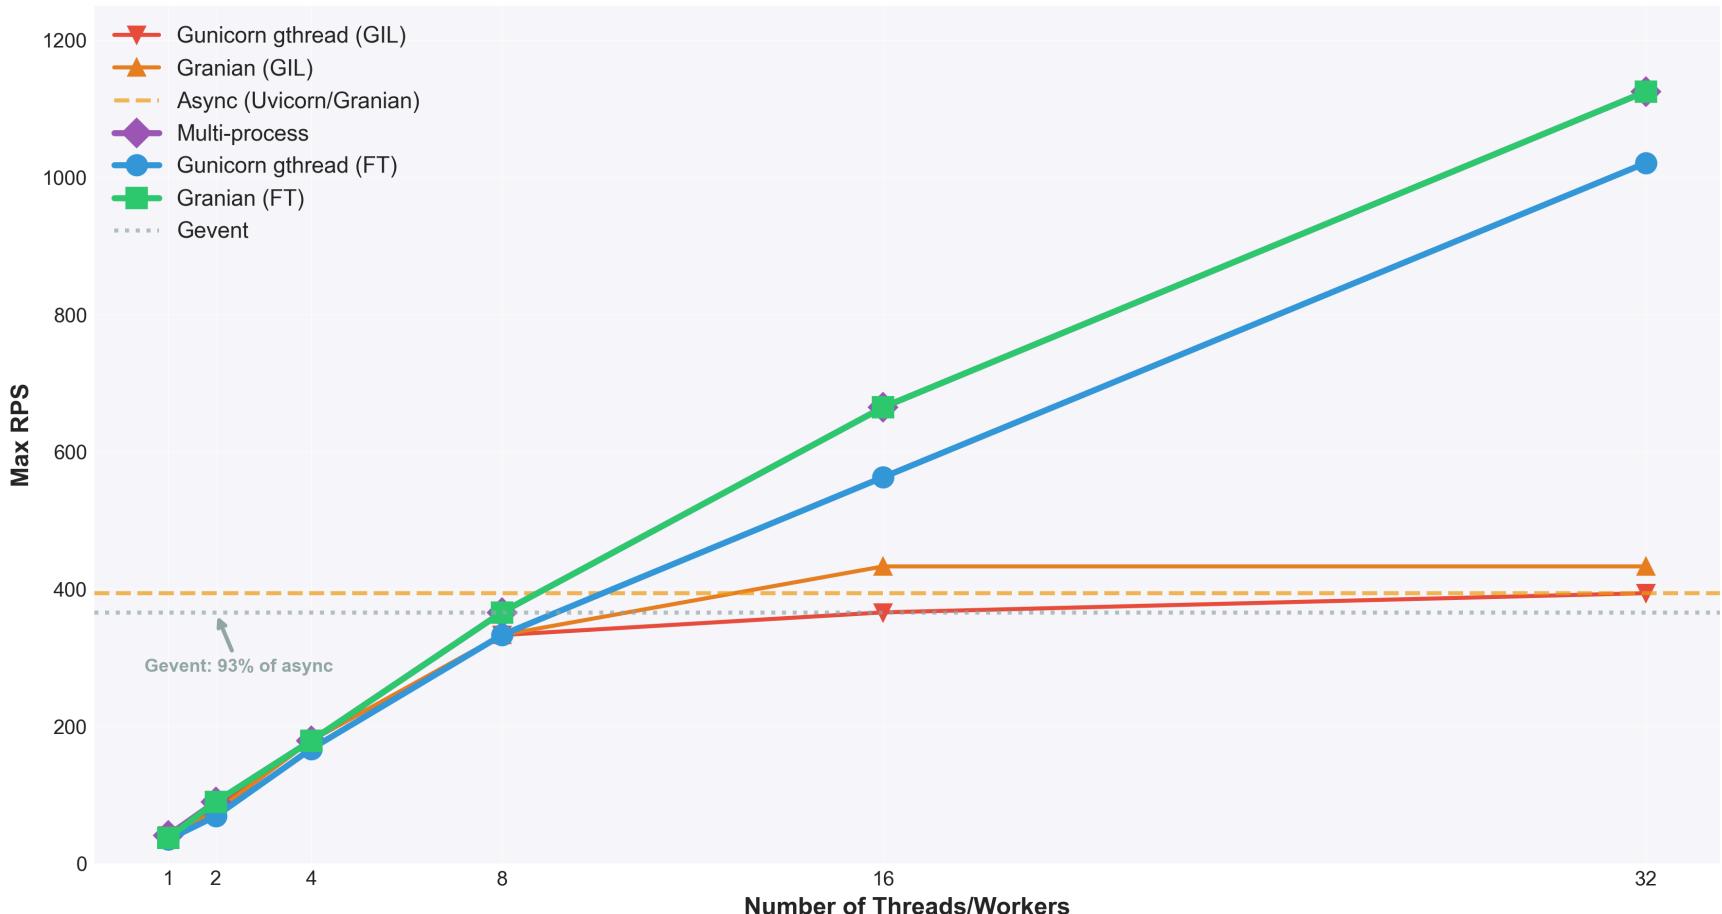
Mixed Workload Scaling (80% I/O, 20% CPU)



Mixed Workload Scaling (80% I/O, 20% CPU)



Mixed Workload Scaling (80% I/O, 20% CPU)



Key Takeaways

- GIL creates a scaling ceiling regardless of server choice
- Async helps with I/O concurrency but can't escape CPU-bound GIL limits
- Free-threading enables linear scaling: 2.5x improvement at 32 threads
- For I/O-heavy APIs: async/Gevent remain excellent choices today
- Modern webservers like Granian can better leverage threads

Conclusions

- The GIL was the right trade-off in 1997, but hardware has changed
- Free-threading delivers true parallelism without memory overhead
- 2.5x scaling improvements for mixed workloads in production-ready Python
- For new projects: consider free-threading. For existing: async/multi-process work well

The future of Python performance is multi-threaded