

# The v4 tape in the Unix history repo

Diomidis Spinellis

Department of Management Science and Technology

Athens University of Economics and Business

&

Department of Software Technology

Delft University of Technology



[www.spinellis.gr](http://www.spinellis.gr)



DSpinellis



@CoolSWEng

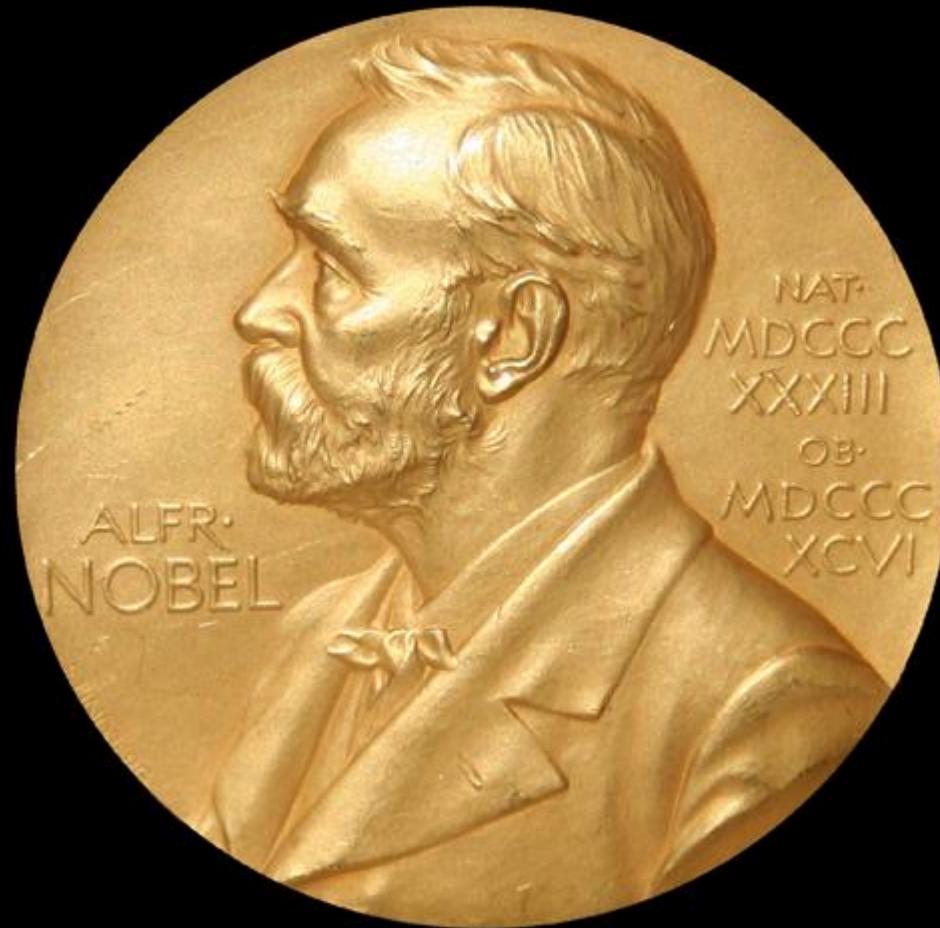


@CoolSWEng@mastodon.acm.org



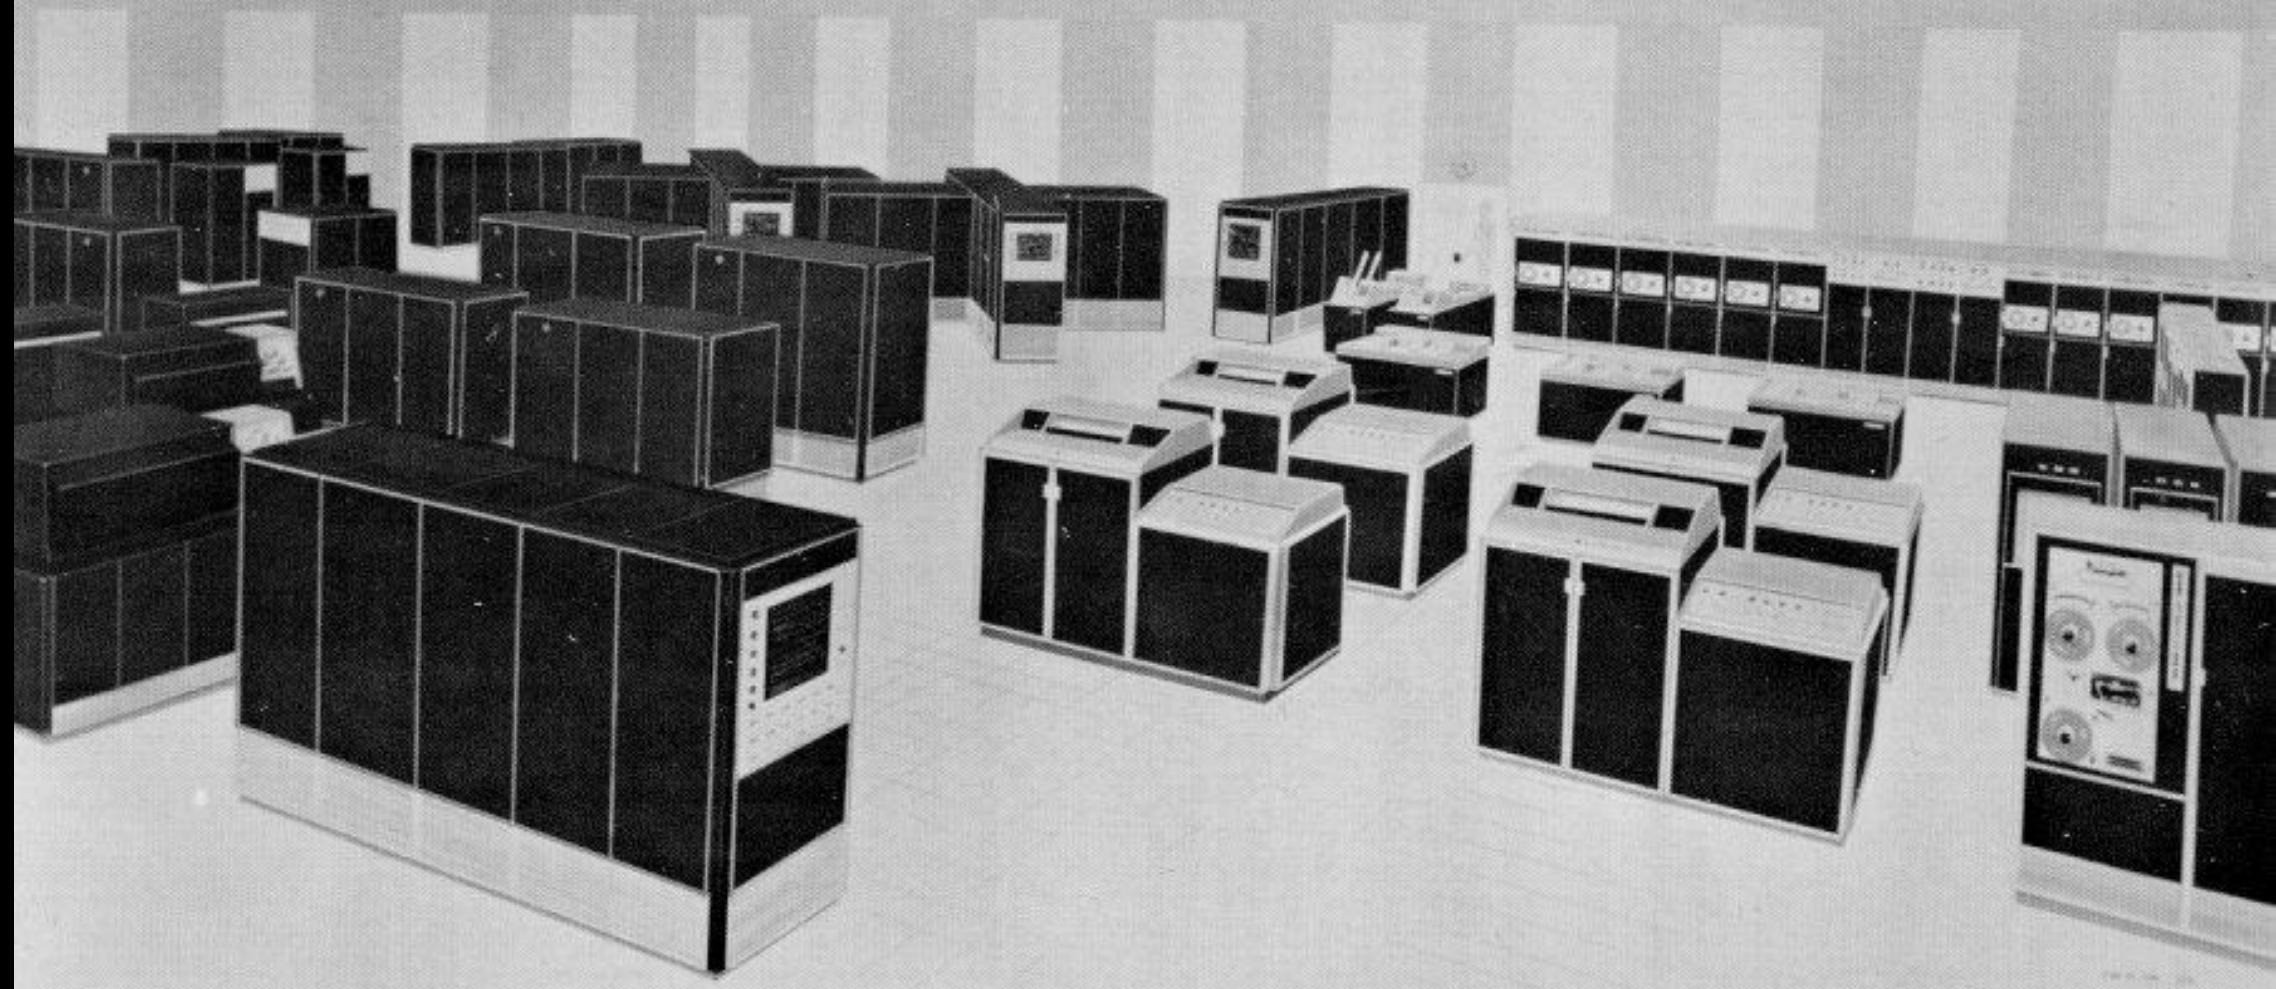
dds@aeub.gr



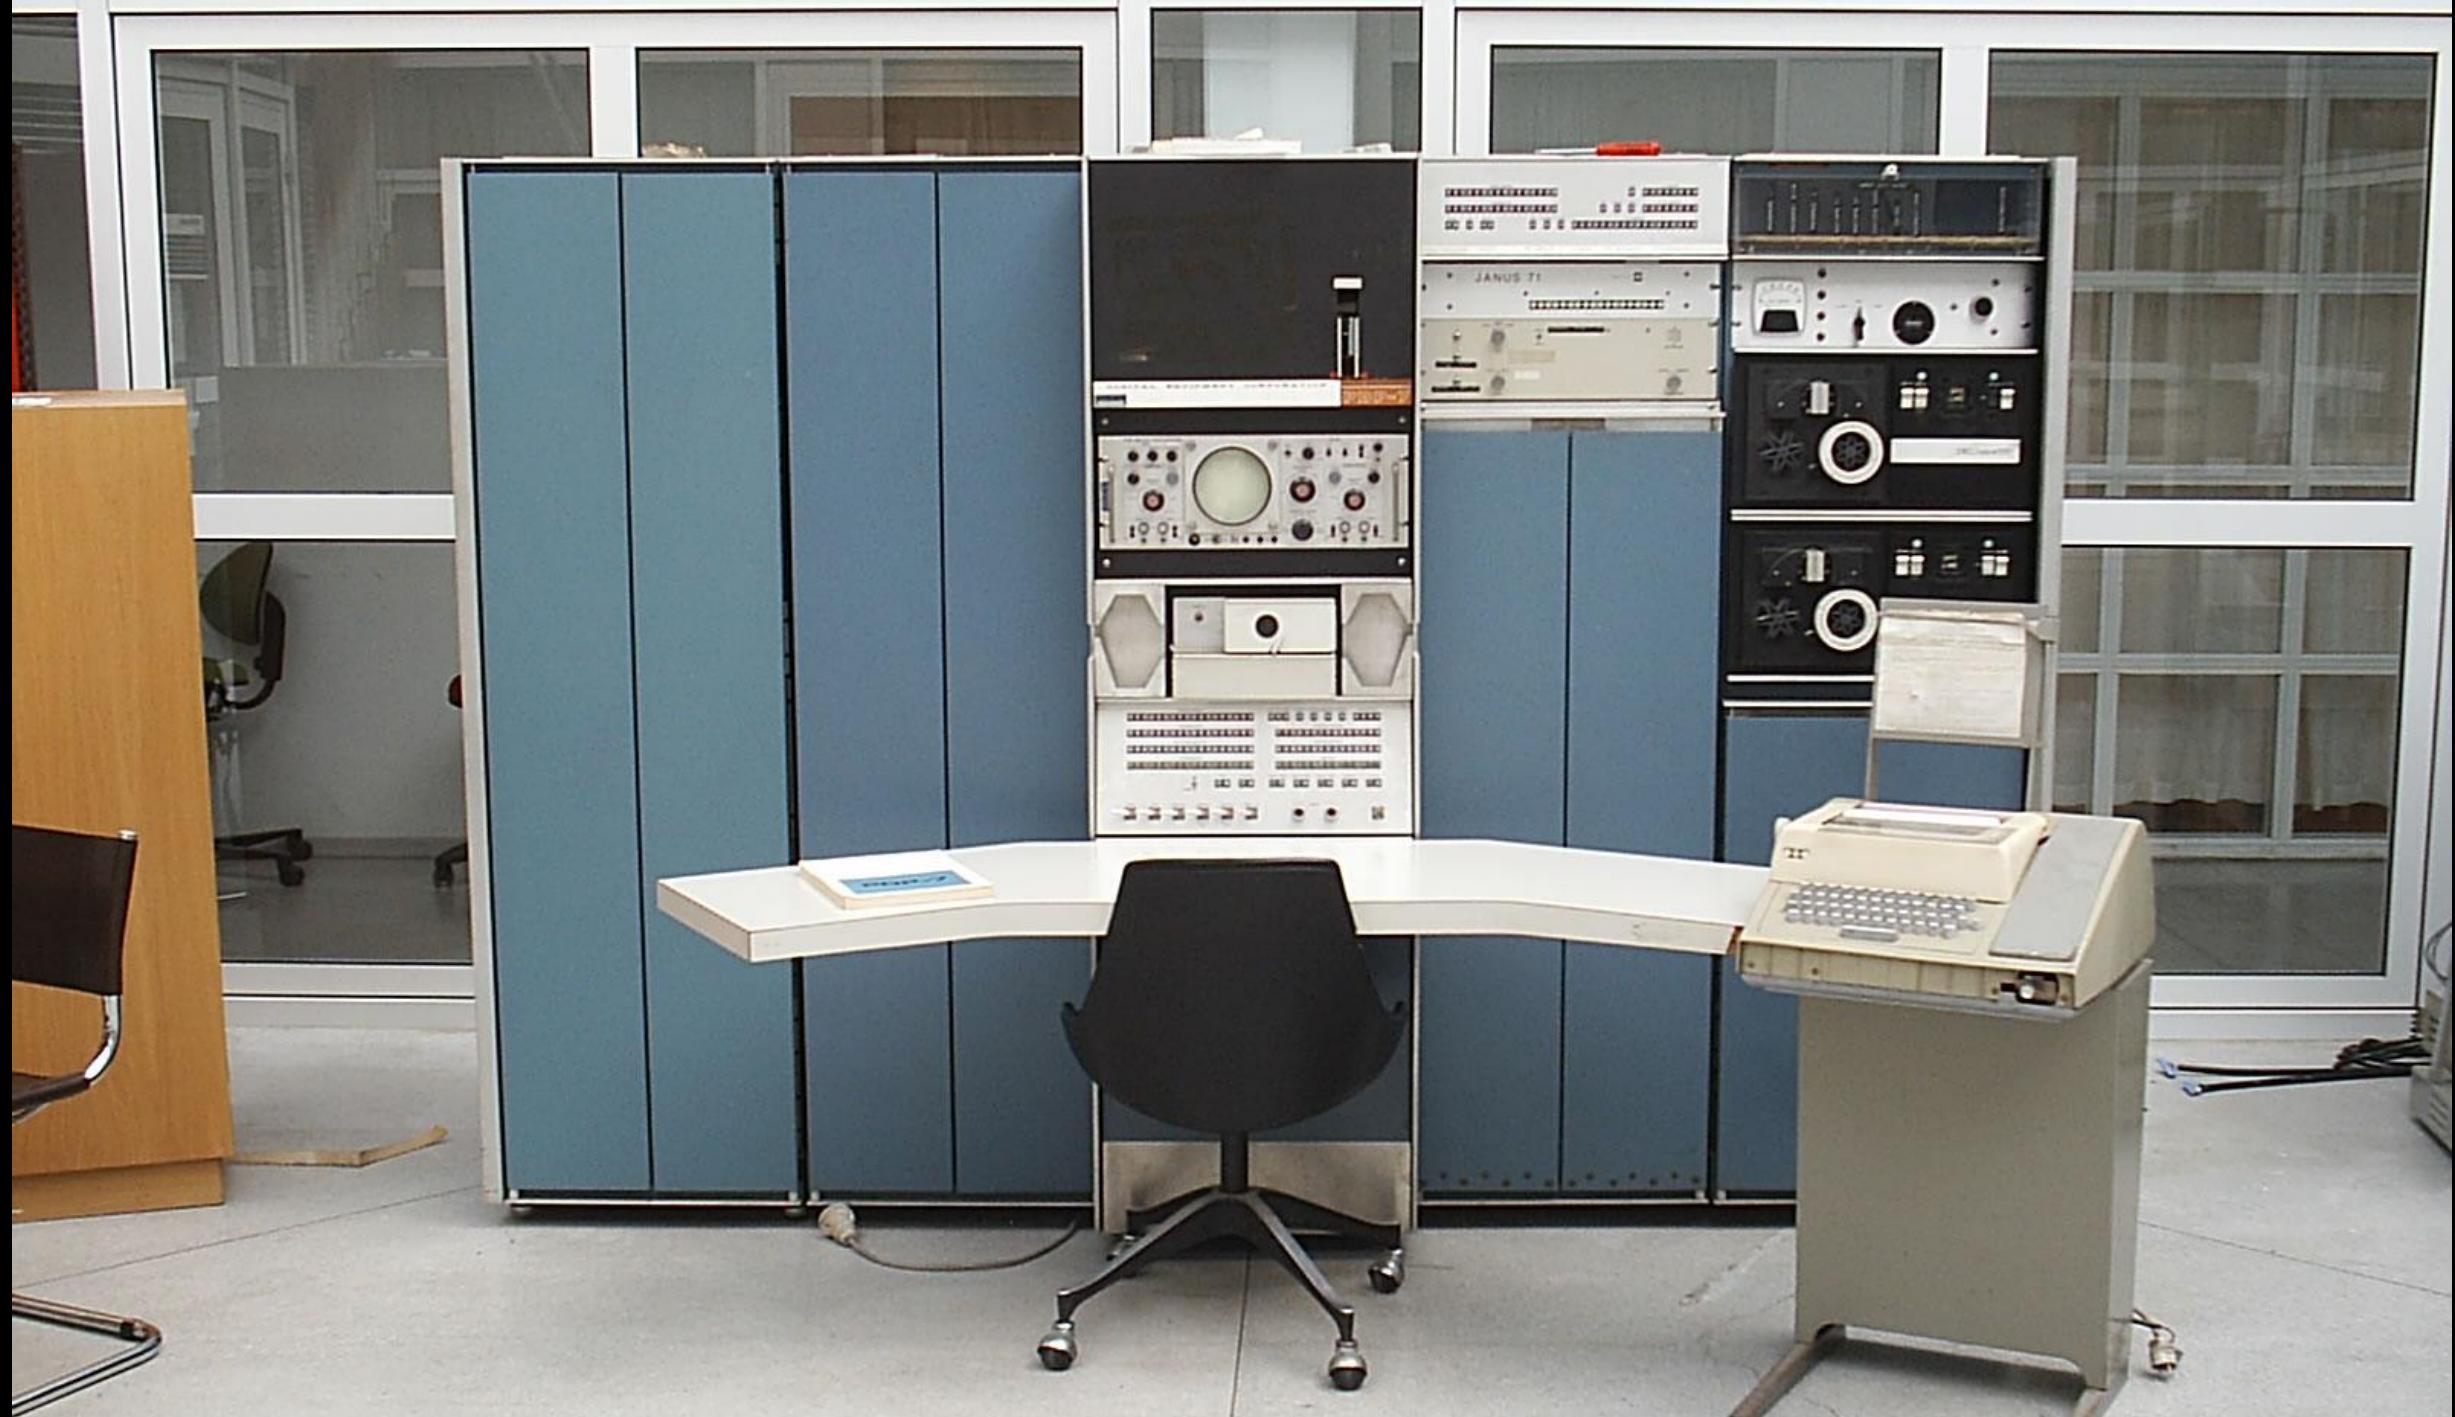




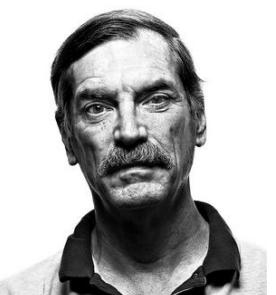
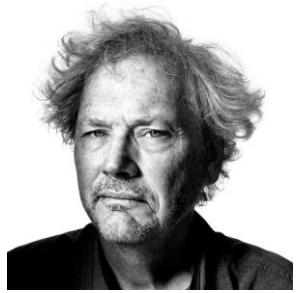
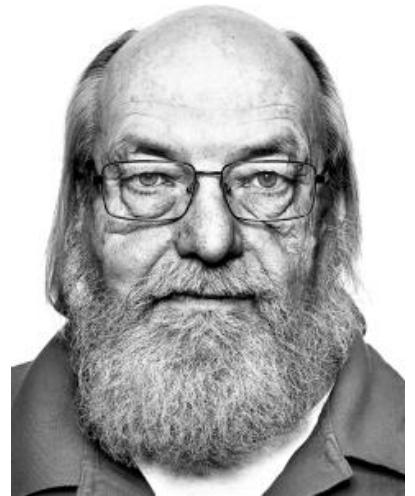
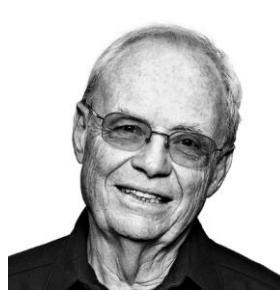
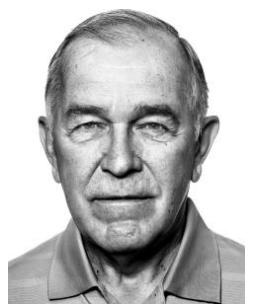
ACM  
A.M. TURING AWARD



**GE-645 SYSTEM**







UNIX PROGRAMMER'S MANUAL

K. Thompson  
D. M. Ritchie

November 3, 1971

UNIX PROGRAMMER'S MANUAL  
Second Edition

K. Thompson  
D. M. Ritchie

June 12, 1972

Copyright © 1972  
Bell Telephone Laboratories, Inc.  
No part of this document may be reproduced,  
or distributed outside the Laboratories, without  
the written permission of Bell Telephone Laboratories.

There is no warranty of merchantability nor any warranty  
of fitness for a particular purpose, either expressed or implied,  
in respect to this document or to the accuracy of the  
enclosed materials or as to their suitability for any  
particular purpose. Accordingly, Bell Telephone  
Laboratories assumes no responsibility for their use by  
the recipient. Furthermore, Bell Telephone Laboratories has no obligation  
to furnish any assistance of any kind whatsoever, or to  
furnish any additional information or documentation.

UNIX PROGRAMMER'S MANUAL

Fifth Edition

K. Thompson  
D. M. Ritchie

June, 1974

Copyright © 1972, 1973, 1974  
Bell Telephone Laboratories, Incorporated

UNIX PROGRAMMER'S MANUAL  
Third Edition

K. Thompson  
D. M. Ritchie

February, 1973

Copyright © 1972  
Bell Telephone Laboratories, Inc.  
No part of this document may be reproduced,  
or distributed outside the Laboratories, without  
the written permission of Bell Telephone Laboratories.

UNIX PROGRAMMER'S MANUAL

Sixth Edition

K. Thompson  
D. M. Ritchie

May, 1975

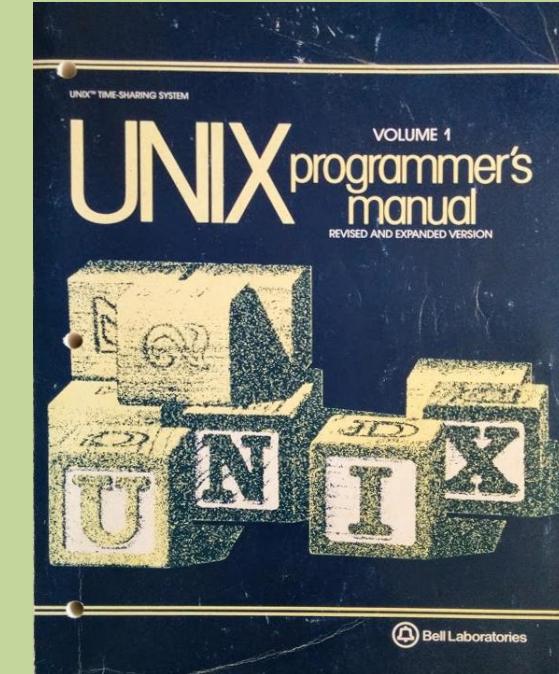
UNIX PROGRAMMER'S MANUAL

Fourth Edition

K. Thompson  
D. M. Ritchie

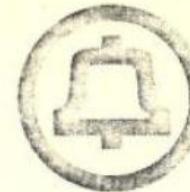
November, 1973

Copyright © 1972, 1973  
Bell Telephone Laboratories, Inc.  
No part of this document may be reproduced,  
or distributed outside the Laboratories, without  
the written permission of Bell Telephone Laboratories.



# Surviving materials

	v0	v1	v2	v3	v4
Kernel	✓	✓		✓	
Commands	○	○	○		
Documentation				✓	✓
C compiler			✓	✓	



## Bell Laboratories

600 Mountain Avenue  
Murray Hill, New Jersey 07974  
Phone (201) 582-3000

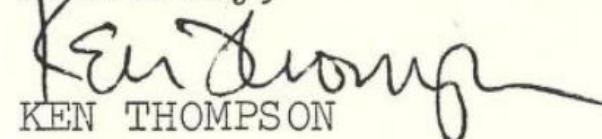
May 31, 1974

MR. MARTIN E. NEWELL  
The University of Utah  
College of Engineering  
Salt Lake City, Utah 84112

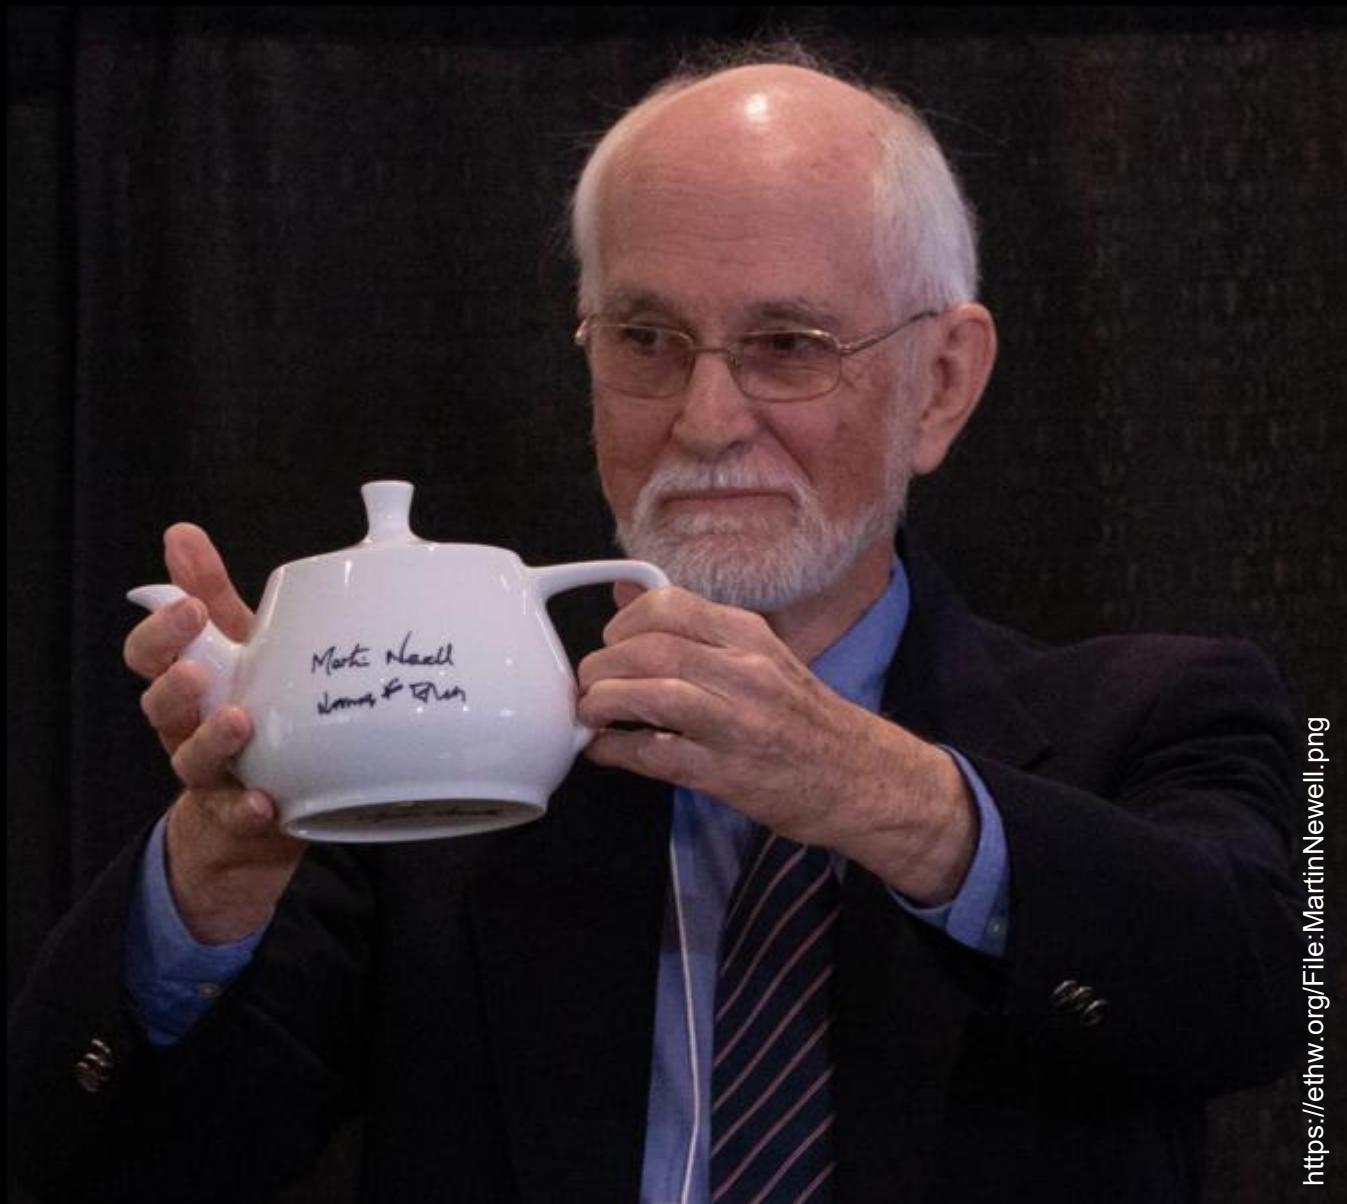
Dear Mr. Newell:

I am sorry for the delay in delivering UNIX. Unfortunately we've run out of copies of documentation. As soon as a new batch comes from the printers I will send you the system.

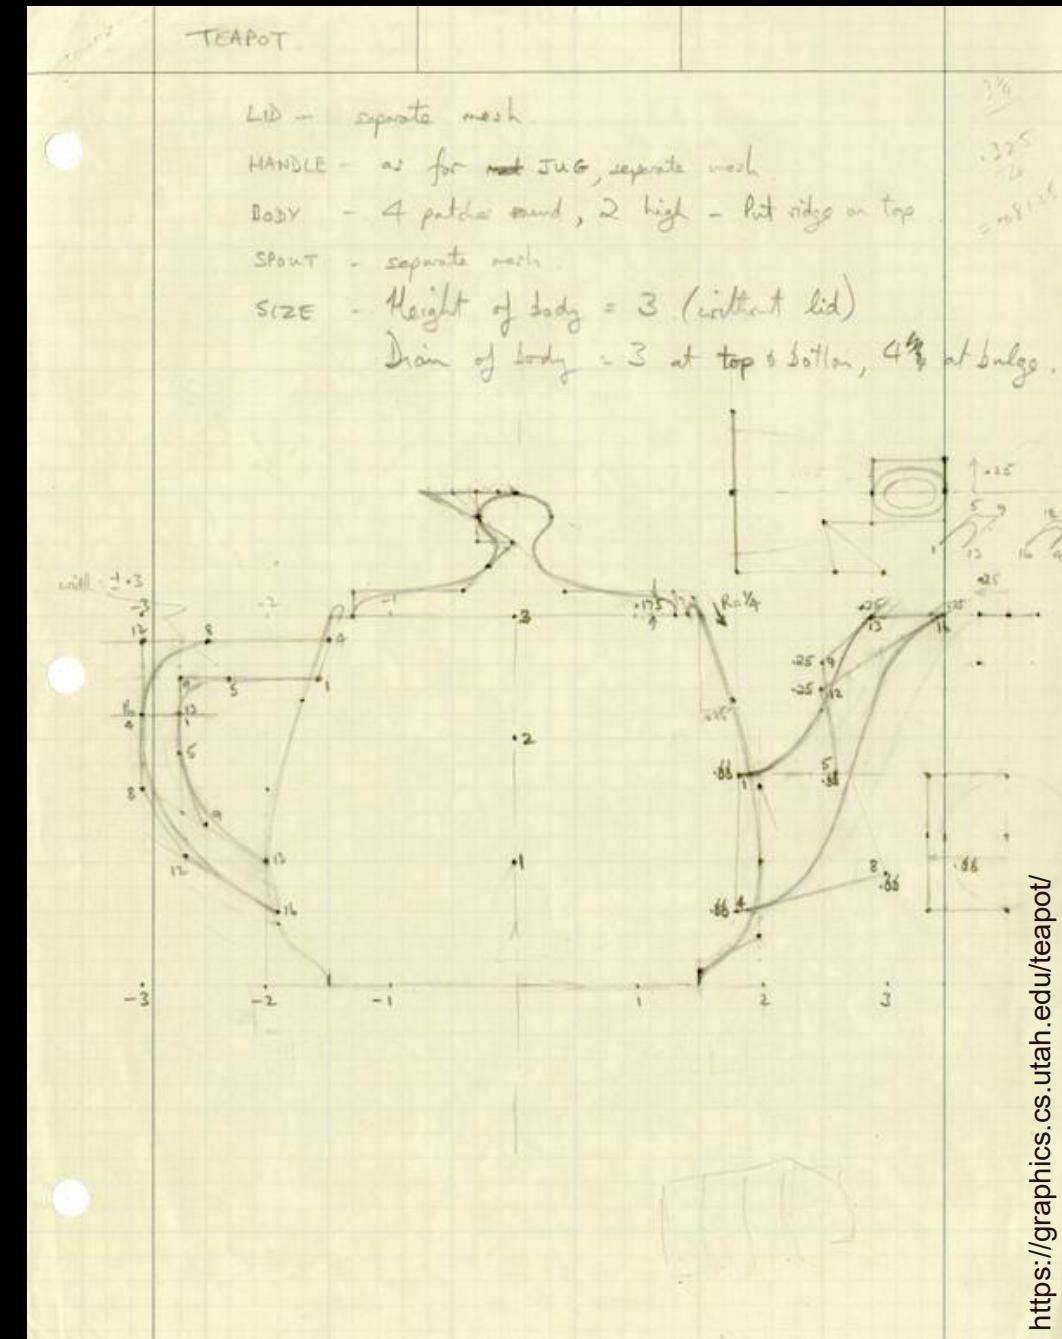
Sincerely,

  
KEN THOMPSON

KT-gam



<https://ethw.org/File:MartinNewell.png>



<https://graphics.cs.utah.edu/teapot/>

UNIX Original  
from Bell Labs v4  
(See Manual for format)

SHL LA1  
UNIVERSITY OF  
MICHIGAN  
SO CENTER



CONTENTS  
ACCEPTED

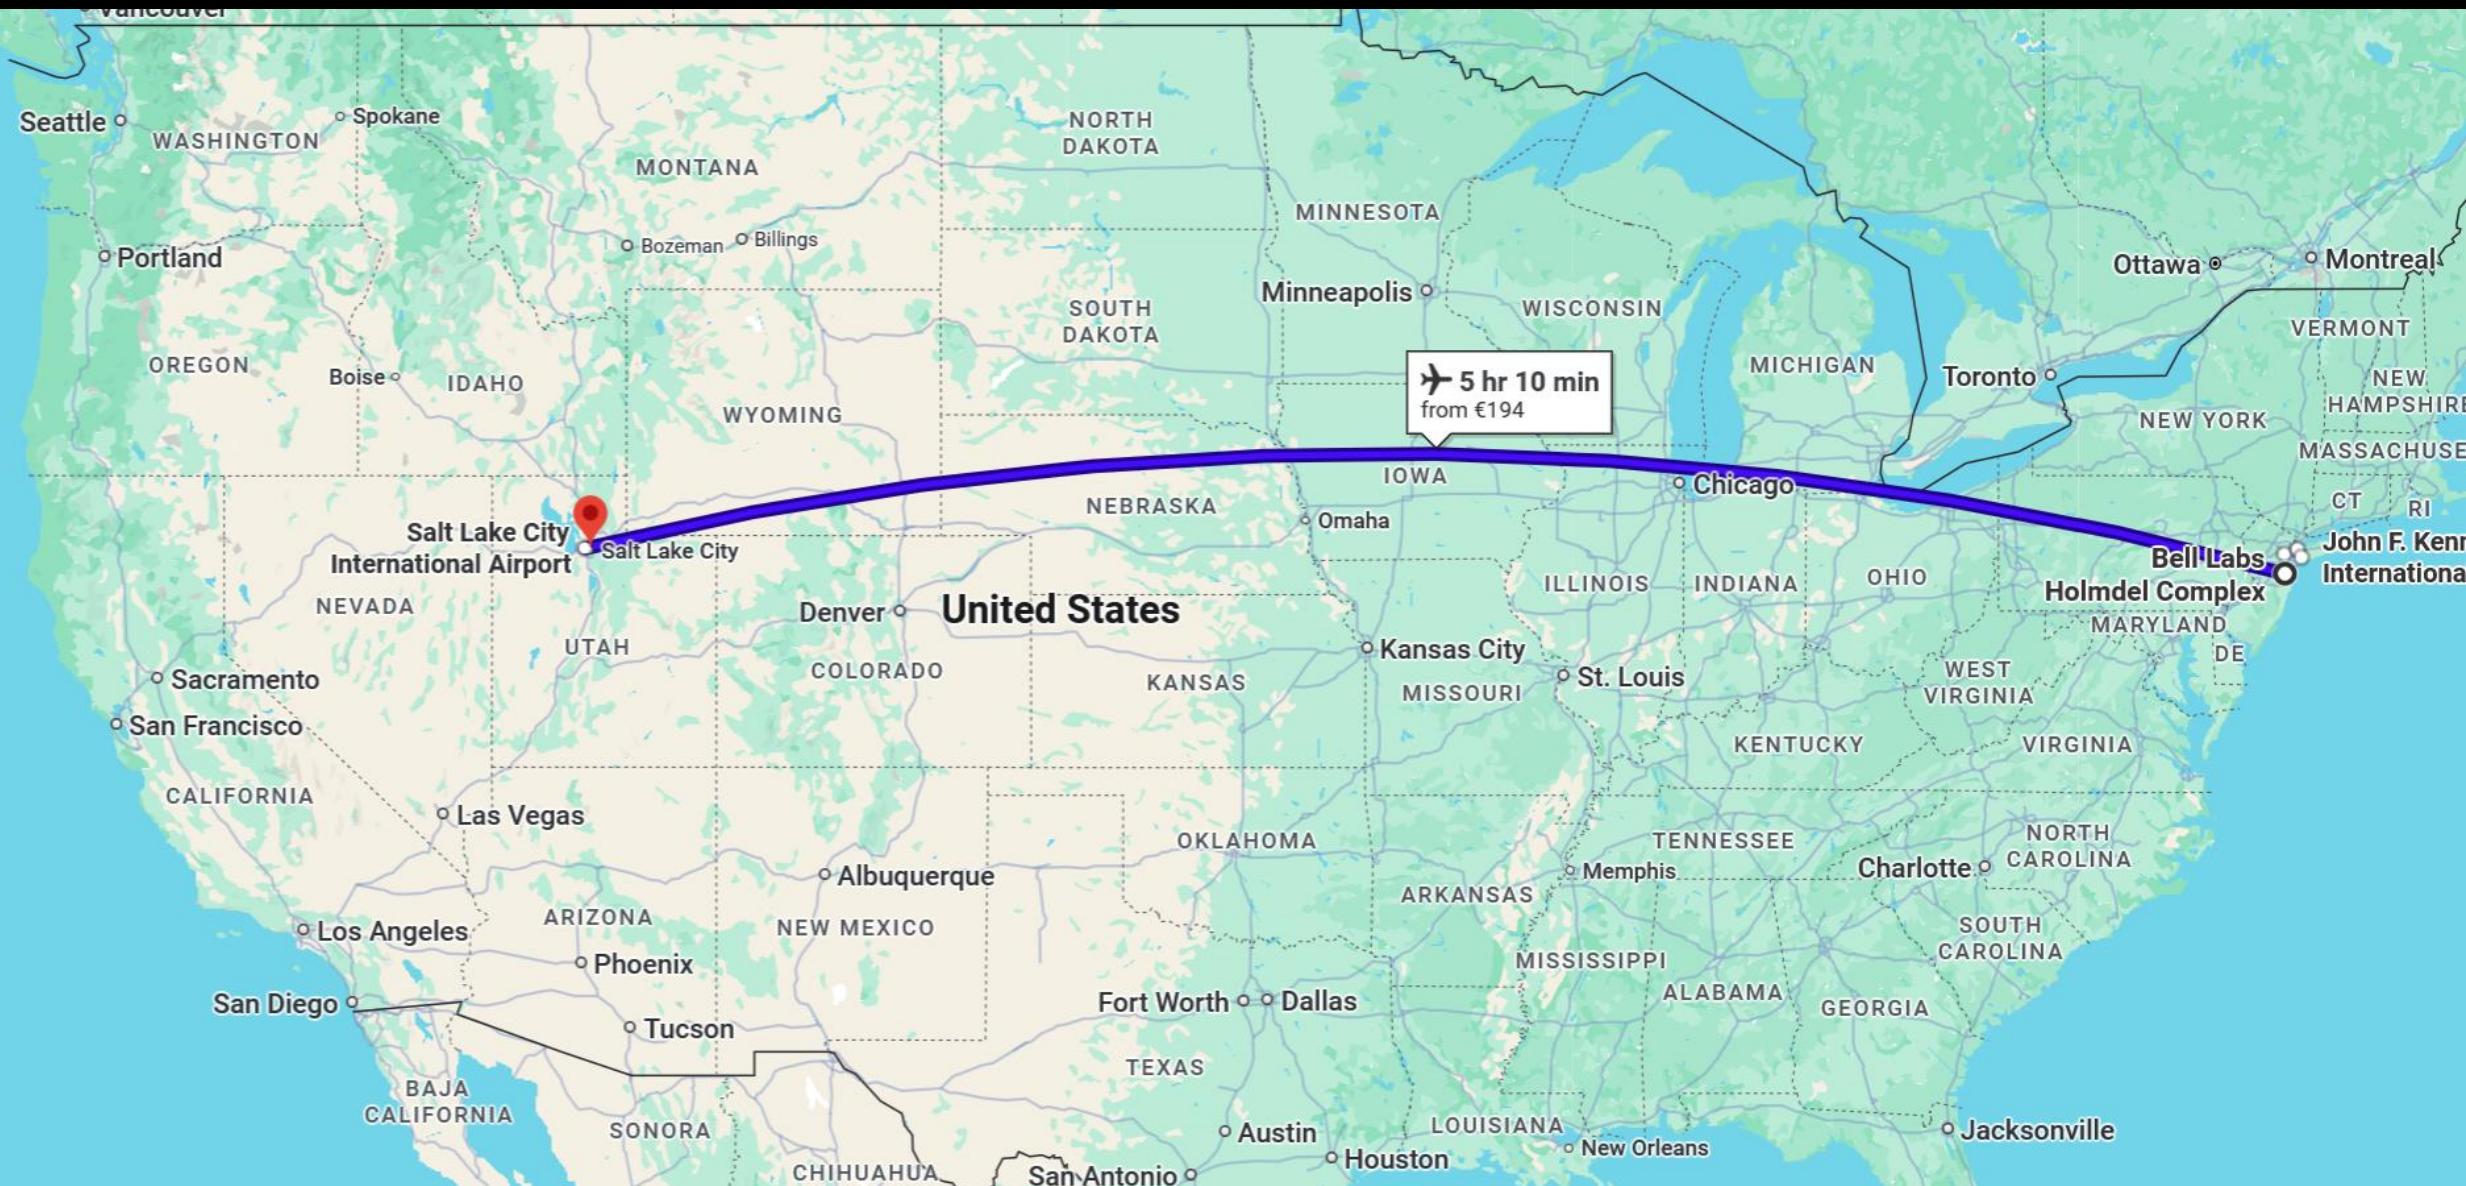
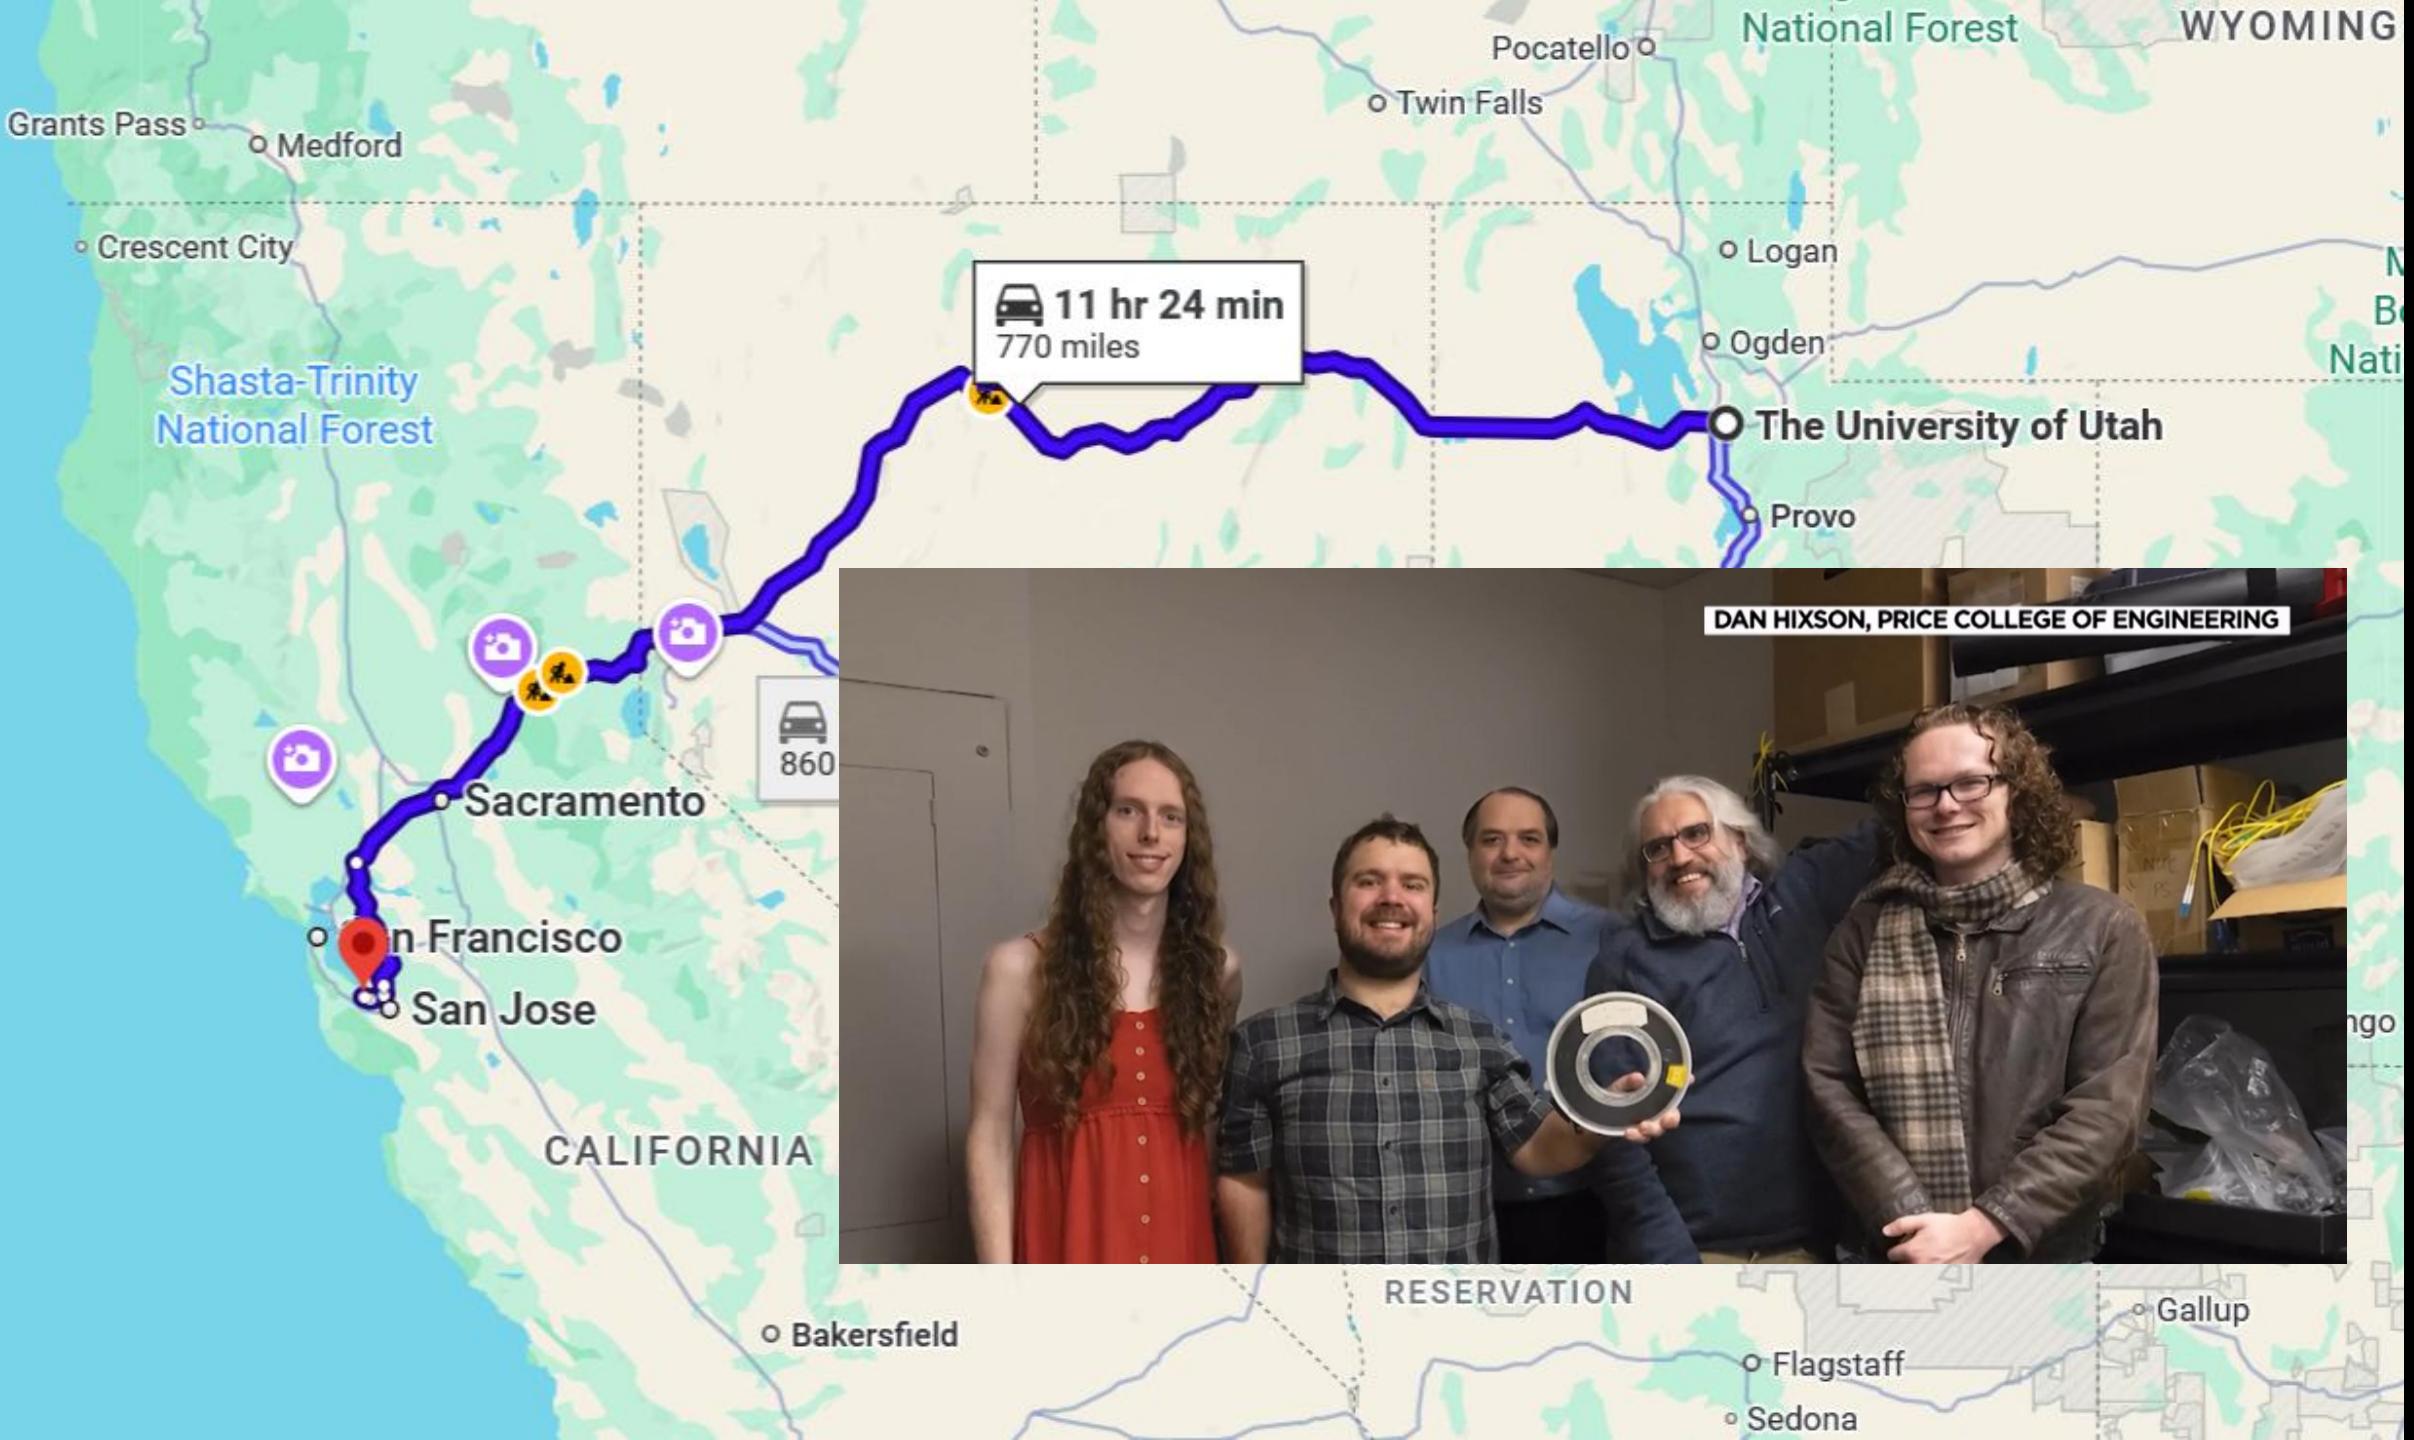




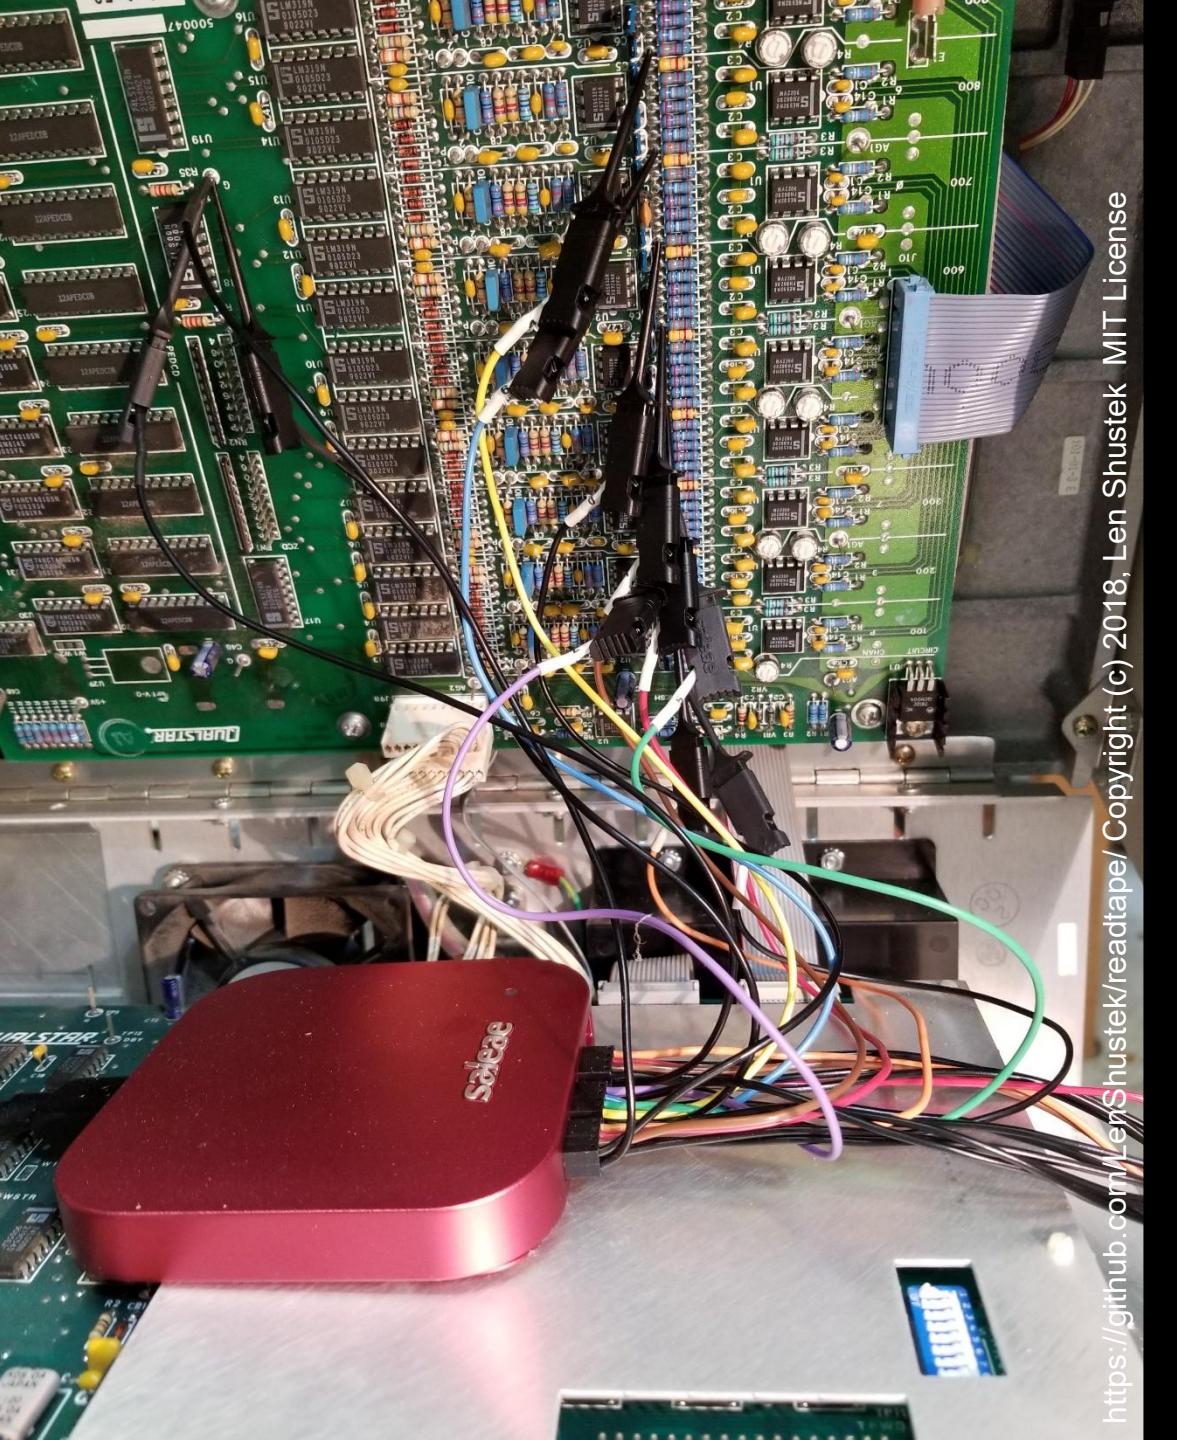


Photo Credit: Dan Hixson



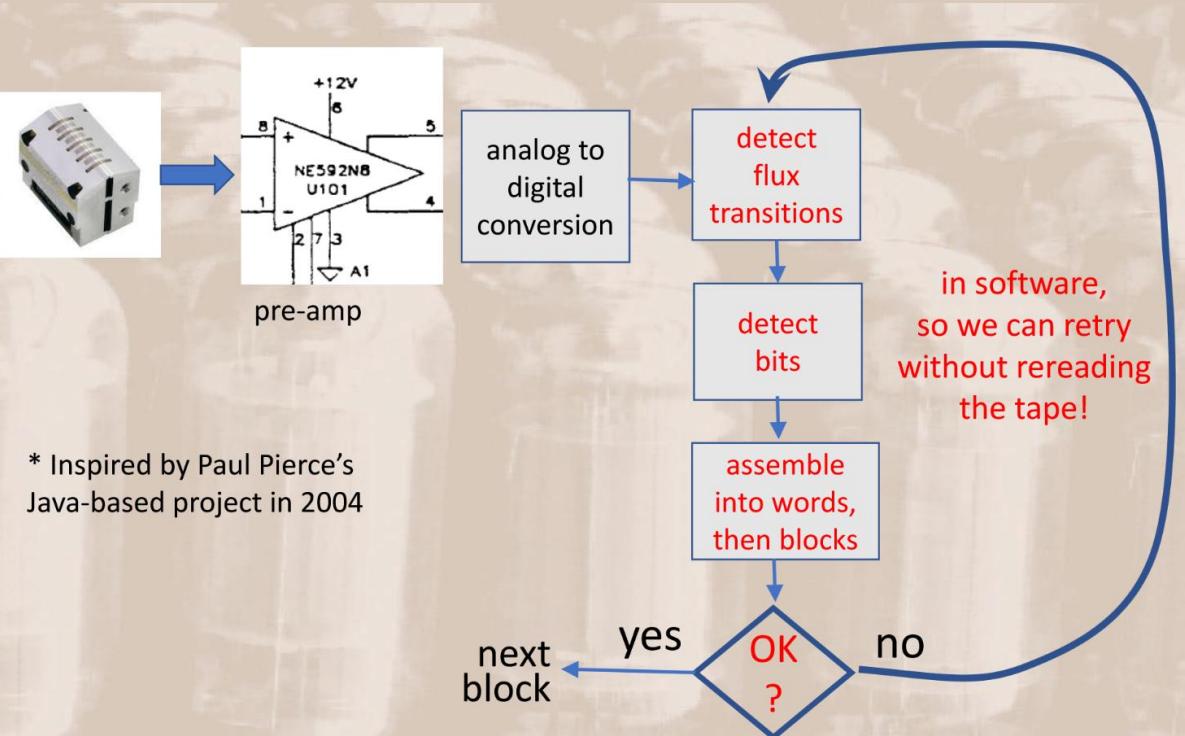
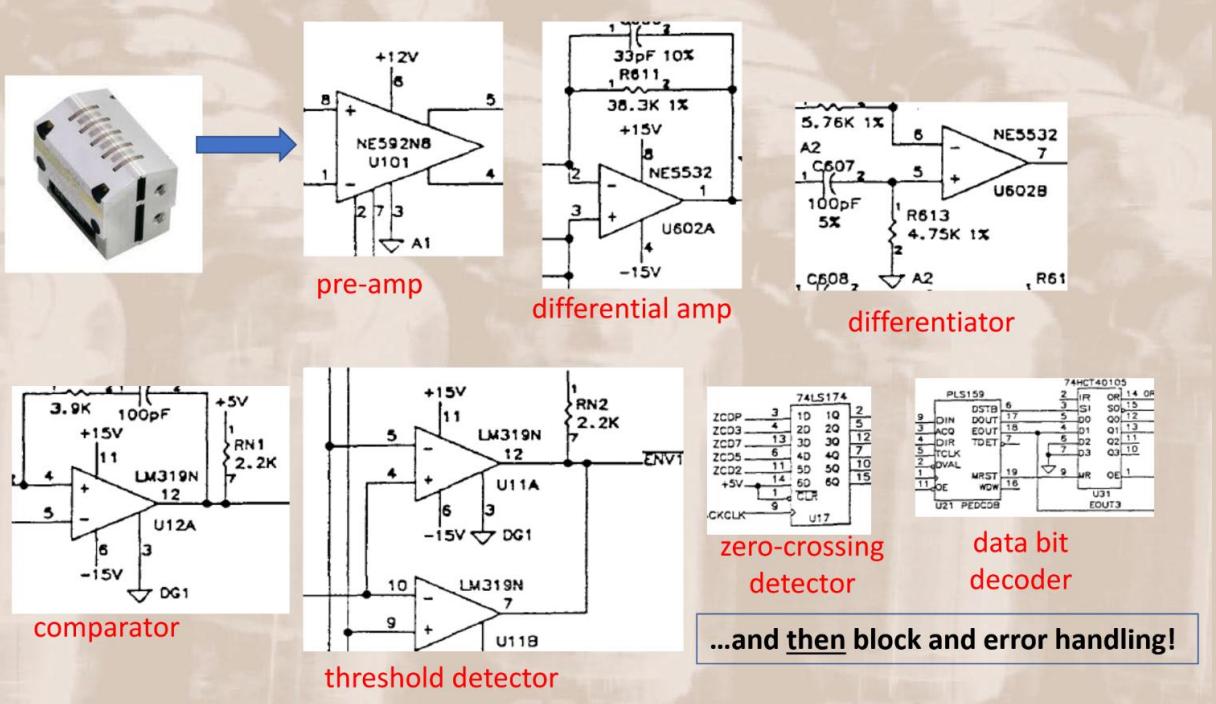






[https://github.com/\\_er/Shustek/readtape/](https://github.com/_er/Shustek/readtape/) Copyright (c) 2018, Len Shustek MIT License





\* Inspired by Paul Pierce's Java-based project in 2004

trk 0

1

2

trk 1

trk 2

trk 3

trk 4

trk 5

trk 6

L 25.52794 s, -201.63152 ms  
R 26.19322 s, 463.64464 ms  
1: 25.72957 s, 0  
2: 25.82500 s, 95.42454 ms  
3: xxx  
4: xxx  
5: xxx  
6: xxx  
7: xxx  
8: xxx  
9: xxx

trk 0

trk 1

trk 2

trk 3

trk 4

trk 5

trk 6

5

-4.746181

L 23.81328 s, 0

R 23.81386 s, 572.16000 us

1: xxx

2: xxx

3: xxx

4: xxx

5: 23.81358 s, 293.96228 us

6: xxx

7: xxx

8: xxx

9: xxx

Modified	Size	Filename
2002-01-23	12.0K	LICENSE-CALDERA.pdf
2002-01-23 (transcription)	2.8K	LICENSE-CALDERA.txt
2025-12-19 11:04:36 -0800	1.1G	v4/v4.sal
2025-12-19 11:14:40 -0800	1.6G	v4/analog.tbin
2025-12-19 11:14:40 -0800	540B	v4/analog.csvtbin.log
2025-12-19 11:23:08 -0800	3.6K	v4/analog.peakstats_deskew.csv
2025-12-19 11:26:36 -0800	2.5M	v4/analog.tap
2025-12-19 11:26:36 -0800	5.7K	v4/analog.log
2025-12-19 11:26:36 -0800	3.7K	v4/analog.peakstats.csv



**Obverse label** (handwritten, probably by Jay Lepreau):

UNIX Original  
from Bell Labs V4  
(See Manual for format)

**Obverse label** (printed, color):

Scotch®

BRAND

700

GP

3200 FCI

[https://archive.org/details/utah\\_unix\\_v4\\_raw](https://archive.org/details/utah_unix_v4_raw)



240 West Center Street  
Orem, Utah 84057  
801-765-4999 Fax 801-765-4481

January 23, 2002

Dear UNIX® enthusiasts,

Caldera International, Inc. hereby grants a fee free license that includes the rights use, modify and distribute this named source code, including creating derived binary products created from the source code. The source code for which Caldera International, Inc. grants rights are limited to the following UNIX Operating Systems that operate on the 16-Bit PDP-11 CPU and early versions of the 32-Bit UNIX Operating System, with specific exclusion of UNIX System III and UNIX System V and successor operating systems:

32-bit 32V UNIX  
16 bit UNIX Versions 1, 2, 3, 4, 5, 6, 7

Caldera International, Inc. makes no guarantees or commitments that any source code is available from Caldera International, Inc.

The following copyright notice applies to the source code files for which this license is granted.

Copyright(C) Caldera International Inc. 2001-2002. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

Redistributions of source code and documentation must retain the above copyright notice, this list of conditions and the following disclaimer. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

All advertising materials mentioning features or use of this software must display the following acknowledgement:

This product includes software developed or owned by Caldera International, Inc.

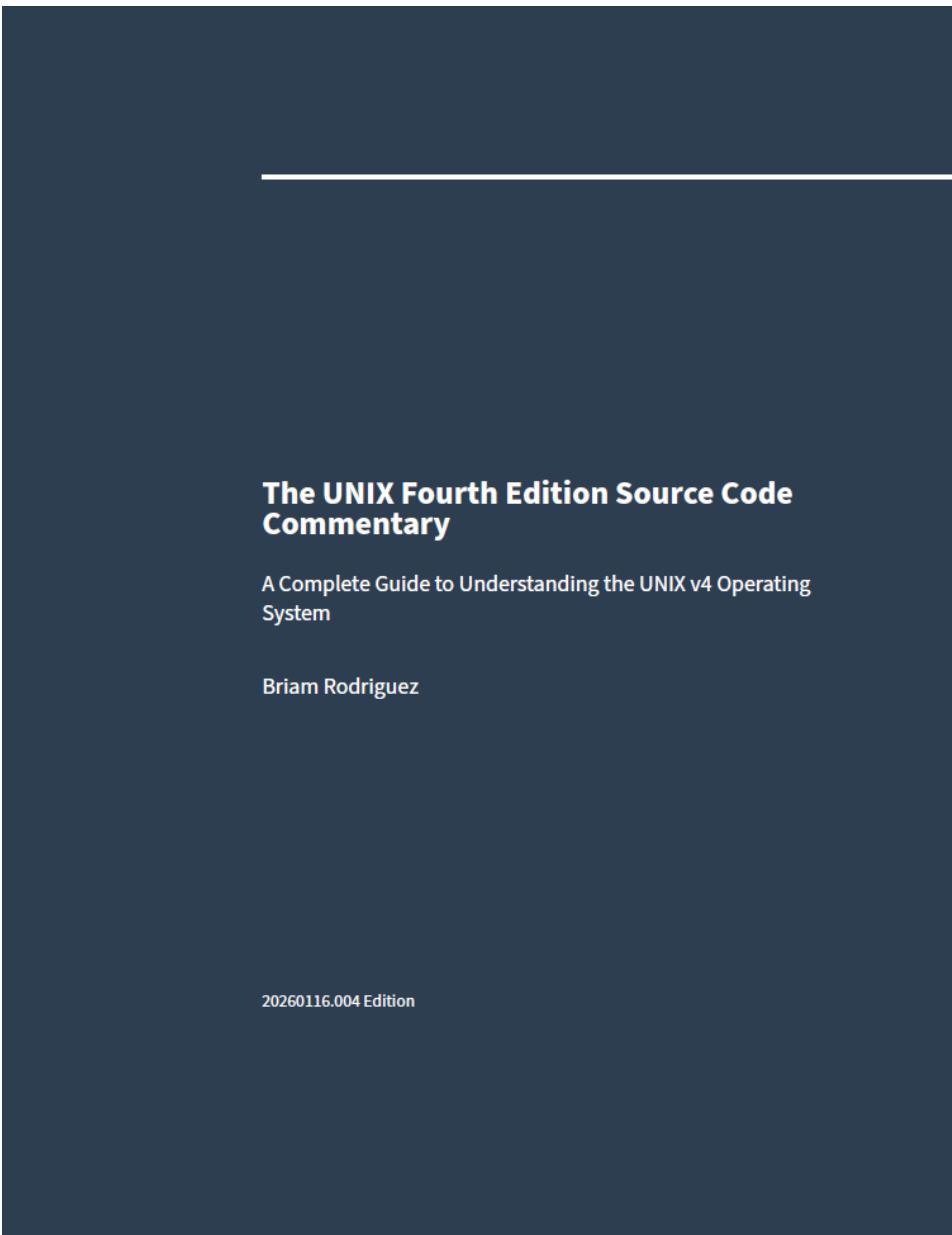
Neither the name of Caldera International, Inc. nor the names of other contributors may be used to endorse or promote products derived from this software without specific prior written permission.

USE OF THE SOFTWARE PROVIDED FOR UNDER THIS LICENSE BY CALDERA INTERNATIONAL, INC. AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL CALDERA INTERNATIONAL, INC. BE LIABLE FOR ANY DIRECT, INDIRECT INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Very truly yours,

/signed/ Bill Broderick

Bill Broderick  
Director, Licensing Services



Contents

20260116.004 Edition

## Contents

<b>The UNIX Fourth Edition Source Code Commentary</b>	2
About This Book . . . . .	2
How to Use This Book . . . . .	7
<b>I Foundation</b>	<b>8</b>
<b>1 Chapter 1: Introduction</b>	<b>9</b>
1.1 Overview . . . . .	9
1.2 Prerequisites . . . . .	9
1.3 The Birth of UNIX . . . . .	9
1.4 The Bell Labs Environment . . . . .	11
1.5 Design Philosophy . . . . .	12
1.6 The Cast of Characters . . . . .	13
1.7 What We'll Study . . . . .	14
1.8 Summary . . . . .	15
1.9 Further Reading . . . . .	15
<b>2 Chapter 2: The PDP-11 Architecture</b>	<b>16</b>
2.1 Overview . . . . .	16
2.2 Source Files . . . . .	16
2.3 Prerequisites . . . . .	16
2.4 The PDP-11 Family . . . . .	16
2.5 Registers . . . . .	17
2.6 The Processor Status Word (PS) . . . . .	18
2.7 Memory Layout . . . . .	19
2.8 The Memory Management Unit (MMU) . . . . .	19
2.9 Trap and Interrupt Mechanism . . . . .	21
2.10 Key Machine Instructions . . . . .	22
2.11 Addressing Modes . . . . .	24
2.12 The User Structure Address . . . . .	25
2.13 Key Machine-Dependent Functions . . . . .	25
2.14 Summary . . . . .	26



Sixth  
Research  
Edition  
(May 1975)

A  
**COMMENTARY  
ON THE  
UNIX  
OPERATING  
SYSTEM**

JOHN LIONS



The University of New South Wales



39c3



[github.com/dspinellis/unix-history-repo/](https://github.com/dspinellis/unix-history-repo/)



Bell Laboratories

BSD 1  
Unix 32/V

BSD 2

BSD 3



BSD SCCS

BSD 4.0  
BSD 4.1

BSD 4.2

BSD 4.3

BSD 4.3-Tahoe

BSD 4.3-Reno

BSD 4.4

BSD 4.3 Net/1

BSD 4.3 Net/2

BSD 4.4 Lite1



386BSD 0.0  
386BSD 0.1

386BSD patch kit



FreeBSD

FreeBSD Git

VCS import

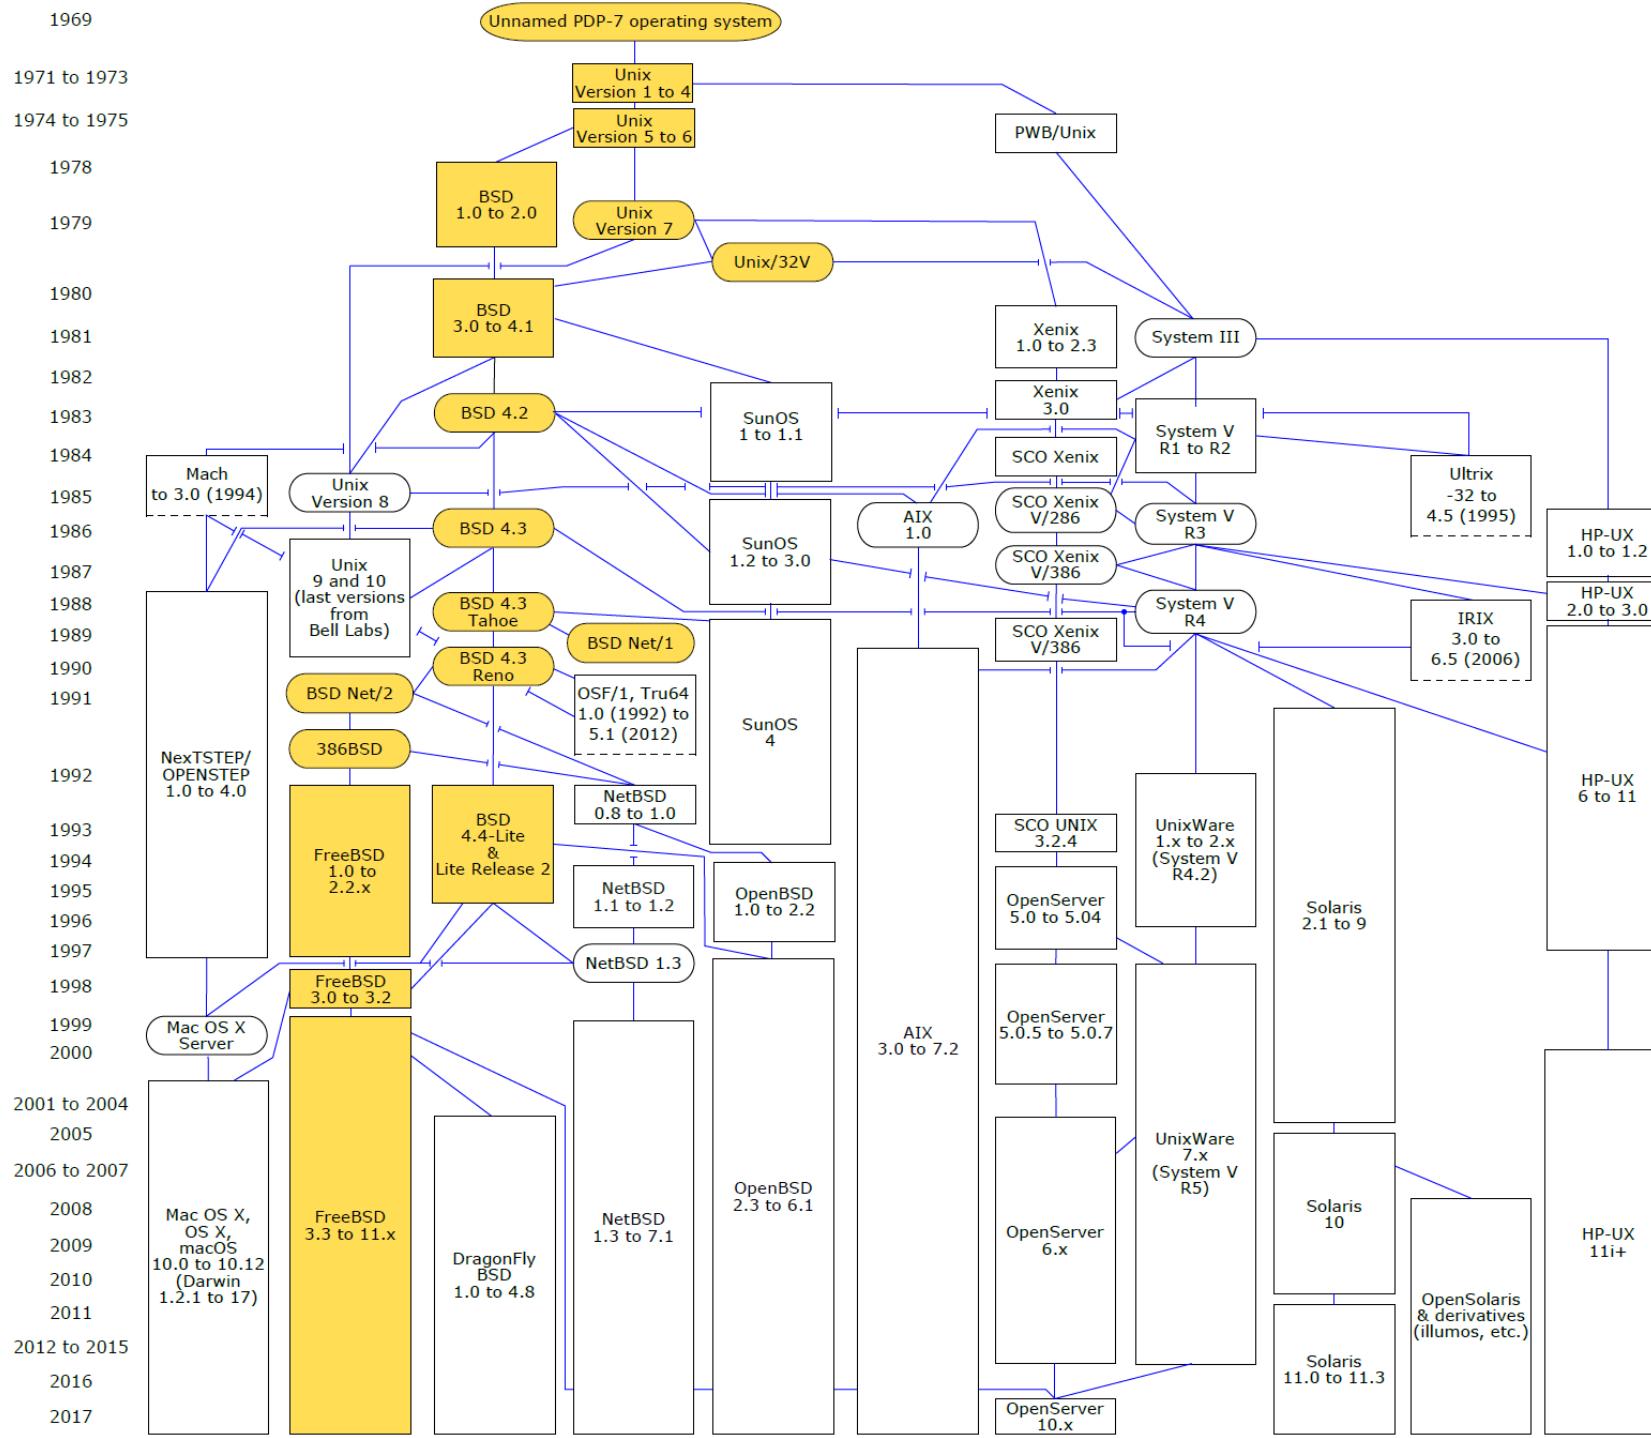
Snapshot import

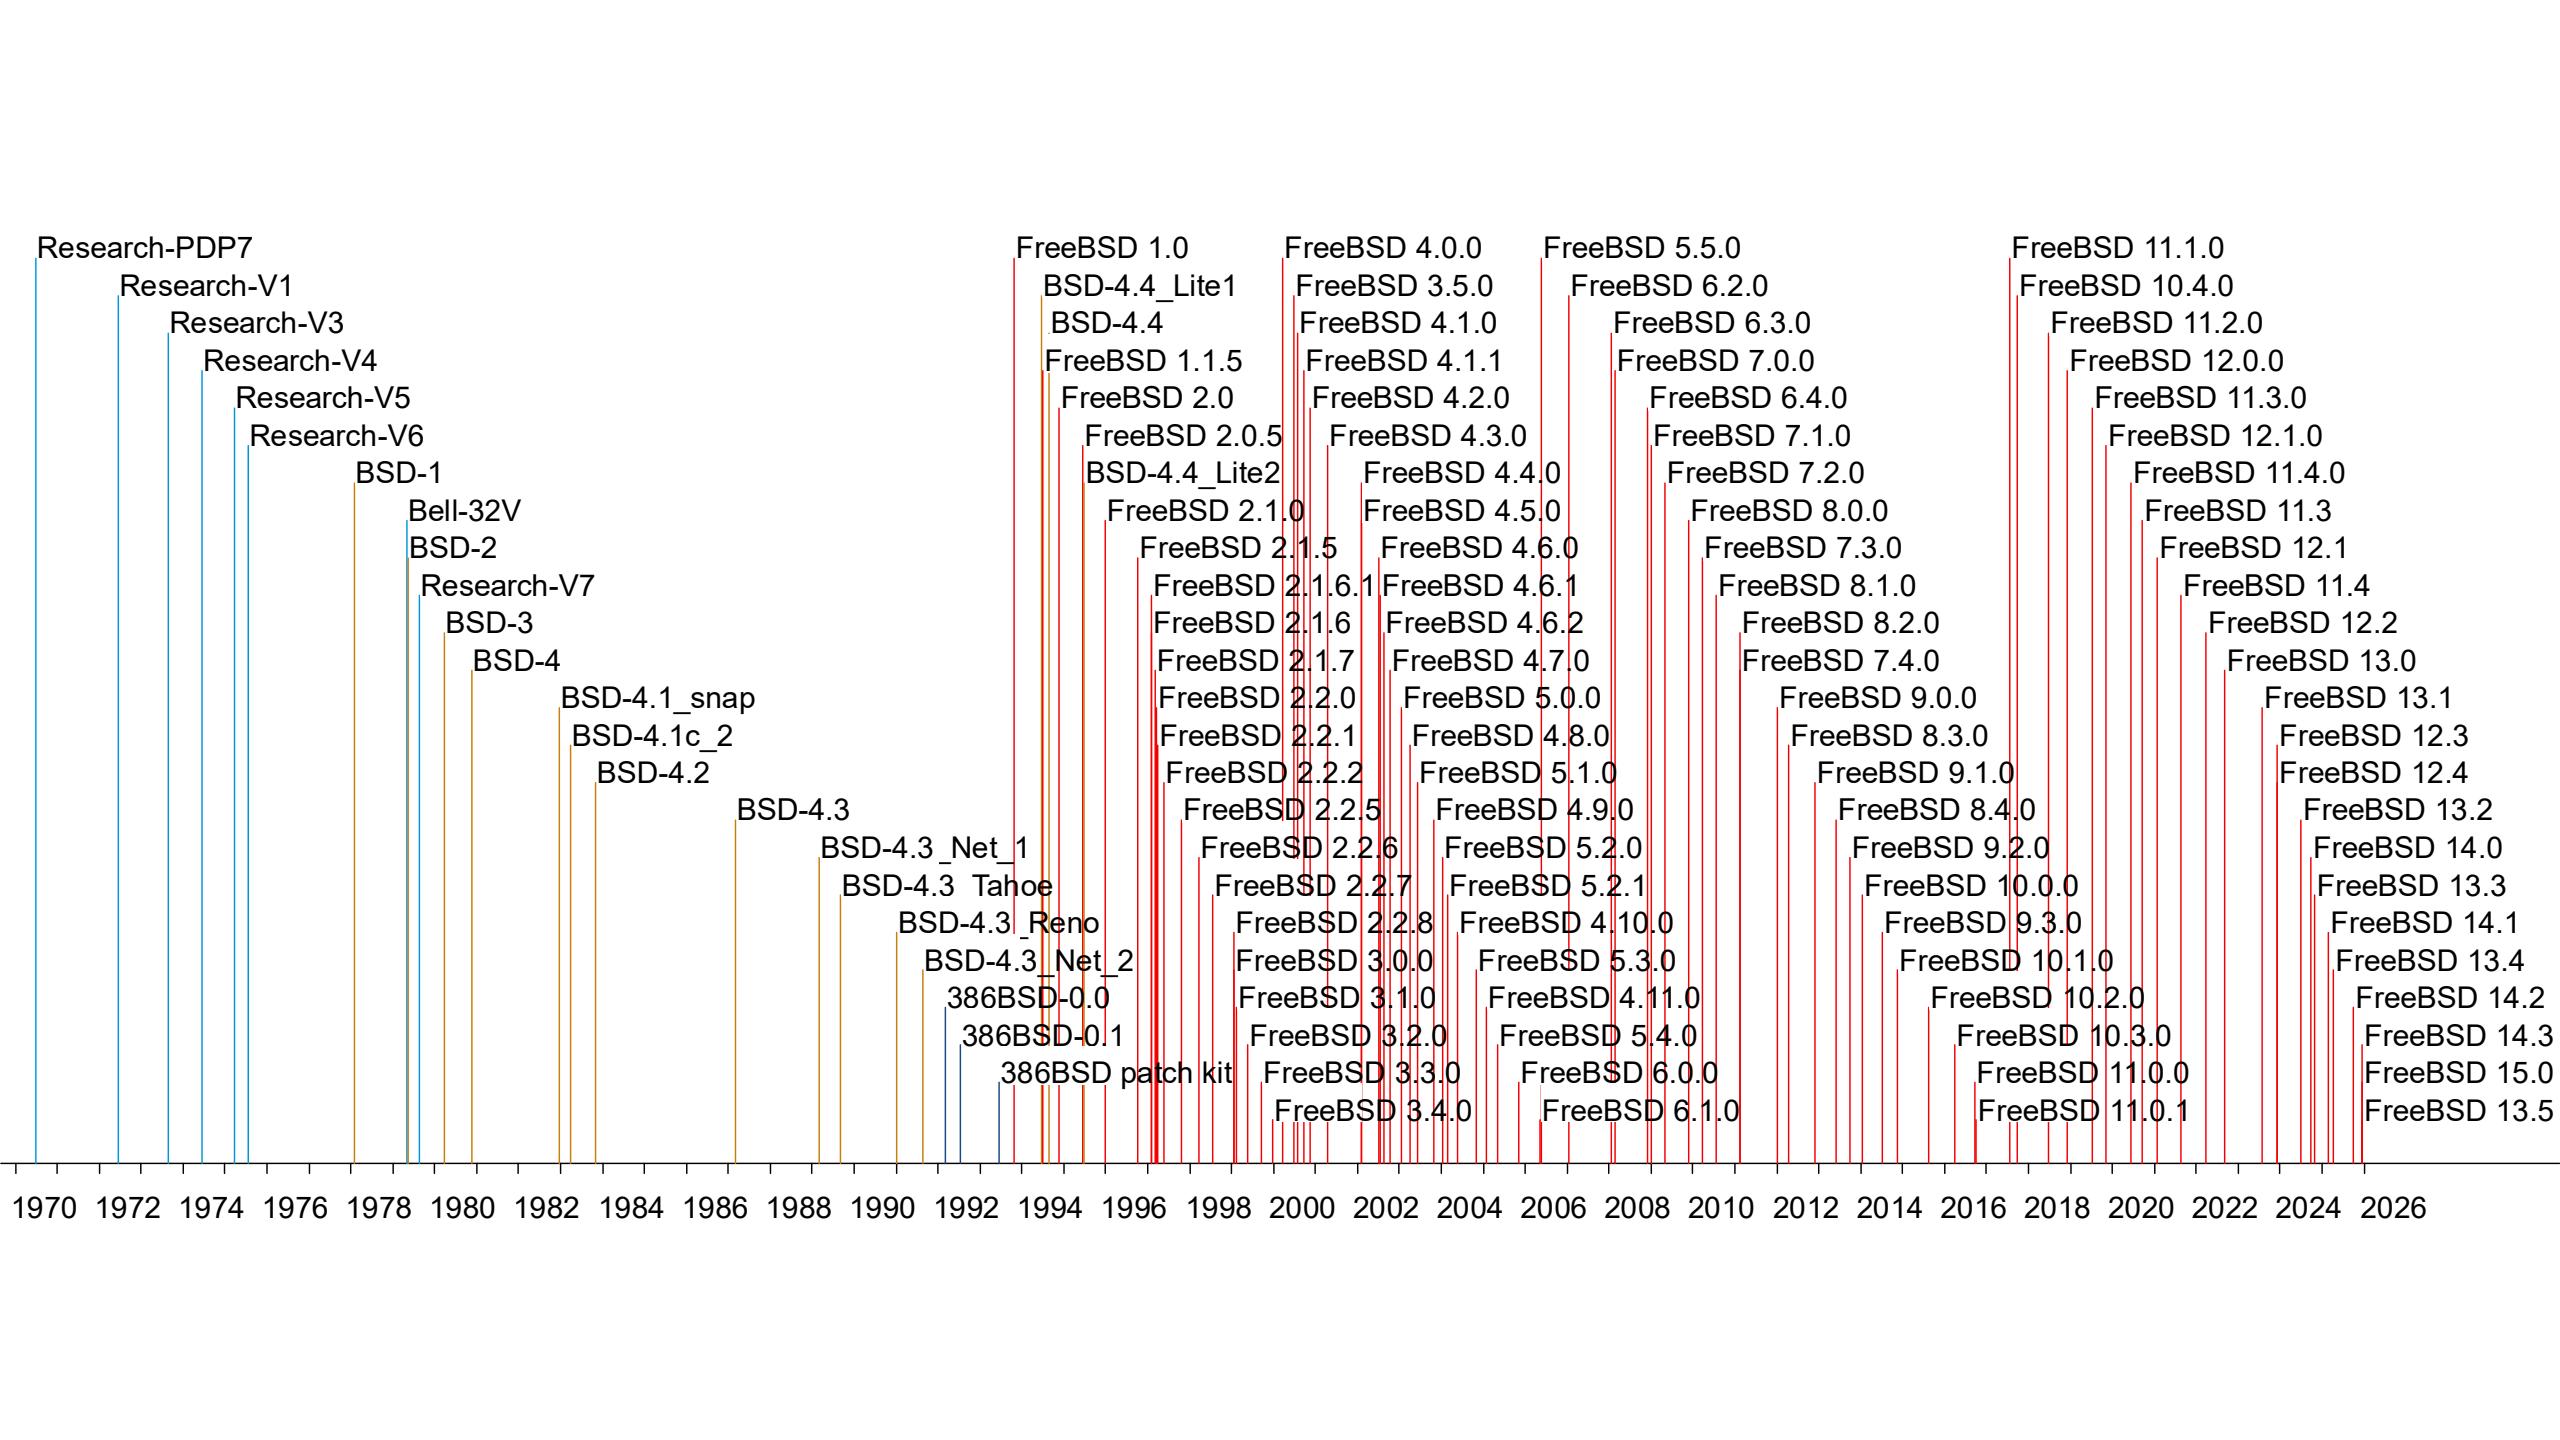
Bell Labs

Berkeley

Open source  
(including Berkeley)







```
$ git checkout FreeBSD-release/10.0.0
$ git blame -M -M -C -C ./lib/libc/gen/timezone.c
```

```
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 76) static struct zone {
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 77)     int offset;
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 78)     char *stdzone;
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 79)     char *dlzone;
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 80) } zonetab[] = {
lib/libc/gen/timezone.c (Jordan K. Hubbard 1996-07-12 18:57:58 +0000 81)     {-1*60, "MET", "MET DST"}, [...]
lib/libc/gen/timezone.c (Jordan K. Hubbard 1996-07-12 18:57:58 +0000 96)     {-1}
usr/src/lib/libc/gen/timezone.c (Bill Joy 1980-12-22 00:40:25 -0800 97) };
usr/src/lib/libc/gen/timezone.c (Bill Joy 1980-12-22 00:40:25 -0800 98)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 106) char *
lib/libc/gen/timezone.c (Ed Schouten 2009-12-05 19:31:38 +0000 107) _tztab(int zone, int dst)
lib/libc/gen/timezone.c (Rodney Grimes 1994-05-27 05:00:24 +0000 108) {
lib/libc/gen/timezone.c (David E. O'Brien 2002-02-01 01:08:48 +0000 109)     struct zone *zp;
lib/libc/gen/timezone.c (David E. O'Brien 2002-02-01 01:08:48 +0000 110)     char sign;
usr/src/lib/libc/gen/timezone.c (Bill Joy 1980-12-22 00:40:25 -0800 111)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 112)
-1; ++zp) /* static tables */
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 113)
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 114)
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 115)
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 116)
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 117)
usr/src/libc/gen/timezone.c (Dennis Ritchie 1979-01-10 14:58:45 -0500 118)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 119)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 120)
/* create one */
usr/src/lib/libc/gen/timezone.c (Bill Joy 1980-12-22 00:40:25 -0800 121)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 122)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 123)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 124)
usr/src/lib/libc/gen/timezone.c (Keith Bostic 1987-03-28 19:27:07 -0800 125)
lib/libc/gen/timezone.c (Warner Losh 1998-01-21 21:46:36 +0000 126)
sizeof(czone),
lib/libc/gen/timezone.c (Warner Losh 1998-01-21 21:46:36 +0000 127)
60, zone % 60);
lib/libc/gen/timezone.c (Rodney Grimes 1994-05-27 05:00:24 +0000 128)
lib/libc/gen/timezone.c (Rodney Grimes 1994-05-27 05:00:24 +0000 129) }
```

usr/source/s1/form.\.s rhm,llc  
usr/source/s1/getty\.s dmr,ken,jfo  
usr/source/s1/glob\.c dmr  
usr/source/s1/goto\.c dmr  
usr/source/s1/grep\.s ken,lem  
usr/source/s1/ld.\.s dmr  
usr/source/s2/mail\.c ken  
usr/source/s2/sa\.c ken  
usr/source/s2/sh\.c ken  
usr/source/s2/sort\.c ken,doug  
usr/source/s2/strip\.s dmr  
usr/source/s2/sum\.s ken  
usr/source/s2/tee\.c doug  
usr/source/s2/tr\.c doug  
usr/source/s2/uniq\.c ken  
usr/source/s2/wc\.c jfo

# Demo time!

<https://unixv4.dev/>



# First available **complete** Unix distribution

	v1	v2	v3	v4
Released	Nov 71	Jun 72	Feb 73	Nov 73
Kernel	✓		✓	✓
Commands	●	●		✓
Man pages			✓	✓
C compiler		✓	✓	✓

# Structured Programming

- Kernel implemented in “New B”
- 7883 lines New B
- 1054 lines PDP-11 assembly

```
main()
{
    extern schar;
    extern char end[], data[], etext[];
    int i, i1, *p;

    /*
     * zero and free all of core
     */

    UIISA->r[0] = KISA->r[6] + USIZE;
    UISD->r[0] = 6;
    for(; fubyte(0) >= 0; UIISA->r[0]++) {
        clearseg(UIISA->r[0]);
        mfree(coremap, 1, UIISA->r[0]);
    }
    mfree(swapmap, NSWAP, SWPLO);

    /*
     * set up system process
     */

    proc[0].p_addr = KISA->r[6];
    proc[0].p_size = USIZE;
    proc[0].p_stat = SRUN;
    proc[0].p_flag = SLOAD|SSYS;
    u.u_procp = &proc[0];

    /*
     * set up 'known' i-nodes
     */

    sureg();
    LKS->integ = 0115;
    cinit();
    binit();
    iinit();
    rootdir = igin(ROOTDEV, ROOTINO);
    rootdir->i_flag = ~ILOCK;
    u.u_cdir = igin(ROOTDEV, ROOTINO);
    u.u_cdir->i_flag = ~ILOCK;
```

# Language-Independent API

PIPE ( II )

8/5/73

PIPE ( II )

## NAME

pipe – create a pipe

## SYNOPSIS

```
(pipe = 42.)  
sys pipe  
(read file descriptor in r0)  
(write file descriptor in r1)  
  
pipe(fildes)  
int fildes[2];
```

## DESCRIPTION

The *pipe* system call creates an I/O mechanism called a pipe. The file descriptors returned can be used in read and write operations. When the pipe is written using the descriptor returned in r1 (resp. fildes[1]), up to 4096 bytes of data are buffered before the writing process is suspended. A read using the descriptor returned in r0 (resp. fildes[0]) will pick up the data.

It is assumed that after the pipe has been set up, two (or more) cooperating processes (created by subsequent *fork* calls) will pass data through the pipe with *read* and *write* calls.

The shell has a syntax to set up a linear array of processes connected by pipes.

Read calls on an empty pipe (no buffered data) with only one end (all write file descriptors closed) return an end-of-file. Write calls under similar conditions are ignored.

# Data Structure Definition & Reuse

**buf.h    file.h    inode.h    proc.h    seg.h    text.h    user.h**  
**conf.h    filsys.h    param.h    reg.h    systm.h    tty.h**

# Device Driver Abstraction

#### IV. SPECIAL FILES

# Driver Interface

```
struct {
    int      (*d_open)();
    int      (*d_close)();
    int      (*d_strategy)();
    int      *d_tab;
} bdevsw[];
```

```
struct {
    int      (*d_open)();
    int      (*d_close)();
    int      (*d_read)();
    int      (*d_write)();
    int      (*d_sgtty)();
} cdevsw[];
```

UNIX PROGRAMMER'S MANUAL

Third Edition

K. Thompson

D. M. Ritchie

February, 1973

Copyright 8c9 1972  
Bell Telephone Laboratories, Inc.

No part of this document may be reproduced,  
or distributed outside the Laboratories, without  
the written permission of Bell Telephone Laboratories.

**UNIX PROGRAMMER'S MANUAL**

*Fourth Edition*

*K. Thompson*

*D. M. Ritchie*

November, 1973

Copyright © 1972, 1973  
Bell Telephone Laboratories, Inc.

No part of this document may be reproduced,  
or distributed outside the Laboratories, without  
the written permission of Bell Telephone Laboratories.

**NAME**

sno – Snobol interpreter

**SYNOPSIS**

**sno [ file ]**

**DESCRIPTION**

*Sno* is a Snobol III (with slight differences) compiler and interpreter. *Sno* obtains input from the concatenation of *file* and the standard input. All input through a statement containing the label ‘end’ is considered program and is compiled. The rest is available to ‘syspt’.

*Sno* differs from Snobol III in the following ways.

There are no unanchored searches. To get the same effect:

a ** b	unanchored search for b
a *x* b = x c	unanchored assignment

There is no back referencing.

x = "abc"	
a *x* x	is an unanchored search for 'abc'

Function declaration is different. The function declaration is done at compile time by the use of the label ‘define’. Thus there is no ability to define functions at run time and the use of the name ‘define’ is preempted. There is also no provision for automatic variables other than the parameters. For example:

**definef()**

or

**define f(a,b,c)**

All labels except ‘define’ (even ‘end’) must have a non-empty statement.

If ‘start’ is a label in the program, program execution will start there. If not, execution begins with the first executable statement. ‘define’ is not an executable statement.

There are no builtin functions.

Parentheses for arithmetic are not needed. Normal precedence applies. Because of this, the arithmetic operators ‘/’ and ‘\*’ must be set off by space.

The right side of assignments must be non-empty.

Either ‘ ’ or “ ” may be used for literal quotes.

The pseudo-variable ‘syspt’ is not available.

**SEE ALSO**

Snobol III manual. (JACM; Vol. 11 No. 1; Jan 1964; pp 21)

**BUGS**

I think writing it was just  
a quick entertainment for  
Ken. The "application" that  
has survived is a 1-page  
program that solves the Soma  
(or Instant Insanity)  
puzzle.

Dennis

```

define remove
front */'1'* = "
back */'1'* = "
right */'1'* = "
left */'1'* = " /(return)

```

```

define test(s)
s */'2'* *ft/'1'* *bk/'1'* *lt/'1'* *rt/'1'* /f(xxx)
front ** ft /s(freturn)
back ** bk /s(freturn)
left ** lt /s(freturn)
right ** rt /s(freturn)
front = ft front
right = rt right
left = lt left
back = bk back /(return)

```

```

define build(s)
\$s */'1'* *2/'1'* *3/'1'* *4/'1'* *5/'1'* *6/'1'* /f(xxx)
/(s '1') = 1 2 3 4 5 6
/(s '2') = 1 2 6 5 3 4
/(s '3') = 1 2 4 3 6 5
/(s '4') = 1 2 5 6 4 3
/(s '5') = 3 4 1 2 6 5

```



```

[...]
$(s '24') = 6 5 3 4 1 2 /(return)

start a = 'wbrrrg'
b = 'wbbrrgg'
c = 'wbgrwg'
d = 'wgwrrb'

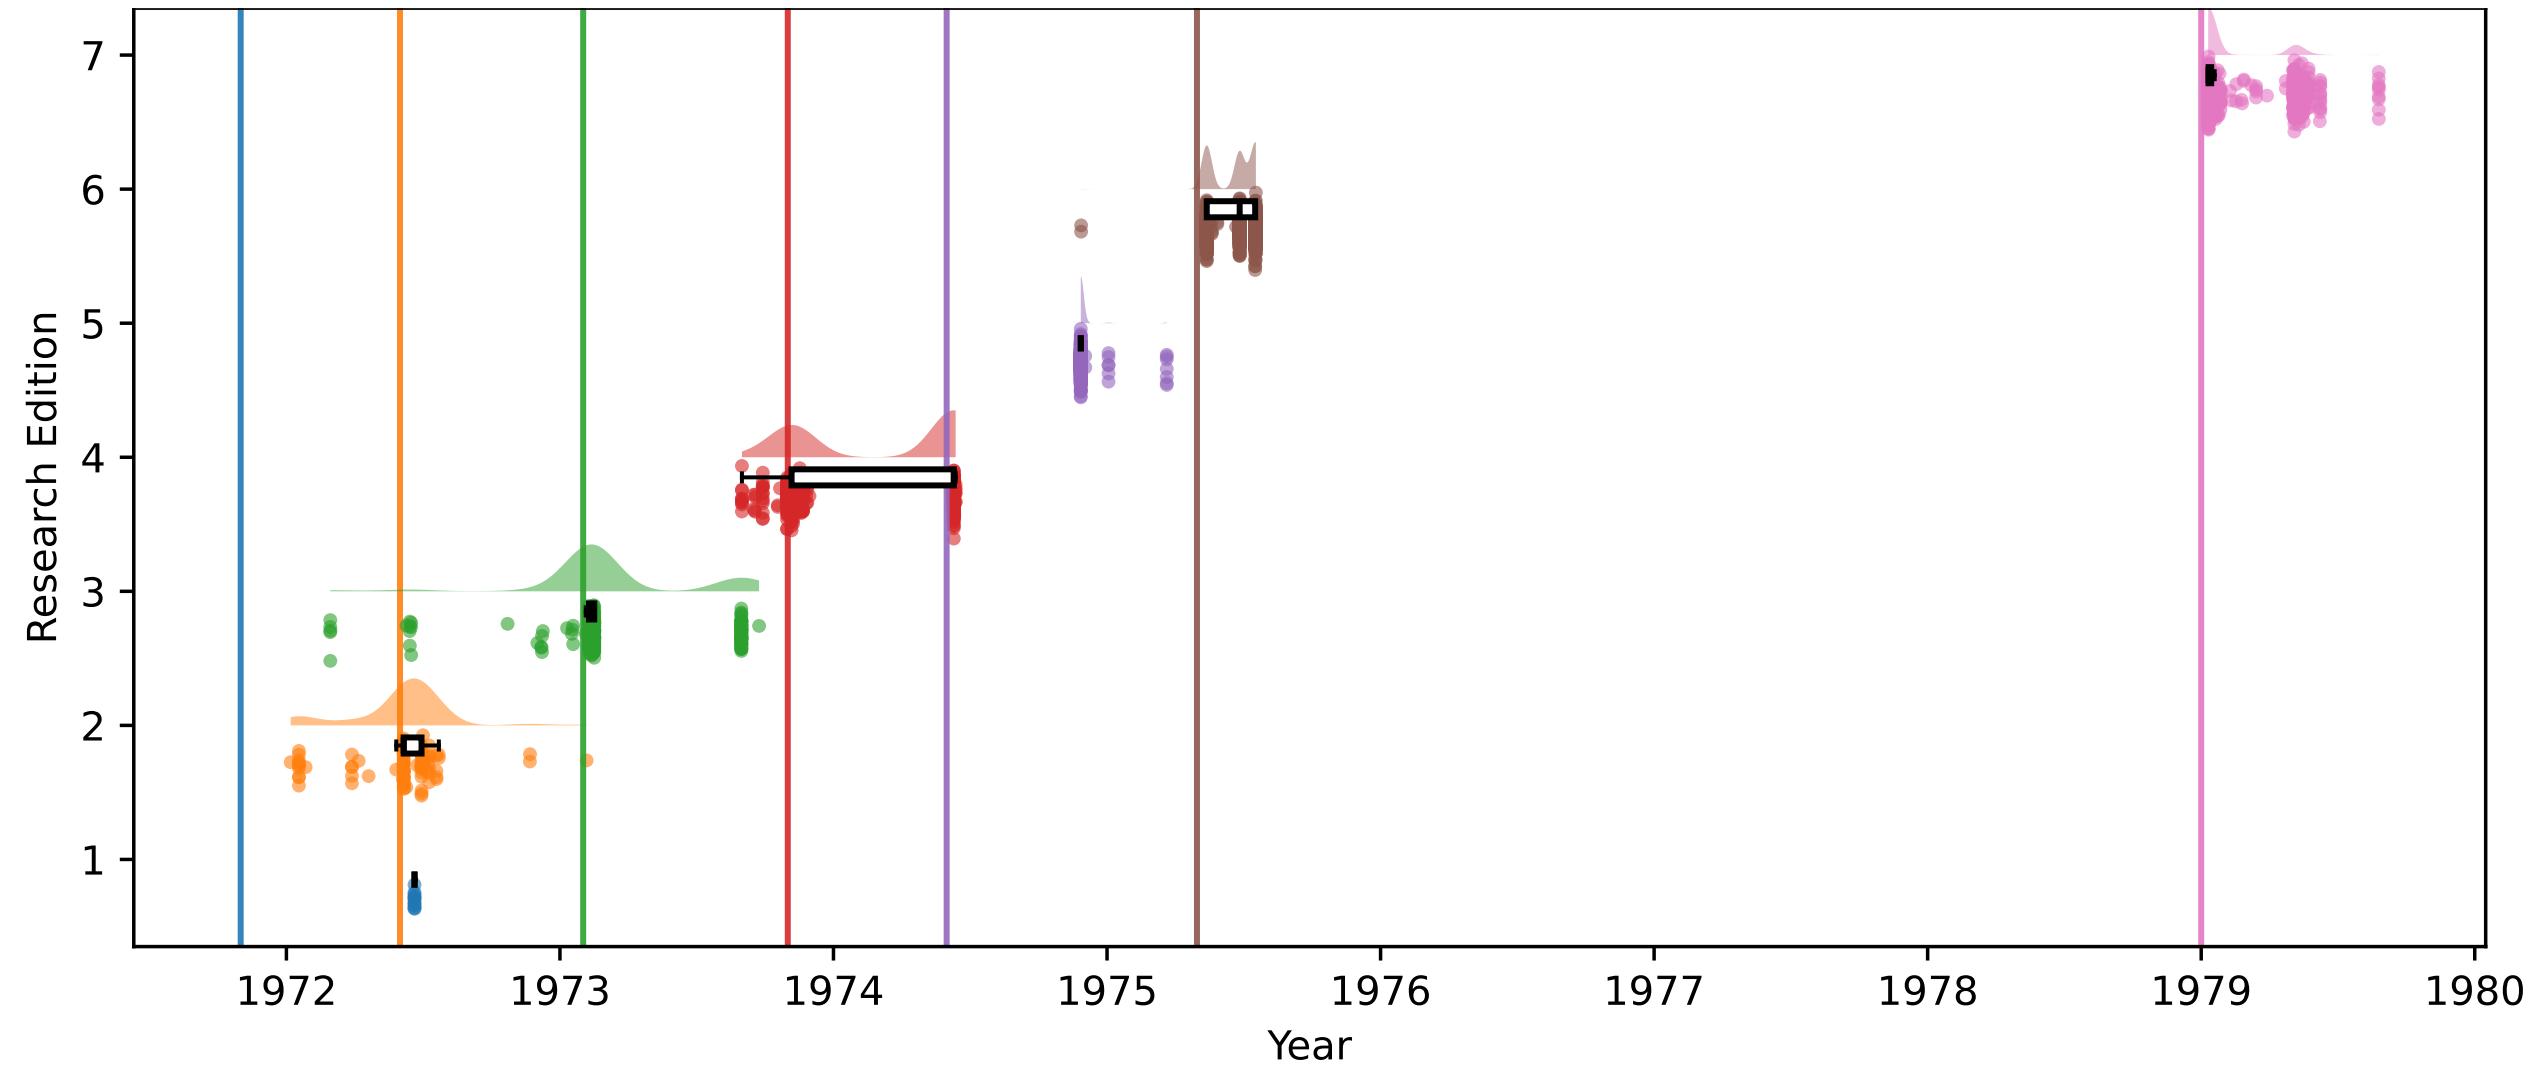
build('a')
build('b')
build('c')
build('d')

a = '1'
l1 x1 = $('a' a)
test(x1) /f(xxx)
b = '1'
l2 x2 = $('b' b)
test(x2) /f(t2)
c = '1'
l3 x3 = $('c' c)
test(x3) /f(t3)
d = '1'
l4 x4 = $('d' d)
test(x4) /f(t4)

syspot = x1
syspot = x2
syspot = x3
syspot = x4
remove()
t4 d = d + '1'
d '25' /f(l4)
remove()
t3 c = c + '1'
c '25' /f(l3)
remove()
t2 b = b + '1'
b '25' /f(l2)
remove()
t1 a = a + '4'
a '13' /f(l1)
end syspot = 'Done...'

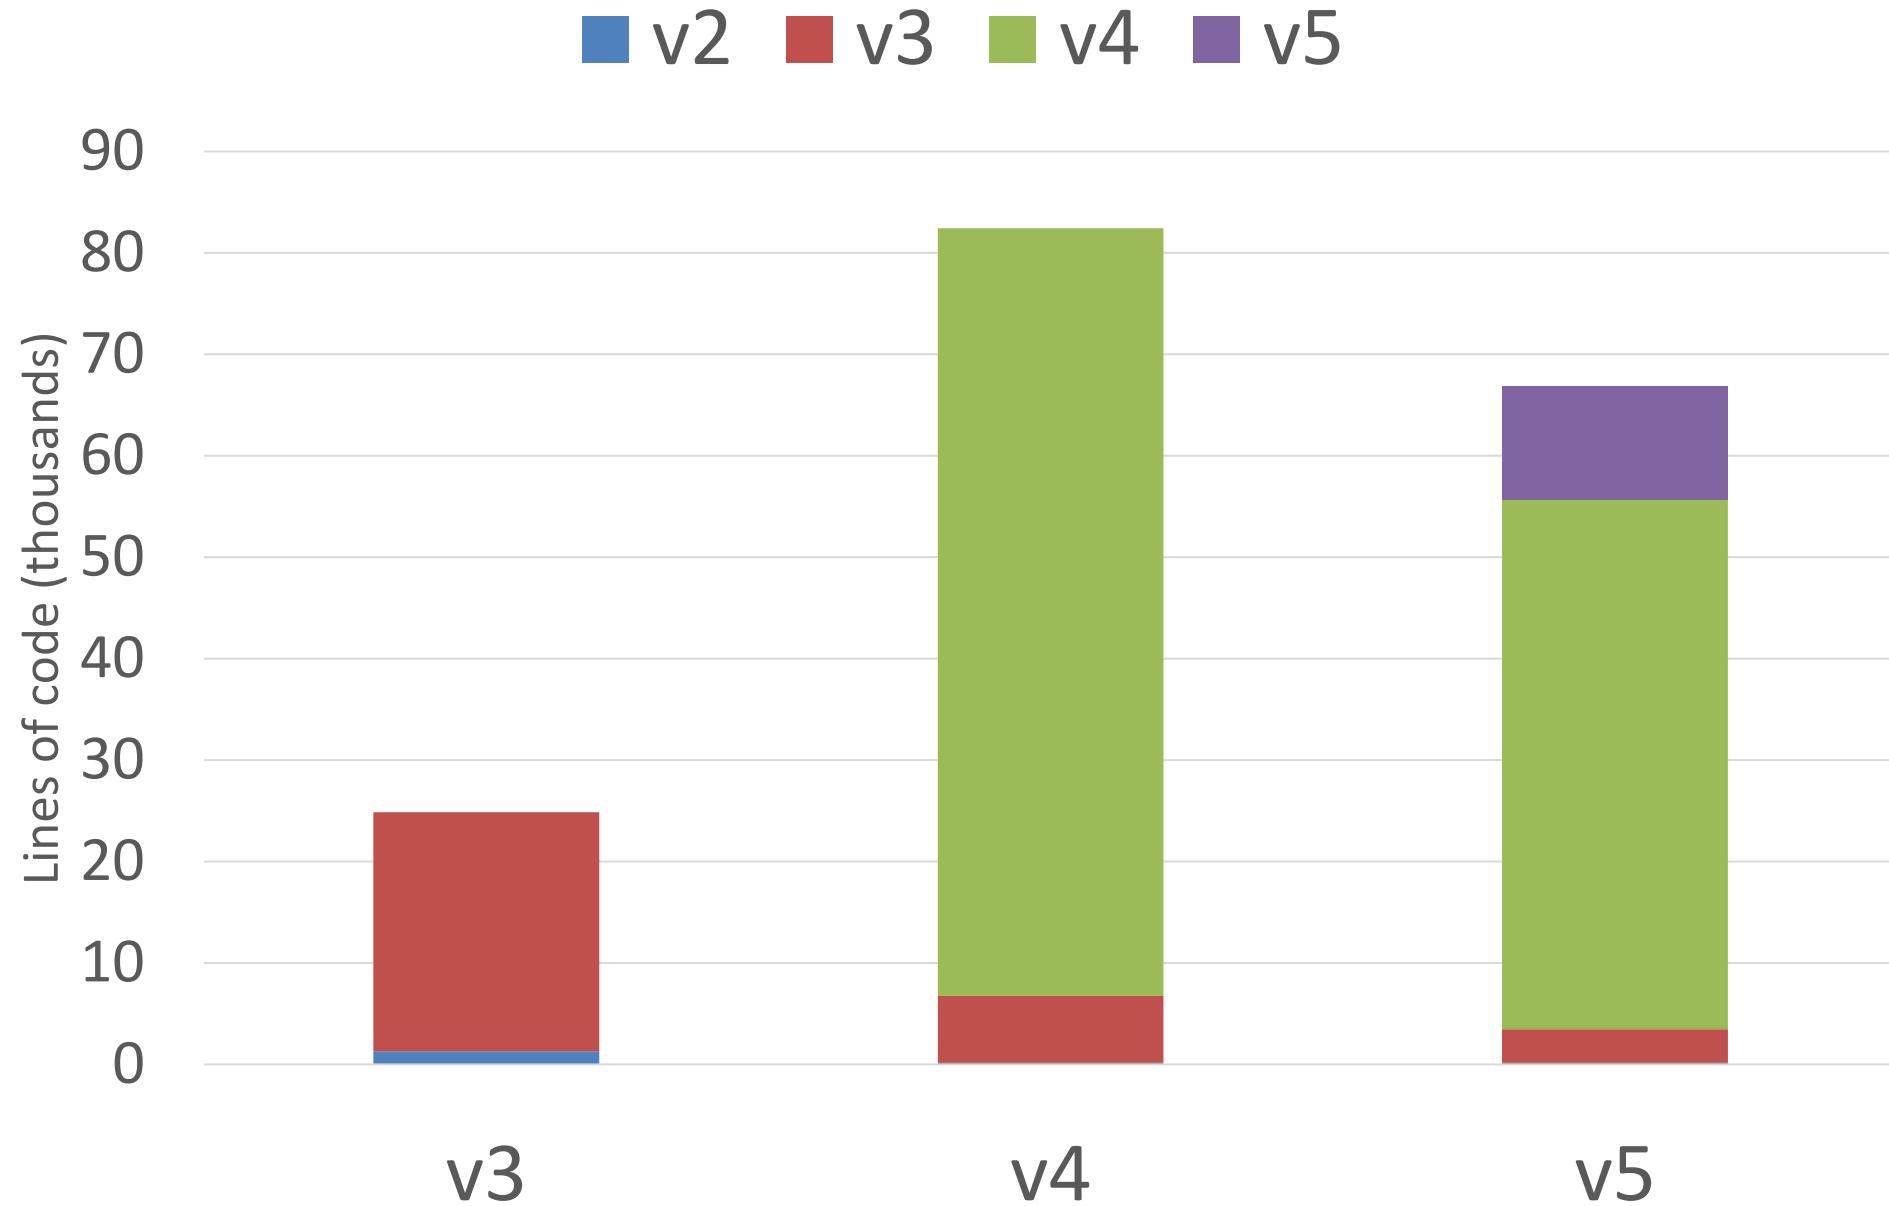
```

# Dating the tape



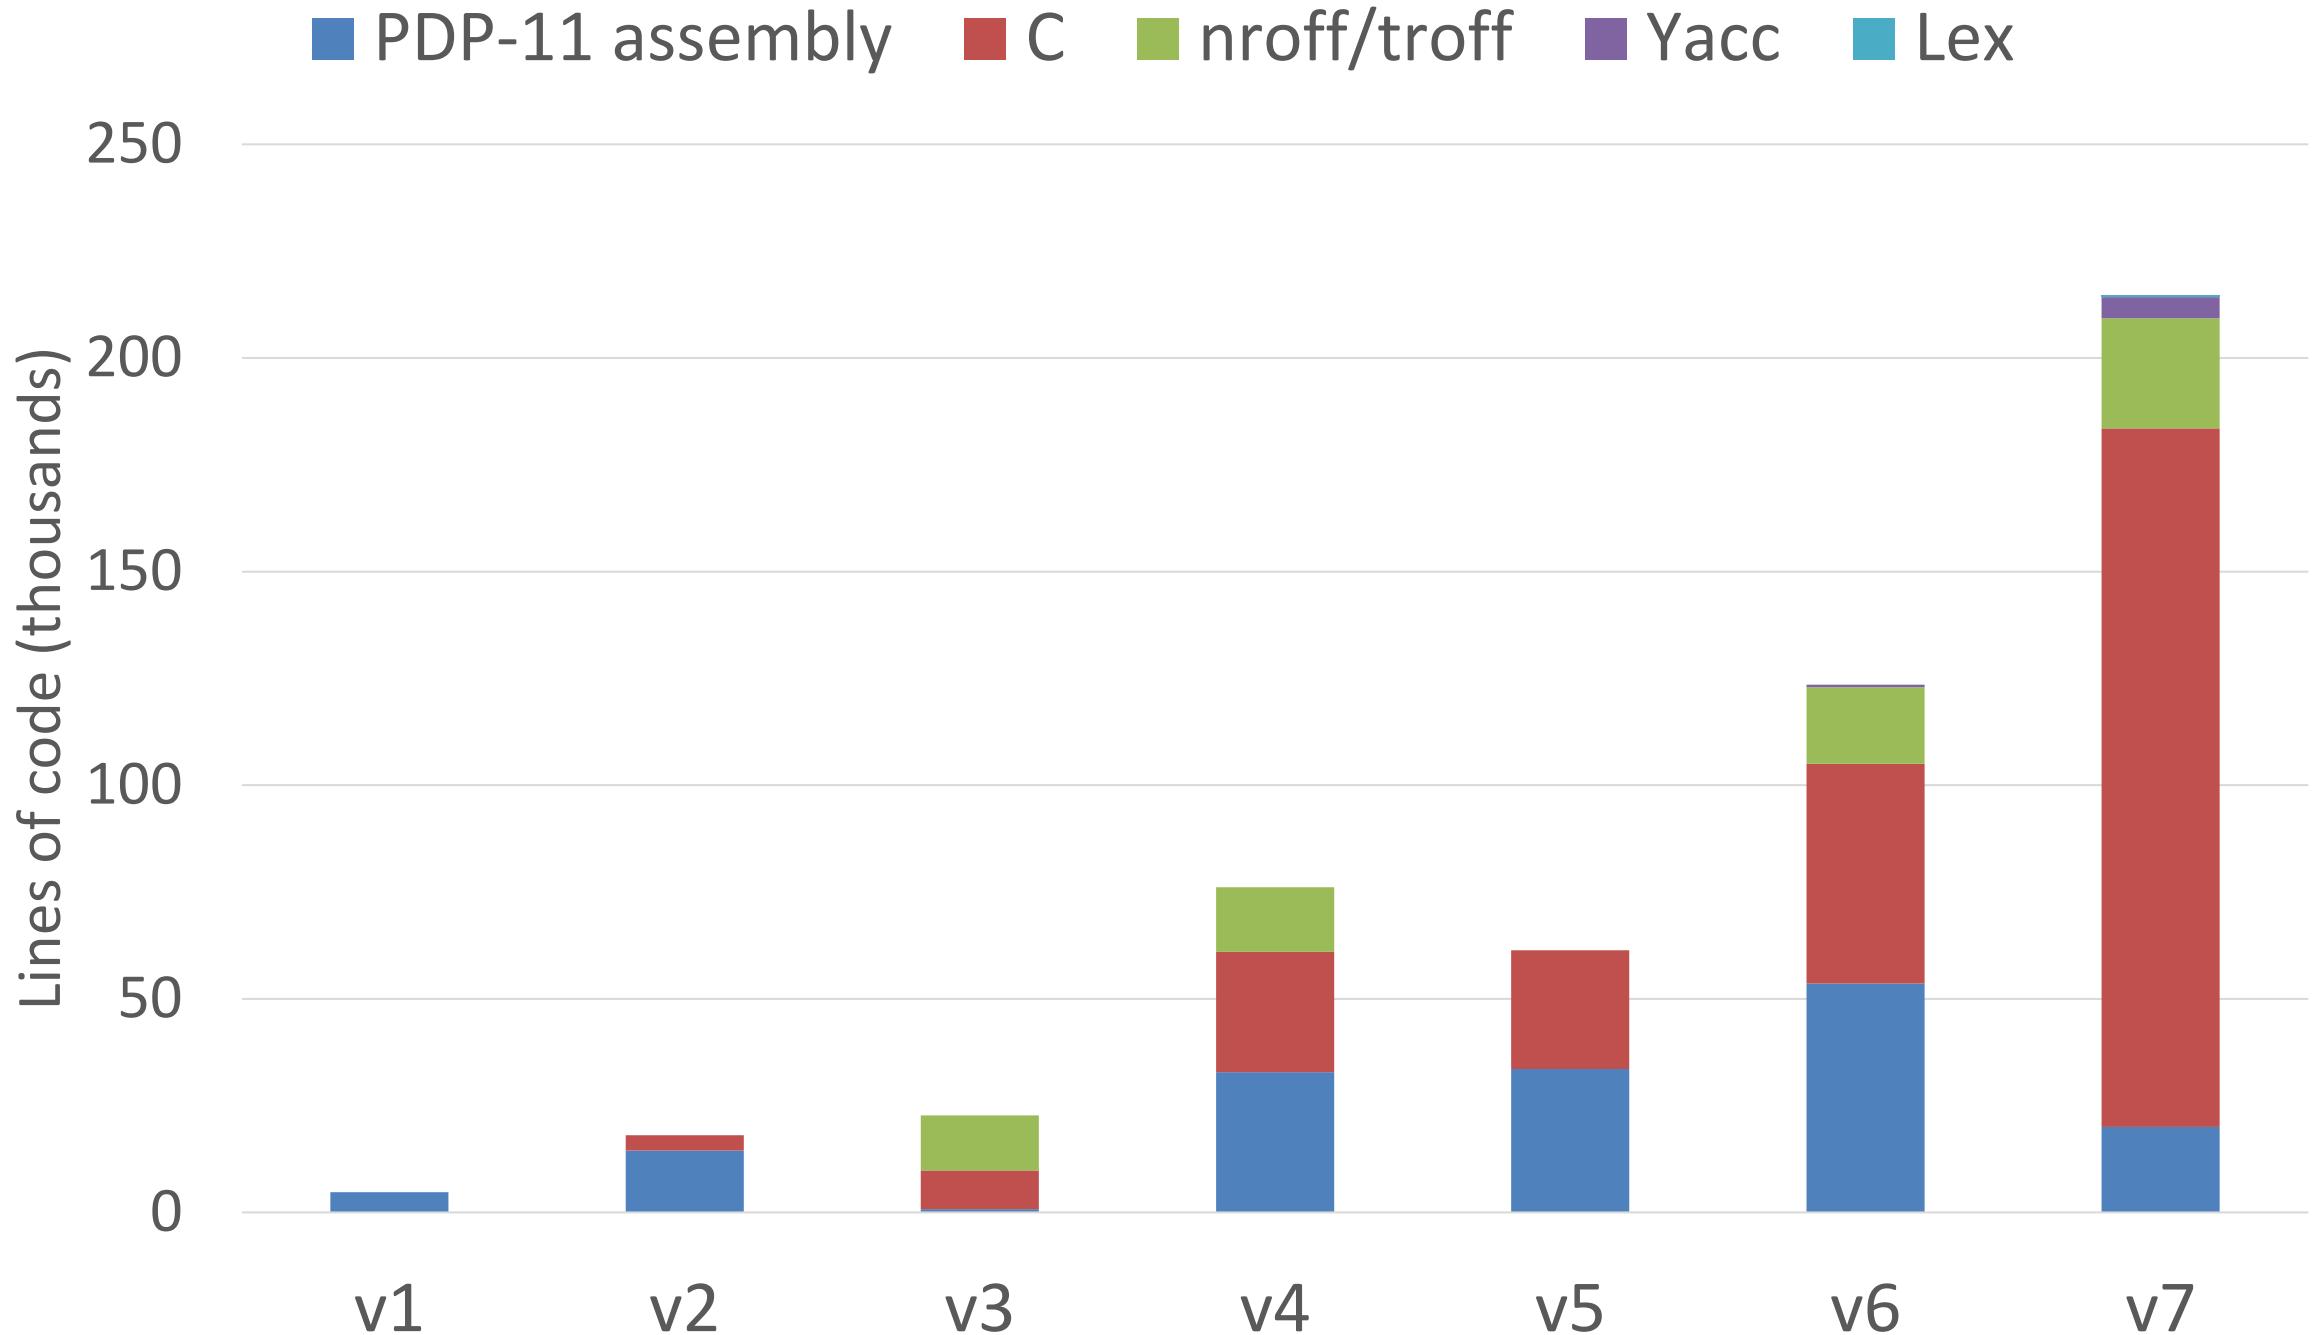
# Inherited code

```
for ref in Research-V{3,4,5}-Snapshot-Development ; do
    echo $ref
    # For all the edition's files
    git ls-tree -r --name-only $ref |
        # Exclude administrative files introduced in the repo.
        grep -Ev 'README|LICENSE|\.pdf|\.ref' |
        # Run git-blame on each.
        xargs -I '{}' git blame -M -M -C -C $ref -- '{}' |
        awk '{print $1}' |
        sort |
        # Sum lines for each commit
        uniq -c |
        # Obtain lines and provenance of each commit; output totals.
        awk '{("git show " $2 "| awk '\"/Synthesized-from:/ {print $2}'\"") |
            getline ver; total[ver] += $1 }
            END {for (v in total) print v, total[v]}'
done
```



# Evolution of language use

```
for ref in Research-V{1,2,3,4,5,6,7}-Snapshot-Development ; do
    echo -n $ref | sed 's/Research-V/v/;s/-.*/\t/'
    # For all the edition's files
    git ls-tree -r --name-only $ref |
        # Exclude administrative files introduced in the repo.
        grep -Ev 'README|LICENSE|\.\.pdf|\.\.ref' |
        # Run git-show on each.
        while read path ; do
            echo -n "$path "
            git show $ref:$path | wc -l
        done |
        awk '$1 ~ /\./ {
            n = split($1, arr, /\./);
            suff = arr[n];
            if (suff ~ /^[0-9][sm]?$/) suff = "r";
            if (suff ~ /^[ch]$/) suff = "C";
            if (suff !~ /\V/)
                lines[suff] += $2
        }
        #END {for (s in lines) if (lines[s] > 500) print s, lines[s]}
        END { OFS = "\t"; print lines["s"], lines["C"], lines["r"], lines["y"], lines["l"]; }'
done
```



# Thank you!



[github.com/dspinellis/unix-history-repo/](https://github.com/dspinellis/unix-history-repo/)

- 🌐 [www.spinellis.gr](http://www.spinellis.gr)
- 🦋 [@CoolSWEng](https://@CoolSWEng)
- Ⓜ [@CoolSWEng@mastodon.acm.org](https://@CoolSWEng@mastodon.acm.org)
- ✉️ [dds@aueb.gr](mailto:dds@aueb.gr)

# Backup slides

# Evolution of the Unix System Architecture: An Exploratory Case Study

Diomidis Spinellis, Senior Member, IEEE, and Paris Avgeriou, Senior Member, IEEE

**Abstract**—Unix has evolved for almost five decades, shaping modern operating systems, key software technologies, and development practices. Studying the evolution of this remarkable system from an architectural perspective can provide insights on how to manage the growth of large, complex, and long-lived software systems. Along main Unix releases leading to the FreeBSD lineage we examine core architectural design decisions, the number of features, and code complexity, based on the analysis of source code, reference documentation, and related publications. We report that the growth in size has been uniform, with some notable outliers, while cyclomatic complexity has been religiously safeguarded. A large number of Unix-defining design decisions were implemented right from the very early beginning, with most of them still playing a major role. Unix continues to evolve from an architectural perspective, but the rate of architectural innovation has slowed down over the system's lifetime. Architectural technical debt has accrued in the forms of functionality duplication and unused facilities, but in terms of cyclomatic complexity it is systematically being paid back through what appears to be a self-correcting process. Some unsung architectural forces that shaped Unix are the emphasis on conventions over rigid enforcement, the drive for portability, a sophisticated ecosystem of other operating systems and development organizations, and the emergence of a federated architecture, often through the adoption of third-party subsystems. These findings have led us to form an initial theory on the architecture evolution of large, complex operating system software.

**Index Terms**—Unix, Software Architecture, Software Evolution, Architecture Design Decisions, Operating Systems.

## 1 INTRODUCTION

UNIX<sup>1</sup> has a long and celebrated history. Its evolution spans five decades and is a result of the work by thousands of developers, including several distinguished pioneers. As an operating system, it has left an undeniable mark on the history of computing, while it has influenced tremendously the current state of the art in software, network, and hardware engineering.

Studying the evolution of operating system software is not just significant from a historical perspective; it can provide valuable insights into evolvability best practices and anti-patterns, for large, complex, and long-lived systems. Unix is a unique case among all operating systems, both due to its longevity, and its impact on the operating systems that followed. The evolution of a system of this size, complexity and age can shed light on how similar systems can sustainably grow without the perils of software aging like soaring technical debt or uncontrolled architectural decay.

In this paper we study the evolution of Unix along the FreeBSD lineage from a software architecture perspective. While there have been studies on how Unix evolved (see Section 2), these have mostly focused at the source code level and were limited to the kernel. On the contrary, we turn our attention to the system architecture and study a) the core architectural design decisions across the main

releases, and b) the evolution in the number of the system's features (obtained from the Unix reference documentation) and in the code's complexity. The former entails qualitative analysis, while the latter quantitative. These analyses subsequently lead to forming an initial theory on the architecture evolution of large and complex operating systems, regarding their form, pace, driving forces, as well as the accumulation of architectural technical debt.

The rest of the paper is structured as follows: In Section 2 we present related work, whereas in Section 3 we elaborate on the case study design. In sections 4 and 5, we present the qualitative results (main architectural design decisions), and the quantitative results (evolution of size and complexity) respectively. Next, in Section 6 we discuss the main findings, and in Section 7 the threats to this study's validity. Finally, in Section 8 we conclude the paper with a summary and discussion of our findings.

## 2 RELATED WORK

The work reported here covers mainly two areas: a) software evolution in general, which has been intensely studied, and b) the evolution of Unix in particular, where related work is more thin on the ground.

### 2.1 Software Evolution

There have been several studies on the longitudinal evolution of large systems. The seminal work of Lehman [1] and its subsequent refinements attempted to establish laws of software evolution, not unlike those of biological evolution. Those laws have been the subject of much discussion and research work [2]: their validity has been long debated, their

- D. Spinellis is with the Athens University of Economics and Business, Greece.  
E-mail: see <http://www.dmsr.aueb.gr/dds/>
- P. Avgeriou is with the University of Groningen.

Manuscript received December 19, 2016.

1. UNIX® is a registered trademark of The Open Group. For the sake of simplicity, in this paper we use the word "Unix" to refer both to UNIX systems developed at Bell Labs and to Unix-like systems, such as FreeBSD, that descended from them.



**Diomidis Spinellis and Paris Avgeriou.**  
**Evolution of the Unix system**  
**architecture: An exploratory case study.**  
*IEEE Transactions on Software Engineering, 47:1134–1163, June 2021.*  
[doi:10.1109/TSE.2019.2892149](https://doi.org/10.1109/TSE.2019.2892149)

## A repository of Unix history and evolution

Diomidis Spinellis<sup>1</sup> 

Published online: 11 August 2016  
© Springer Science+Business Media New York 2016

**Abstract** The history and evolution of the Unix operating system is made available as a revision management repository, covering the period from its inception in 1972 as a five thousand line kernel, to 2016 as a widely-used 27 million line system. The 1.1GB repository contains 496 thousand commits and 2,523 branch merges. The repository employs the commonly used Git version control system for its storage, and is hosted on the popular GitHub archive. It has been created by synthesizing with custom software 24 snapshots of systems developed at Bell Labs, the University of California at Berkeley, and the 386BSD team, two legacy repositories, and the modern repository of the open source FreeBSD system. In total, 973 individual contributors are identified, the early ones through primary research. The data set can be used for empirical research in software engineering, information systems, and software archaeology.

**Keywords** Software archeology · Unix · Configuration management · Git

### 1 Introduction

The Unix operating system stands out as a major engineering breakthrough due to its exemplary design, its numerous technical contributions, its impact, its development model, and its widespread use (Gehani 2003, pp. 27–29). The design of the Unix programming

---

Communicated by: Romain Robbes, Martin Pinzger and Yasutaka Kamei

The work has been partially funded by the Research Centre of the Athens University of Economics and Business, under the Original Scientific Publications framework (project code EP-2279-01) and supported by computational time granted from the Greek Research & Technology Network (GRNET) in the National HPC facility — ARIS — under project ID PA003005-CDOLPOT.

---

✉ Diomidis Spinellis  
dds@auer.gr

<sup>1</sup> Department of Management Science and Technology, Athens University of Economics and Business, Patisson 76, 104 34 Athens, Greece



**Diomidis Spinellis. A repository of Unix History and evolution. *Empirical Software Engineering*, 22(3):1372–1404, 2017. [doi:10.1007/s10664-016-9445-5](https://doi.org/10.1007/s10664-016-9445-5)**

## Documented Unix Facilities Over 48 Years

Diomidis Spinellis

Department of Management Science and Technology  
Athens University of Economics and Business  
Athens, Greece  
[dds@aueb.gr](mailto:dds@aueb.gr)

### ABSTRACT

The documented Unix facilities data set provides the details regarding the evolution of 15 596 unique facilities through 93 versions of Unix over a period of 48 years. It is based on the manual transcription of early scanned documents, on the curation of text obtained through optical character recognition, and on the automatic extraction of data from code available on the Unix History Repository. The data are categorized into user commands, system calls, C library functions, devices and special files, file formats and conventions, games et. al., miscellanea, system maintenance procedures and commands, and system kernel interfaces. A timeline view allows the visualization of the evolution across releases. The data can be used for empirical research regarding API evolution, system design, as well as technology adoption and trends.

### CCS CONCEPTS

- Software and its engineering → Software evolution;

### ACM Reference Format:

Diomidis Spinellis. 2018. Documented Unix Facilities Over 48 Years. In *MSR '18: 15th International Conference on Mining Software Repositories, May 28–29, 2018, Gothenburg, Sweden*. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3196398.3196476>

### 1 INTRODUCTION

The Unix operating system is being continuously developed from the same code base for almost over half a century. It stands out as a major engineering artefact due to its exemplary design, its numerous technical contributions, its impact, its development model, and its widespread use [2, pp. 27–29], [9]. The design of the Unix programming environment, which nowadays offers thousands of tools and libraries, has been characterized as offering unusual simplicity, power, and elegance [5, 7]. Consequently, empirical data regarding how the facilities Unix provides grew and changed over time can be used for empirical research on API evolution, system design, as well as technology adoption and trends.

Although one can study a system's evolution through its source code [1, 10], the very large size of modern systems can hinder the recognition of the relevant parts. Fortunately, another avenue is

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyright for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](http://permissions.acm.org).

MSR '18, May 28–29, 2018, Gothenburg, Sweden

© 2018 Copyright held by the owner/author(s). Publication rights licensed to the Association for Computing Machinery.  
ACM ISBN 978-1-4503-5716-6/18/05...\$15.00  
<https://doi.org/10.1145/3196398.3196476>

available for studying Unix systems, namely their documentation. From the first version of the Unix system until today, every release is accompanied by a complete reference manual, where all provided facilities (commands, APIs, file formats, and device drivers) are neatly organized into several corresponding sections (see Table 1). The central role of the reference manual in the Unix system is evidenced by the fact that early Unix versions coming out of AT&T Bell Labs were named after the edition of the accompanying manual. Some early editions of the manuals have not survived in a machine-readable format, but most are available in text markup that can be processed through scripts to extract relevant data.

The data set presented here is based on the printed and machine-readable Unix reference manuals released over a period of 48 years. It documents the evolution of 15 596 facilities through 93 versions of Unix. Section 2 outlines the provided data, Section 3 describes how the data were produced, and Section 4 sketches two examples of how the data can be used for quantitative and qualitative empirical studies.

### 2 UNIX RELEASES AND THEIR FACILITIES

The primary data are made available in the form of 93 text files containing 405 726 records. The files are named after the associated Unix release, following the tags and branches nomenclature established in the Unix History Repository [9]. A record is a text line with tab-separated fields. Each record contains the number of the Unix manual section associated with a facility (1–9; see Table 1) followed by the facility's name, optionally followed by a URI identifying the facility's documentation *troff* markup [6]. In total, the set contains data about 15 596 facilities pointing to 193 781 unique URIs, identifying 48 250 distinct manual page instances.

As an example, the following lines show the documentation files associated with the label command (:), the archiver (*ar*), and the assembler (*as*), as documented in Section I of the 1973 Third Edition Unix manual.

```
1      :      Research-V3/man/man1/:.1
1      ar     Research-V3/man/man1/ar.1
1      as     Research-V3/man/man1/as.1
```

By prepending the Unix History Repository GitHub permalink base URL "<https://github.com/dspinellis/unix-history-repo/blob/>" to an entry's URI, one can obtain a URL for viewing the documentation markup source code for the corresponding entry.

A separate text file, named *timeline* associates each of the data files with the year, month, and day of the corresponding release. For instance, the following entries of the *timeline* file list the dates associated with the Sixth and Seventh Research Editions and the first Berkeley Software Distribution.

```
Research-V6 1975 07 18
BSD-1 1978 02 01
Research-V7 1979 08 26
```

**Diomidis Spinellis. Documented Unix facilities over 48 years. In MSR '18: Proceedings of the 15th Conference on Mining Software Repositories, pages 58–61, New York, NY, USA, May 2018. Association for Computing Machinery.**  
[doi:10.1145/3196398.3196476](https://doi.org/10.1145/3196398.3196476)



# Free open edX course (MOOC): Unix Tools: Data, Software and Production Engineering

---

Grow from being a Unix novice to Unix wizard status!  
Process big data, analyze software code, run DevOps tasks and excel in your everyday job through the amazing power of the Unix shell and command-line tools.

<https://www.spinellis.gr/unix>

