

TRUST THE MATH, FEAR THE COMPILER:

How Optimisations undermine
Cryptographic Software

René Meusel
FOSDEM 2026: /dev/random

ROHDE & SCHWARZ

Make ideas real





René Meusel

Rohde & Schwarz Cybersecurity

 @reneme

 Co-Maintaining the Botan* crypto library
Contributing to Open Source since 2011

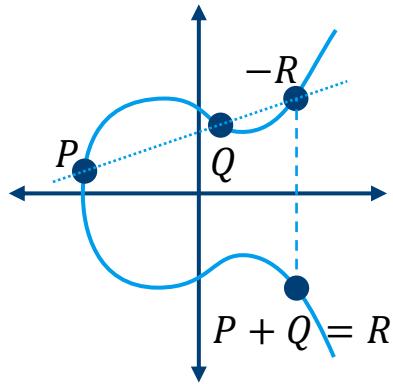
WHY DO WE TRUST THE MATH?

Prime Factorization Problem

15251262345256836781

$$= p_1? \times p_2?$$

(Elliptic Curve) Discrete Logarithm Problem



Learning with Errors Problem

$$\begin{bmatrix} b \\ \vdots \end{bmatrix} = \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & A & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} * \begin{bmatrix} s \\ \vdots \end{bmatrix} + \begin{bmatrix} e \\ \vdots \end{bmatrix}$$

All these problems are **believed** to be hard enough for cryptography.

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

Enter your password

* * * * *

POST /auth

```
if (equal(req.password, db.password))  
    okay();  
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

Enter your password

* * * * *

POST /auth

```
if (equal(req.password, db.password))
    okay();
else
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database


```
bool equal(std::string pw1, std::string pw2)
{
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        if (a != b)
        {
            // early return for efficiency!
            return false;
        }
    }
    return pw1.size() == pw2.size();
}
```

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

Enter your password

* * * * *

POST /auth

```
if (equal(req.password, db.password))  
    okay();
```

```
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

```
bool equal(std::string pw1, std::string pw2)  
{  
    for (auto [a, b] : std::views::zip(pw1, pw2))  
    {  
        if (a != b)  
        {  
            // early return for efficiency!  
            return false;  
        }  
    }  
    return pw1.size() == pw2.size();  
}
```

This is a Timing Side Channel!

Program behaviour depends on a secret

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

fosdem26

Enter your password

* * * * *

POST /auth



"a" → HTTP 401 after 1ms

```
if (equal(req.password, db.password))  
    okay();
```

```
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

```
bool equal(std::string pw1, std::string pw2)  
{  
    for (auto [a, b] : std::views::zip(pw1, pw2))  
    {  
        if (a != b)  
        {  
            // early return for efficiency!  
            return false;  
        }  
    }  
    return pw1.size() == pw2.size();  
}
```

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

fosdem26

Enter your password

* * * * *

POST /auth



"a" -> HTTP 401 after **1ms**
"b" -> HTTP 401 after **1ms**
"c" -> HTTP 401 after **1ms**
"d" -> HTTP 401 after **1ms**
"e" -> HTTP 401 after **1ms**
"f" -> HTTP 401 after **2ms**

```
if (equal(req.password, db.password))  
    okay();
```

```
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

```
bool equal(std::string pw1, std::string pw2)  
{  
    for (auto [a, b] : std::views::zip(pw1, pw2))  
    {  
        if (a != b)  
        {  
            // early return for efficiency!  
            return false;  
        }  
    }  
    return pw1.size() == pw2.size();  
}
```

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

fosdem26

Enter your password

* * * * *

POST /auth



"a" -> HTTP 401 after **1ms**
"b" -> HTTP 401 after **1ms**
"c" -> HTTP 401 after **1ms**
"d" -> HTTP 401 after **1ms**
"e" -> HTTP 401 after **1ms**
"f" -> HTTP 401 after **2ms**
"fa" -> HTTP 401 after **2ms**
"fb" -> HTTP 401 after **2ms**
"fc" -> HTTP 401 after **2ms**
...
...

```
if (equal(req.password, db.password))  
    okay();
```

```
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

```
bool equal(std::string pw1, std::string pw2)  
{  
    for (auto [a, b] : std::views::zip(pw1, pw2))  
    {  
        if (a != b)  
        {  
            // early return for efficiency!  
            return false;  
        }  
    }  
    return pw1.size() == pw2.size();  
}
```

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

fosdem26

Enter your password

* * * * *

POST /auth



"a" -> HTTP 401 after **1ms**
"b" -> HTTP 401 after **1ms**
"c" -> HTTP 401 after **1ms**
"d" -> HTTP 401 after **1ms**
"e" -> HTTP 401 after **1ms**
"f" -> HTTP 401 after **2ms**
"fa" -> HTTP 401 after **2ms**
"fb" -> HTTP 401 after **2ms**
"fc" -> HTTP 401 after **2ms**
...
"fosdem26" -> **HTTP 200** after **8ms**

```
if (equal(req.password, db.password))  
    okay();
```

```
else  
    unauthorized();
```

Let's ignore that we are saving plain passwords in our database

```
bool equal(std::string pw1, std::string pw2)  
{  
    for (auto [a, b] : std::views::zip(pw1, pw2))  
    {  
        if (a != b)  
        {  
            // early return for efficiency!  
            return false;  
        }  
    }  
    return pw1.size() == pw2.size();  
}
```

The math is solid and trustworthy...

IT'S THE IMPLEMENTATIONS THAT FAIL!

"Constant-time" implementation

```
bool equal(std::string pw1, std::string pw2)
{
    bool match = true;
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        // no side-channel: just keep comparing
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}
```

Secret-dependent control flow

```
bool equal(std::string pw1, std::string pw2)
{
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        if (a != b)
        {
            // early return for efficiency!
            return false;
        }
    }
    return pw1.size() == pw2.size();
}
```

IT'S THE IMPLEMENTATIONS THAT FAIL!

“Constant-time” implementation

```
bool equal(std::string pw1, std::string pw2)
{
    bool match = true;
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        // no side-channel: just keep comparing
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}
```

=> GCC 15.2, with -std=c++23 -O3

```
[ Loop setup ]
    mov    ecx, 1
.loop:
    test   cl, cl
    je     .way_out
    movzx ecx, BYTE PTR [rdx]
    cmp    BYTE PTR [rax], cl
    sete  cl
    [ if needed, goto .loop ]
.way_out:
    cmp    r9, r8
    sete  al
    and   eax, ecx
    ret
```

IT'S THE IMPLEMENTATIONS THAT FAIL!

“Constant-time” implementation

```
bool equal(std::string pw1, std::string pw2)
{
    bool match = true;
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        // no side-channel: just keep comparing
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}
```

=> GCC 15.2, with -std=c++23 -O3

```
[ Loop setup ]
mov    ecx, 1
.loop:
test   cl, cl
je     .way_out
movzx ecx, BYTE PTR [rdx]
cmp    BYTE PTR [rax], cl
sete  cl
[ if needed, goto .loop ]
.way_out:
cmp    r9, r8
sete  al
and    eax, ecx
ret
```

match = (a == b);

IT'S THE IMPLEMENTATIONS THAT FAIL!

“Constant-time” implementation

```
bool equal(std::string pw1, std::string pw2)
{
    bool match = true;
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        // no side-channel: just keep comparing
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}
```

=> GCC 15.2, with `-std=c++23 -O3`

<pre>[Loop setup] mov ecx, 1 .loop: test cl, cl je .way_out movzx ecx, BYTE PTR [rdx] cmp BYTE PTR [rax], cl sete cl [if needed, goto .loop] .way_out: cmp r9, r8 sete al and eax, ecx ret</pre>	 <pre>if (!match) break; match = (a == b);</pre>
---	---

GCC “optimized” the side-channel back into our safe implementation!

Use `uint32_t` instead of `bool`

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    bool match = true;
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}

/// @returns 0x00000000 if v is zero,
///          0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    return ((~v & (v - 1)) >> 31) - 1;
}
```

Optimizing boolean logic is easy for the compiler.

... let's try to hide booleans behind ordinary integers.



15

René Meusel | FOSDEM 2026 | How Optimisations undermine Cryptographic Software

<https://godbolt.org/z/cKGxP6T75>

Use `uint32_t` instead of `bool`

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        match = match && (a == b);
    }
    return match && pw1.size() == pw2.size();
}

/// @returns 0x00000000 if v is zero,
///          0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    return ((~v & (v - 1)) >> 31) - 1;
}
```

Optimizing boolean logic is easy for the compiler.

... let's try to hide booleans behind ordinary integers.



16

René Meusel | FOSDEM 2026 | How Optimisations undermine Cryptographic Software

<https://godbolt.org/z/cKGxP6T75>

Use `uint32_t` instead of `bool`

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}

/// @returns 0x00000000 if v is zero,
///          0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    return ((~v & (v - 1)) >> 31) - 1;
}
```

Optimizing boolean logic is easy for the compiler.

... let's try to hide booleans behind ordinary integers.



17

René Meusel | FOSDEM 2026 | How Optimisations undermine Cryptographic Software

<https://godbolt.org/z/cKGxP6T75>

Use `uint32_t` instead of `bool`

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }

    return match & expand(pw1.size() == pw2.size());
}

[ Loop setup ]
mov    ecx, -1
.loop:
    movzx esi, BYTE PTR [rdx]
    cmp    BYTE PTR [rax], sil
    setne sil
    movzx esi, sil
    sub    esi, 1
    and    ecx, esi
[ if needed, goto .loop ]
.after_loop:
    xor    eax, eax
    cmp    r9, r10
    setne al
    sub    eax, 1
    test   eax, ecx
    setne al
    ret
```



Use `uint32_t` instead of `bool`

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return match & expand(pw1.size() == pw2.size());
}

[ Loop setup ]
mov    ecx, -1
.loop:
    movzx esi, BYTE PTR [rdx]
    cmp    BYTE PTR [rax], sil
    setne sil
    movzx esi, sil
    sub    esi, 1
    and    ecx, esi
[ if needed, goto .loop ]
.after_loop:
    xor    eax, eax
    cmp    r9, r10
    setne al
    sub    eax, 1
    test   eax, ecx
    setne al
    ret
```



esi = $(a \neq b)$

mask &= $(esi-1) ((~v \& (v-1)) \gg 31) - 1$

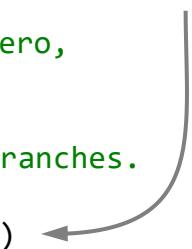
*Technically, okay.
But GCC continues
to reason about
the boolean!*

HIDING VALUES AND SEMANTICS FROM THE COMPILER

Due to inlining, the compiler might know that the parameter is actually a boolean!

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}
```

```
/// @returns 0x00000000 if v is zero,
///           0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    return ((~v & (v - 1)) >> 31) - 1;
}
```



HIDING VALUES AND SEMANTICS FROM THE COMPILER

Due to inlining, the compiler might know that the parameter is actually a boolean!

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}
```

```
/// @returns 0x00000000 if v is zero,
///           0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    const uint32_t v2 = obfuscate(v);
    return ((~v2 & (v2 - 1)) >> 31) - 1;
}
```

Use advanced tricks to obfuscate integer value ranges

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}
```

*Due to inlining, the compiler might know
that the parameter is actually a boolean!*

```
/// @returns 0x00000000 if v is zero,
///           0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    const uint32_t v2 = obfuscate(v);
    return ((~v2 & (v2 - 1)) >> 31) - 1;
}
```

*Extracts the most-significant bit.
Compilers might recognize that as boolean!*

Compilers infer plausible integer ranges and use that for optimization.

Use advanced tricks to obfuscate integer value ranges

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}
```

*Due to inlining, the compiler might know
that the parameter is actually a boolean!*

```
/// @returns 0x00000000 if v is zero,
///           0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    const uint32_t v2 = obfuscate(v);
    return obfuscate((~v2 & (v2 - 1)) >> 31) - 1;
}
```

*Extracts the most-significant bit.
Compilers might recognize that as boolean!*

Compilers infer plausible integer ranges and use that for optimization.

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```

[ preamble, Loop setup ]
.loop:
    movzx  eax, BYTE PTR [r8]
    xor    edx, edx
    cmp    BYTE PTR [rcx], al
    sete   dl
    /// @returns 0x00000000 if v is zero,
    ///          0xFFFFFFFF otherwise
    /// Uses only bitwise logic, no branches.

a == b

expand(uint32_t v)
{
    ((~v&(v-1))>>31)-1
}

mask &= ...

[ if needed, goto .Loop ]
.way_out:
[ epilogue ]

```

mov eax, edx
sub edx, 1
not eax
and eax, edx
shr eax, 31
sub eax, 1
and r9d, eax

///
inline uint32_t expand(uint32_t v)
{
const uint32_t v2 = obfuscate(v);
return obfuscate((~v2 & (v2 - 1)) >> 31) - 1;
}



Yay! No more optimization!

HIDING VALUES AND SEMANTICS FROM THE COMPILER

```
/// Hide the value of @p v from the compiler
/// and return it unchanged.
///
inline uint32_t obfuscate(uint32_t v)
{
    asm("") : "+r"(v) :;
    return v;
}
```



```
/// @returns 0x00000000 if v is zero,
///           0xFFFFFFFF otherwise
///
/// Uses only bitwise logic, no branches.
///
inline uint32_t expand(uint32_t v)
{
    const uint32_t v2 = obfuscate(v);
    return obfuscate((~v2 & (v2 - 1)) >> 31) - 1;
}
```

Compilers cannot reason about values produced by inline assembly.

```
bool equal(std::string pw1, std::string pw2)
{
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        if (a != b)
        {
            // early return for efficiency!
            return false;
        }
    }
    return pw1.size() == pw2.size();
}
```

```

bool equal(std::string pw1, std::string pw2)
{
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        if (a != b)
        {
            // early return for efficiency!
            return false;
        }
    }
    return pw1.size() == pw2.size();
}

```

This is a mess! How to maintain this?



```

inline uint32_t obfuscate(uint32_t v)
{
    asm("") : "+r"(v) :);
    return v;
}

inline uint32_t expand(uint32_t v)
{
    const uint32_t v2 = obfuscate(v);
    return obfuscate((~v2 & (v2 - 1)) >> 31) - 1;
}

bool equal(std::string pw1, std::string pw2)
{
    uint32_t mask = expand(true);
    for (auto [a, b] : std::views::zip(pw1, pw2))
    {
        mask = mask & expand(a == b);
    }
    return mask & expand(pw1.size() == pw2.size());
}

```

SURPRISINGLY: VALGRIND CAN HELP

Valgrind warns if program behaviour depends on “uninitialized memory”.

```
const std::string pw = get_password_from_user();

VALGRIND_MAKE_MEM_UNDEFINED(pw.data(), pw.size());
const bool ok = bad_equal(pw, the_password);
VALGRIND_MAKE_MEM_DEFINED(&ok, sizeof(ok));

if(ok) // this branch is fine!
    okay();
else
    unauthorized();
```

SURPRISINGLY: VALGRIND CAN HELP

Valgrind warns if program behaviour depends on “uninitialized memory”.

... let's just claim that our “secret memory” is “uninitialized”.



```
const std::string pw = get_password_from_user();
VALGRIND_MAKE_MEM_UNDEFINED(pw.data(), pw.size());
const bool ok = bad_equal(pw, the_password);
VALGRIND_MAKE_MEM_DEFINED(&ok, sizeof(ok));

if(ok) // this branch is fine!
    okay();
else
    unauthorized();
```

```
#> valgrind ./check_password fosdem2026
[...]
Conditional jump or move depends on uninitialized value(s)
  at: bad_equal(...) (equal.cpp:41)
  by: main (equal.cpp:68)
```

SURPRISINGLY: VALGRIND CAN HELP

```
const std::string pw = get_password_from_user();

VALGRIND_MAKE_MEM_UNDEFINED(pw.data(), pw.size());
const bool ok = bad_equal(pw, the_password);
VALGRIND_MAKE_MEM_DEFINED(&ok, sizeof(ok));

if(ok) // this branch is fine!
    okay();
else
    unauthorized();
```

Valgrind warns if program behaviour depends on “uninitialized memory”.

... let's just claim that our “secret memory” is “uninitialized”.



Valgrind even finds that results were calculated from “uninitialized” / “secret” data.

... we have to “declassify” inferred values before using them to determine program flow.

```
#> valgrind ./check_password fosdem2026
[...]
Conditional jump or move depends on uninitialized value(s)
  at: bad_equal(...) (equal.cpp:41)
  by: main (equal.cpp:68)
```

SURPRISINGLY: VALGRIND CAN HELP

```
const std::string pw = get_password_from_user();
VALGRIND_MAKE_MEM_UNDEFINED(pw.data(), pw.size());
const bool ok = bad_equal(pw, the_password);
VALGRIND_MAKE_MEM_DEFINED(&ok, sizeof(ok));
```

```
if(ok) // this branch is fine!
    okay();
else
    unauthorized();
```

```
#> valgrind ./check_password fosdem2026
```

```
[...]
```

```
Conditional jump or move depends on uninitialized value(s)
at: bad_equal(...) (equal.cpp:41)
by: main (equal.cpp:68)
```

... let's just do
is "uninitializ

... we have to "decla.
using them to deter

Valgrind warns if program behaviour
depends on uninitialized values

Matrix: valgrind

✓ valgrind (clang, -O1)	1h 13m
✓ valgrind (clang, -O2)	1h 5m
✓ valgrind (clang, -O3)	1h 4m
✓ valgrind (clang, -Os)	1h 20m
✓ valgrind (gcc-14, -O1)	1h 8m
✓ valgrind (gcc-14, -O2)	33m 18s
✓ valgrind (gcc-14, -O3)	31m 20s
✓ valgrind (gcc-14, -Os)	1h 17m



Botan's nightly tests

SUMMARY

- ▶ **Compilers make code *efficient*.**
But they don't take other qualitative requirements into account.
- ▶ **Security products must always keep the entire system in mind.**
Depending on the hardware, even individual instructions might have side-channels.¹
- ▶ **Implementing cryptography is tough, not just because of the math.**
Don't do it on your own, contribute to existing projects if you want to learn.

=> GCC 15.2, with `-std=c++23 -O3`

```
[ Loop setup ]
mov    ecx, 1
.loop:
test   cl, cl
je     .way_out
movzx  ecx, BYTE PTR [rdx]
cmp    BYTE PTR [rax], cl
sete   cl
[ if needed, goto .loop ]
.way_out:
cmp    r9, r8
sete   al
and    eax, ecx
ret
```

if (!match)
break;

