


Where should my VIRTIO device live?

FOSDEM 2026

Stefano Garzarella

[<sqarzare@redhat.com>](mailto:sqarzare@redhat.com)






VIRTIO specification

- [Virtual I/O Device \(VIRTIO\) Version 1.3](#)
 - *The purpose of virtio and this specification is that virtual environments and guests should have a **straightforward, efficient, standard and extensible mechanism for virtual devices**, rather than boutique per-environment or per-OS mechanisms.*
 - [VIRTIO 1.4 \(Public Review Draft\)](#)
- <https://github.com/oasis-tcs/virtio-spec>
 - Authoritative source of the VIRTIO (Virtual I/O) Specification
- Virtual I/O devices
 - core components (features, notifications, configuration, virtqueues, etc.)
 - initialization steps
 - transports (PCI, MMIO, Channel I/O)
 - device types (e.g. net, block, vsock, sound, fs, etc.)


VIRTIO core components



- Control path
 - features negotiation
 - configuration space
 - data path setup
- Data path
 - virtqueue
 - split / packed
 - always allocated by the guest
- Notifications
 - kick
 - guest -> host
 - irq
 - host -> guest

VIRTIO device types


- Several device types in the specification
 - virtio-net
 - Network card
 - virtio-blk
 - Block device (HDD, SSD)
 - virtio-vsock
 - Virtual Socket
 - virtio-fs
 - File system (e.g. shared folder)
 - and others
- Modern / Legacy
 - Legacy (pre VIRTIO v1.0)
 - Modern



Where should my VIRTIO device live?

Where the device could be?


- **VMM built-in**
 - e.g. in QEMU process
- user space (host) device
 - vhost-user
 - VDUSE (vDPA in user space)
- kernel (host) device
 - vhost
 - vDPA SW device
- hardware
 - vDPA HW device



Where should my VIRTIO device live?

Where the device could be?


- VMM built-in
 - e.g. in QEMU process
- **user space (host) device**
 - **vhost-user**
 - **VDUSE (vDPA in user space)**
- kernel (host) device
 - vhost
 - vDPA SW device
- hardware
 - vDPA HW device



Where should my VIRTIO device live?


Where the device could be?

- VMM built-in
 - e.g. in QEMU process
- user space (host) device
 - vhost-user
 - VDUSE (vDPA in user space)
- **kernel (host) device**
 - **vhost**
 - **vDPA SW device**
- hardware
 - vDPA HW device




Where the device could be?

- VMM built-in
 - e.g. in QEMU process
- user space (host) device
 - vhost-user
 - VDUSE (vDPA in user space)
- kernel (host) device
 - vhost
 - vDPA SW device
- **hardware**
 - **vDPA HW device**




VIRTIO device emulated by the VMM




- Common scenario
- QEMU
 - de facto reference implementation for VIRTIO devices
- Why?
 - Simplicity
 - Portability
- Drawbacks
 - Reliability risk
 - Performance overhead

VIRTIO device emulated by the VMM




```
$ qemu-system-x86_64 -smp 2 -m 2G \
-M q35,accel=kvm \
-blockdev file,filename=fedora.qcow2,node-name=file \
-blockdev qcow2,file=file,node-name=qcow2 \
-device virtio-blk-pci,drive=qcow2
```

vhost: VIRTIO device in the host kernel




- Initially introduced to increase performance of virtio-net device
 - Control path
 - IOCTLs
 - Data path
 - kthread/vhost_task attaches VMMS address space
- Linux kernel supports
 - vhost-net, vhost-scsi, vhost-vsock
- Why?
 - Performance
 - Easily to integrate with host kernel stacks (e.g. AF_VSOCK)
- Drawbacks
 - Security & reliability risk
 - Upgradability
 - Portability
 - Linux-specific

vhost: VIRTIO device in the host kernel




```
$ qemu-system-x86_64 -smp 2 -m 2G \
-M q35,accel=kvm \
-drive file=fedora.qcow2,format=qcow2,if=virtio \
-device vhost-vsock-pci,guest-cid=42
```

vhost-user: VIRTIO device in an external process



- Inspired by vhost
 - Control path
 - `AF_UNIX`
 - Data path
 - Shared memory through fd sharing (`memfd`, `/dev/shm`, etc.)
- Why?
 - Security & reliability risks mitigated
 - Upgradability
 - Flexibility
 - Different language from VMM (e.g. Rust)
- Drawbacks
 - More resources used
 - More coordination
 - can be hidden by management layer (e.g. libvirt)
 - Portability
 - see [FOSDEM 2025 - Can QEMU and vhost-user devices be used on macOS and *BSD?](#)

vhost-user: VIRTIO device in an external process




```
$ qemu-storage-daemon \
--blockdev file,filename=fedora.qcow2,node-name=file \
--blockdev qcow2,file=file,node-name=qcow2 \
--export vhost-user-blk,id=vbu,node-name=qcow2, \
num-queues=1,writable=on, \
addr.type=unix,addr.path=/tmp/vhost.socket

$ qemu-system-x86_64 -smp 2 \
-M q35,accel=kvm,memory-backend=mem \
-chardev socket,id=char0,path=/tmp/vhost.socket \
-device vhost-user-blk-pci,num-queues=1,chardev=char0 \
-object memory-backend-memfd,id=mem,size="2G"

# -object memory-backend-shm,id=mem,size="2G"
# can eventually be used on any POSIX host OS (available from QEMU 9.1)
```


vDPA: VIRTIO device (also) in hardware

- virtio Data Path Acceleration
 - VIRTIO compliant data path
 - vendor specific control path
 - small vDPA driver for the control path
- Why?
 - Performance
 - Designed for hardware accelerators
 - software accelerators also possible
 - Support both VMs and containers workloads
 - vhost-vdpa
 - interface for userspace/guest virtio driver
 - virtio-vdpa
 - interface for host virtio driver
 - bare metal or containerized applications



vDPA: virtio Data Path Acceleration

- Unified software stack for vDPA devices
 - Hardware device
 - small parent driver needed
 - Software device
 - in-kernel
 - in-userspace
 - **VDUSE**: vDPA device in Userspace
- Drawbacks
 - Portability
 - Linux-specific
 - Maturity
 - supported by few hardware devices
 - support few virtio types
 - net, block (VDUSE)
 - Cost (HW devices)




vDPA: virtio Data Path Acceleration


```
$ qemu-storage-daemon \
--blockdev file,filename=fedora.qcow2,node-name=file \
--blockdev qcow2,file=file,node-name=qcow2 \
--export vduse-blk,id=vduse1,node-name=qcow2, \
num-queues=1,writable=on,name=vduse1

# instantiate the `vduse1` device (same name used in QSD)
$ vdpa dev add name vduse1 mgmtdev vduse

# the device is identified as `vhost-vdpa-0` in the host
$ ls /sys/bus/vdpa/devices/vduse1/
driver driver_override power subsystem uevent vhost-vdpa-0

$ qemu-system-x86_64 -smp 2 \
-M q35,accel=kvm,memory-backend=mem \
-object memory-backend-memfd,id=mem,size="2G" \
-device vhost-vdpa-device-pci,vhostdev=/dev/vhost-vdpa-0
```





Thank you!

Stefano Garzarella <sqarzare@redhat.com>

<https://stefano-garzarella.github.io/>

 [linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)

 [facebook.com/redhatinc](https://www.facebook.com/redhatinc)

 [youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)

 twitter.com/RedHat

vhost-user protocol

- <https://qemu-project.gitlab.io/qemu/interop/vhost-user.html>
 - *control plane needed to establish virtqueue sharing with an user space process on the same host*
 - **frontend**
 - application that shares its virtqueues (i.e. VMM like QEMU)
 - **backend**
 - consumer of the virtqueues (i.e. virtio device emulation)
- Key components
 - UNIX domain socket (**AF_UNIX**)
 - + ancillary data support to exchange file descriptors
 - shared memory, notifications (irqfd, kickfd), etc.
 - **shared memory** represented by a file descriptor
 - so it can be passed over a UNIX domain socket and then mapped by the other process
 - notifications
 - **eventfd** or **pipe/pipe2**
 - on platforms where eventfd is not available, QEMU will automatically fall back to pipe2 or, as a last resort, pipe.

