



# XDP Virtual Server

An eBPF Load Balancer library

David Coles, February 2026


# Background


- UK's largest commercial broadcast radio company
- Legacy hardware LB appliances handling all traffic
- Replace with one solution for high bandwidth applications
- Separate solution for low bandwidth applications
- 95% long lived TCP connections (**Icecast**)
- IPv4 only, layer 4 LB, Direct Server Return mode at layer 2

## Layer 2 Direct Server Return

## Direct Routing

## nPath (f5)





# Solutions considered

- Google Maglev
- Cloudflare Unimog
- Github GLB
- Facebook Katran
- IPVS

# Try building something with XDP/eBPF?

- Direct Server Return relatively simple to implement
- Need the VIP on loopback interface on backend servers
- Update MAC addresses on incoming packets and TX
- How hard could it be?

```
SEC("xdp") int xdp_forward_func(struct xdp_md *ctx)
{
    const __u32 vip = 192<<24 | 168<<16 | 100<<8 | 1;
    const __u8 h_dest[][][6] = {{ { 0x00, 0x50, 0x56, 0x90, 0xcd, 0x66 },
        { 0x00, 0x50, 0x56, 0x90, 0x37, 0xcd } }};
    void *data_end = (void *)(long)ctx->data_end;
    struct ethhdr *eth = (void *)(long)ctx->data;
    struct iphdr *ip4 = (void *)(eth + 1);

    if ((eth + 1 > data_end) || (eth->h_proto != bpf_htons(ETH_P_IP)) ||
        (ip4 + 1 > data_end) || (ip4->daddr != bpf_ntohl(vip)))
        return XDP_PASS;

    memcpy(eth->h_source, eth->h_dest, 6);
    memcpy(eth->h_dest, h_dest[bpf_ntohl(ip4->saddr) % (sizeof(h_dest)/6)], 6);

    return XDP_TX;
}
```

# eBPF maps can make this dynamic

- Services and destinations
- VLANs and interfaces
- Backend selection and flow tracking
- Statistics
- Queues

# Required library userland functionality

- Control plane functions to manage services
- Backend MAC address discovery via ICMP/ARP
- IP address/VLAN interface discovery
- Other NIC settings, eg. 802.3ad
- Ability to check backend server health by *VIP*

## type Client

```
type Client interface {
    Info() (Info, error)

    Config() (Config, error)
    SetConfig(Config) error

    Services() ([]ServiceExtended, error)
    Service(Service) (ServiceExtended, error)
    CreateService(Service) error
    UpdateService(Service) error
    RemoveService(Service) error

    Destinations(Service) ([]DestinationExtended, error)
    CreateDestination(Service, Destination) error
    UpdateDestination(Service, Destination) error
    RemoveDestination(Service, Destination) error
}
```

```
package main

import (
    "log"
    "net/netip"
    "os"

    "github.com/davidcoles/xvs"
)

// go run balancer.go eth0 10.1.2.0/24 192.168.100.1 10.1.2.3 10.1.2.4
func main() {
    eth := os.Args[1]
    pfx := os.Args[2]
    vip := os.Args[3]
    rip := os.Args[4:]

    options := xvs.Options{IPv4VLANS: map[uint16]netip.Prefix{1: netip.MustParsePrefix(pfx)}}
    client, err := xvs.NewWithOptions(options, eth)
    if err != nil {
        log.Fatal(err)
    }

    service := xvs.Service{Address: netip.MustParseAddr(vip), Port: 80, Protocol: xvs.TCP}
    if err := client.CreateService(service); err != nil {
        log.Fatal(service, err)
    }

    for _, r := range rip {
        destination := xvs.Destination{Address: netip.MustParseAddr(r)}
        if err := client.CreateDestination(service, destination); err != nil {
            log.Fatal(service, destination, err)
        }
    }
}
```

<b>192.168.100.1</b>	<b>10.1.2.3</b>	<b>255.0.0.1</b>
<b>192.168.100.1</b>	<b>10.1.2.4</b>	<b>255.0.0.2</b>
<b>192.168.100.2</b>	<b>10.1.2.4</b>	<b>255.0.0.3</b>
<b>192.168.100.2</b>	<b>10.1.2.5</b>	<b>255.0.0.4</b>
<b>192.168.100.2</b>	<b>10.1.2.6</b>	<b>255.0.0.5</b>

Each unique virtual and real server address pair is mapped to a NAT address

# Health checks via NAT addresses

- xvs interface is created with a route to NAT addresses
- To obtain NAT addr, use `nat := client.NAT(vip, rip)`
- Use standard TCP/IP stack to build checks
  - `http.Get("http://" + nat.String())`
  - Build a custom `http.Client` which always uses NAT IP

# Minimum Viable Product load balancer:

- Health checks
- Advertise VIPs via BGP - “Route Health Injection”
- Metrics
- Logging
- Flow state sharing
- Configuration, eg. YAML/JSON

# Minimum Viable Product load balancer:

- Health checks
- Advertise VIPs via BGP - “Route Health Injection”
- Metrics
- Logging
- Flow state sharing
- Configuration, eg. YAML/JSON

There's an app for that

# vc5

vc5.example.com/

http://vc5.example.com

[config](#) | [status](#) | [metrics](#)

Latency	Traffic	Packets	Connection rate	Active connections	BGP: 127.0.0.1
493ns	36.0 Mbits/s in	60.7 Kpackets/s in	11.3 Kconns/s	223	ESTABLISHED 2m25s

VIP	State	Traffic	Packets	Rate	Active
<a href="#">192.168.100.1</a>	2m15s UP	36.0 Mbits/s in	60.7 Kpackets/s in	11.3 Kconns/s	223

```
31/12/2025, 13:48:33: VIP 192.168.100.1 went up
31/12/2025, 13:48:33: Backend 10.1.2.3:80 for service 192.168.100.1:80:tcp went up
31/12/2025, 13:48:33: Backend 10.1.2.3:443 for service 192.168.100.1:443:tcp went up
31/12/2025, 13:48:33: Backend 10.1.2.4:443 for service 192.168.100.1:443:tcp went up
31/12/2025, 13:48:33: Backend 10.1.2.4:80 for service 192.168.100.1:80:tcp went up
31/12/2025, 13:48:24: Learn timer expired
31/12/2025, 13:48:23: Initialised
31/12/2025, 13:48:23: Backend 10.1.2.5:80 for service 192.168.100.1:80:tcp went down
31/12/2025, 13:48:23: Backend 10.1.2.4:80 for service 192.168.100.1:80:tcp went down
31/12/2025, 13:48:23: Backend 10.1.2.3:80 for service 192.168.100.1:80:tcp went down
31/12/2025, 13:48:23: Backend 10.1.2.6:443 for service 192.168.100.1:443:tcp went down
31/12/2025, 13:48:23: Backend 10.1.2.5:443 for service 192.168.100.1:443:tcp went down
```

<a href="#">192.168.100.1</a>	2m15s UP	36.0 Mbits/s in	60.7 Kpackets/s in	11.3 Kconns/s	223
-------------------------------	----------	-----------------	--------------------	---------------	-----

nginx	[2/4 available - 1 required]				Active
192.168.100.1:80:tcp	2m14s UP	28.9 Mbits/s in	51.7 Kpackets/s in	10.3 Kconns/s	95
10.1.2.3:80	2m14s UP (0ms)	14 400 624	25 790	5 156	49
10.1.2.4:80	2m14s UP (0ms)	14 493 080	25 949	5 193	46
10.1.2.5:80	2m22s DOWN (2000ms)	0	0	0	0
10.1.2.6:80	2m22s DOWN (2000ms)	0	0	0	0

nginx	[2/4 available - 1 required]				Active
192.168.100.1:443:tcp	2m13s UP	7.12 Mbits/s in	8.96 Kpackets/s in	948 conns/s	128
10.1.2.3:443	2m14s UP (0ms)	3 697 888	4 654	493	62
10.1.2.4:443	2m14s UP (0ms)	3 421 664	4 303	455	66
10.1.2.5:443	2m22s DOWN (2000ms)	0	0	0	0
10.1.2.6:443	2m22s DOWN (2000ms)	0	0	0	0

```
---  
bgp:  
  as_number: 65001  
  peers:  
    - 10.1.2.254  
  
services:  
  - name: icecast  
    virtual:  
      - 192.168.100.1  
    servers:  
      - 10.1.2.3  
      - 10.1.2.4  
    host: foo.example.com  
    path: /health  
    policy:  
      http:  
      https:  
  
vlans:  
  1: 10.1.2.0/24
```






config.pl →

```
{  
  "bgp" : {  
    "10.1.2.254" : {  
      "as_number" : 65001  
    }  
  },  
  "services" : {  
    "192.168.100.1:443:tcp" : {  
      "name" : "icecast",  
      "need" : 1,  
      "reals" : {  
        "10.1.2.3:443" : {  
          "checks" : [  
            {  
              "host" : "foo.example.com",  
              "path" : "/health",  
              "type" : "https"  
            }  
          ]  
        },  
        "10.1.2.4:443" : {  
          "checks" : [  
            {  
              "host" : "foo.example.com",  
              "path" : "/health",  
              "type" : "https"  
            }  
          ]  
        }  
      },  
      "sticky" : false  
    },  
    "192.168.100.1:80:tcp" : {  
      ...  
    }  
  }
```

# Deployment

2024 -

- Ran for one year on VMs, small services
- Running our tier 1 services since the start of 2025
- Relatively modest servers - 2.30GHz Xeon, 32core, 32GB
- 4x10G Intel X710 NICs, running VLAN trunks with MLAG
- Migrated services via new VIPs and DNS




MAP OF THE INTERNET  
THE IPv4 SPACE, 2006



THIS CHART SHOWS THE IP ADDRESS SPACE ON A PLANE USING A FRACTAL MAPPING WHICH PRESERVES GROUPING -- ANY CONSECUTIVE STRING OF IPs WILL TRANSLATE TO A SINGLE COMPACT, CONTIGUOUS REGION ON THE MAP. EACH OF THE 256 NUMBERED BLOCKS REPRESENTS ONE /8 SUBNET (CONTAINING ALL IPs THAT START WITH THAT NUMBER). THE UPPER LEFT SECTION SHOWS THE BLOCKS SOLD DIRECTLY TO CORPORATIONS AND GOVERNMENTS IN THE 1990's BEFORE THE RIRs TOOK OVER ALLOCATION.

0	1	14	15	16	19 →
3	2	13	12	17	18
4	7	8	11		
5	6	9	10		



= UNALLOCATED  
BLOCK

<https://xkcd.com/195/>

0 IANA	1 APNI	14 APNI	15 ARIN	16 ARIN	19 FORD	20 ARIN	21 ARMY	234 MCST	235 MCST	236 MCST	239 MCST	240 RSVD	241 RSVD	254 RSVD	255 RSVD
3 ARIN	4 APNI	13 ARIN	12 AT&T	17 APPL	18 ARIN	23 ARIN	22 ARMY	233 MCST	232 MCST	237 MCST	238 MCST	243 RSVD	242 RSVD	253 RSVD	252 RSVD
4 ARIN	7 ARIN	8 ARIN	11 ARMY	30 ARMY	29 ARMY	24 ARIN	25 RIPE	230 MCST	231 MCST	226 MCST	225 MCST	244 RSVD	247 RSVD	248 RSVD	251 RSVD
5 RIPE	6 HRM	9 ARIN	10 IANA	31 RIPE	28 ARMY	27 APNI	26 ARMY	229 MCST	228 MCST	227 MCST	224 MCST	245 RSVD	246 RSVD	249 RSVD	250 RSVD
50 APNI	57 RIPE	54 ARIN	53 DAIM	32 ARIN	35 ARIN	36 APNI	37 RIPE	218 APNI	219 APNI	220 APNI	223 APNI	262 APNI	201 LACN	198 ARIN	197 APNI
59 APNI	56 ARIN	55 ARMY	52 ARIN	33 ARMY	34 ARIN	39 APNI	38 RIPE	216 ARIN	221 APNI	222 APNI	203 APNI	200 LACN	199 ARIN	196 APNI	
60 APNI	61 APNI	50 ARIN	51 RIPE	49 RIPE	45 ARIN	40 ARIN	41 APNI	214 ARMY	215 ARMY	210 APNI	209 ARIN	204 ARIN	205 ARIN	194 RIPE	176 RIPE
63 ARIN	64 APNI	49 APNI	48 PRUD	47 ARIN	44 ARIN	43 APNI	42 APNI	312 RIPE	212 RIPE	211 APNI	208 ARIN	207 ARIN	206 ARIN	192 APNI	193 ARIN
64 ARIN	67 ARIN	68 ARIN	69 ARIN	122 APNI	123 APNI	124 APNI	127 LPBK	128 ARIN	131 ARIN	132 ARIN	133 APNI	186 LACN	187 LACN	198 LACN	191 LACN
65 ARIN	66 ARIN	71 ARIN	70 ARIN	121 APNI	120 APNI	125 APNI	126 APNI	129 ARIN	130 ARIN	135 ARIN	134 ARIN	182 APNI	184 ARIN	189 LACN	190 LACN
78 RIPE	79 RIPE	72 ARIN	73 ARIN	128 APNI	119 APNI	114 APNI	113 APNI	142 ARIN	141 RIPE	136 ARIN	137 ARIN	182 APNI	183 APNI	185 RIPE	177 LACN
79 RIPE	76 ARIN	75 ARIN	74 ARIN	117 APNI	116 APNI	115 APNI	112 APNI	143 ARIN	140 ARIN	139 ARIN	138 ARIN	181 LACN	180 APNI	179 LACN	176 RIPE
80 RIPE	81 RIPE	82 RIPE	83 RIPE	95 RIPE	96 ARIN	97 ARIN	916 APNI	111 APNI	144 ARIN	145 RIPE	158 ARIN	159 ARIN	160 ARIN	161 ARIN	174 ARIN
85 RIPE	86 RIPE	87 RIPE	88 RIPE	90 RIPE	99 ARIN	98 ARIN	100 RIPE	108 ARIN	147 ARIN	146 ARIN	157 ARIN	156 ARIN	163 APNI	162 ARIN	173 ARIN
86 RIPE	87 RIPE	88 RIPE	89 RIPE	100 ARIN	103 APNI	104 ARIN	107 ARIN	148 ARIN	151 RIPE	152 ARIN	155 ARIN	164 ARIN	167 ARIN	168 ARIN	171 APNI
89 RIPE	90 RIPE	91 RIPE	92 RIPE	101 APNI	102 APNI	105 APNI	106 APNI	149 ARIN	150 APNI	153 APNI	154 APNI	165 ARIN	166 ARIN	169 ARIN	170 ARIN

# **Added layer 3 and IPv6 support**

- Tunnel traffic over IP: IPIP, GRE, FOU and GUE supported
- IPv4 or IPv6 may be used for VIP or RIP (any combination)
- Need to tell the library about which gateway to use
- Layer 3-only far simpler than layer 2 with multiple VLANs

# **Issues encountered during rollout**

- Intel X710 media negotiation failures
- VMware vmxnet3 buffer sizing bug in driver mode
- Program size increased with layer 3 support - use tail calls
- FOU/GUE termination on IPv6 unsupported?

# Future development

- Config could be centralised ... or
- Self service - teams could run their own instances ... or
- Service discovery - etcd/Consul/ZooKeeper
- Tunnel termination eBPF utility
- We have the flexibility to proceed at our own pace

X ssh

root@load-balancer:~/demo#

X root@client: ~ (ssh)

root@client:~#

X Default (ssh)

root@backend2:~#

**This slide intentionally left blank in case of video playback screwups**

# Enjoy FOSDEM!

**github.com/davidcoles/xvs**

**github.com/davidcoles/vc5**

**github.com/davidcoles/fosdem2026**

