

# Zero-Downtime Upgrades: PostgreSQL and OS/glibc at Global Scale

# Alexander Sosna

Senior Database Reliability Engineer



# Agenda

This talk will showcase:

- How we execute PostgreSQL and OS upgrades at GitLab, with **zero downtime**.

By answering these questions:

- PostgreSQL Upgrades - How do they work, and why are they hard?
- OS Upgrades - How do they work, and why are they hard?
- What did we do to minimize impact to our users?

*To fit the time slot, some aspects are simplified, details and code in the linked resources!*



# README

- Slides with the white triangle in the corner are not included in the presentation
- They are added to provide more context when reading the slides



# Why are PostgreSQL Major Upgrades hard?

- Major releases (can) change the layout of system tables
- Data files can not be used by newer versions
- Rewriting of system tables and metadata is necessary
- Helping structures like indexes might require a rebuild
- Depending on data size and complexity this can take significant time



# Upgrade Method - pg\_dumpall (default)



1.

maintenance mode  
(DB becomes RO)

**User impact starts**

2.

physical to logical  
(binary to SQL\*)

3.

logical to physical  
(SQL\* to binary)

4.

create indexes  
collect statistic

# Upgrade Method - pg\_dumpall



- Data is extracted and brought to a logical representation
  - SQL, or optimized internal format
- Logical data is then imported in the new cluster
- Both operations are resource and time consuming
  - Can be performed in parallel to disk  
OR
  - Piped from old to new cluster
- All data gets validated
- All indexes are freshly created
- No bloat in the new cluster


# Upgrade Method - pg\_dumpall



- **Safest method available**
- Also able to upgrade
  - OS/glibc
  - Hardware architecture, e.g. x86 ⇒ RISC-V
- Some data types like *jsonb* get validated
- Requires **downtime based on data and indexes**
  - Hard to provide simple estimate: our ~40 TiB DB will take > 24h
  - You can easily try it out and measure to get exact timing
- No quick rollback after upgrade!

If this fulfills your needs, it's the safest option! Don't look any further!

# Upgrade Method - pg\_upgrade



1.


Maintenance mode  
(offline / RO with standby)

**User impact starts**

2.

In-place upgrading  
binary data

# Upgrade Method - pg\_upgrade



1.

Maintenance mode  
(offline / RO with standby)

**User impact starts**

2.

In-place upgrading  
binary data



# Upgrade Method - pg\_upgrade

- Quite simple
- Reasonable fast
  - Additional operations like a reindex or tests can take longer!
- Reasonable safe
- No quick rollback after upgrade!
- When I joined GitLab, we used it as well
  - Due to mandatory QA tests total downtime was >4h per upgrade
  - Upgrades were avoided due to downtime

If this fulfills your needs, it's a safe and simple option! Don't look any further!

# How did we perform Upgrades in the Past?

*pg\_upgrade*, with significant downtime

1. Create second cluster from backup, called *Target*
2. Sync *Target* with *Source* cluster via streaming replication
- 3. Put **GitLab.com** into maintenance**
4. Used *pg\_upgrade* to upgrade *Target* cluster primary
5. Re-create all standbys in *Target* cluster
6. Run full QA tests and benchmark on *Target* cluster (multiple hours)
7. Switch application to use new cluster
8. Bring **GitLab.com** back online



# Why can't we use a boring solution for GitLab.com?



# Why can't we use a boring solution for GitLab.com?

- GitLab.com is a globally used SaaS offering
  - > 50 million users around the world
  - > 2,500 team members, all-remote and globally distributed (>65 countries)
  - > 1 Million SQL queries per second on PostgreSQL (US working hours)
  - There is not a single minute, at which a downtime would not impact users and team members!
  - Data Sources [ir.gitlab.com](https://ir.gitlab.com), [about.gitlab.com/company/team](https://about.gitlab.com/company/team)
- No budget for downtime
- We need to be able to roll back if the new DBMS does not perform



# How do you define “Zero Downtime” in SaaS?

- User requests are not handled instantaneously
- When a user presses a button it takes time before the result is shown
- We can't go for “0 ms” downtime :)



# How do you define “Zero Downtime” in SaaS?

- User requests are not handled instantaneously
- When a user presses a button it takes time before the result is shown
- We can't go for “0 ms” downtime :)

“Zero Downtime” ⇒ no user impact!



# How is GitLab measuring User Impact?

- Apdex (Application Performance Index)
  - Open standard for measuring application performance
  - Based on classifying user interactions into
    - “satisfied”
    - “tolerating”
    - “frustrated”
  - Requires tuned thresholds to classify samples
  - Details: [wikipedia.org/wiki/Apdex](https://en.wikipedia.org/wiki/Apdex)

$$\text{Apdex}_t = \frac{\text{SatisfiedCount} + (0.5 \cdot \text{ToleratingCount}) + (0 \cdot \text{FrustratedCount})}{\text{TotalSamples}}$$



# How do we achieve Zero Downtime?



# How do we achieve Zero Downtime?

Logical Replication



# How are we achieving Zero Downtime?

Logical Replication  
(and a lot of automation)



# Logical Replication

- Unlike Streaming Replication, LR can replicate across different PG versions
- We can upgrade a clone of our production database and bring it in sync
- Does it come with restrictions?
  - Yes!
  - Watch my previous talk or read the extended slide deck
  - [How we execute PG major upgrades at GitLab, with zero downtime.](#)  
(PGConf.EU 2023) youtube.com/watch?v=o08kJgkova
  - Important for this talk: Schema changes would break LR!
    - No DDL allowed: CREATE, ALTER, DROP, ...



# Logical Replication - What is the catch?

1. Database schema and DDL commands are not replicated!
2. Sequences are not replicated, but are needed for auto increment values
3. Each table needs a *REPLICA IDENTITY*, to distribute changes
  - o Primary key
  - o Other unique key
  - o *FULL*, last resort, all changes need to be recorded
4. More complex
  - o Prone to human errors
  - o Automation and testing is highly advised



# Logical Replication - DDL is not replicated

- Schema changes would break logical replication!
  - No DDL allowed: CREATE, ALTER, DROP



# Logical Replication - DDL is not replicated

- Schema changes would break logical replication!
  - No DDL allowed: CREATE, ALTER, DROP

Our solution

- Disable all deployments, migration, and maintenance jobs creating DDL
  - GitLab features
    - Database upgrade DDL lock
    - *disallow\_database\_ddl\_feature\_flags*, MR130554
  - You need to check **YOUR** applications DDL usage!
    - Most common software will not erratically execute DDL



# Logical Replication - Sequences are not replicated

- Sequences are vital to PostgreSQL
  - Generates unique sequential numbers wherever they are needed
  - Used for SERIAL (AUTO INCREMENT)



# Logical Replication - Sequences are not replicated

- Sequences are vital to PostgreSQL
  - Generates unique sequential numbers wherever they are needed
  - Used for SERIAL (AUTO INCREMENT)

## Our solution

- Measure the daily growth of all sequences
- Defined a large “sequences buffer value”, e.g. *1 million*
- Increase the sequences on the Target cluster by this value
- Before switchover check that the sequences on OLD, have not grown more than expected (optional)
- Simple solution, only uses up a fraction of the key space of 64-bit integer



# Logical Replication - REPLICA IDENTITY

- Each table needs a *REPLICA IDENTITY* to clearly identify rows, like:
  - Primary key
  - Other unique key
  - *FULL Record*, last resort, all changes need to be recorded



# Logical Replication - REPLICA IDENTITY

- Each table needs a *REPLICA IDENTITY* to clearly identify rows, like:
  - Primary key
  - Other unique key
  - *FULL Record*, last resort, all changes need to be recorded

## Our solution

- Nothing to do, GitLab already used primary keys
- You need to check **YOUR** application's schema!



# Logical Replication - Complexity

- More complex
  - Prone to human errors



# Logical Replication - Complexity


- More complex
  - Prone to human errors

Our solution


- Automation
  - Orchestration via Ansible
  - Process as CR issue which could be executed repetitively
- Excessive testing - "*When it hurts, do it more often*"
  - Intense QA tests before switchover, rollback if not perfect
  - Dry runs in production



# Logical Replication + pg\_upgrade



# Logical Replication + pg\_upgrade



Create and sync Target




# Logical Replication + pg\_upgrade



Upgrade Target  
(no sync during upgrade)




# Logical Replication + pg\_upgrade




Resync via LR

# Logical Replication + pg\_upgrade




Application Switchover

# Logical Replication + pg\_upgrade




# PostgreSQL Upgrade - State 2023



# **What is actual the User Impact?**




# How well did we do? - Web Apdex



- Web Service Apdex - top 1% (0.99 - 1.00 nit-picking view)
- Degradation SLO: 98.8%, red line would be below this graph :D




# How well did we do? - Web Apdex



- Web Service Apdex - top 1% (0.99 - 1.00 nit-picking view)
- Degradation SLO: 98.8%, red line would be below this graph :D

# How well did we do? - API Apdex



Apdex top 4% (0.96 - 1.00)



# Can we improve further?

1. Switchover is a Point of no Return
  - If performance degrades or any problem occurs, we can't go back!
  - Significant business risk!
2. This approach only upgrades PostgreSQL
  - OS or library upgrades are not handled




# (Re)move Point of no Return - Reverse Replication

- After the Switchover we reverse the replication
- Enables swift rollback without data loss




# (Re)move Point of no Return - Reverse Replication



Reverse Replication




# (Re)move Point of no Return - Reverse Replication



Operation and Monitoring


# (Re)move Point of no Return - Reverse Replication



Late Rollback  
(optional)



# PostgreSQL Upgrade - With Reverse Replication



# Why are OS Upgrades hard?

- When upgrading the OS, you will get a new version of glibc (GNU C Library)
  - This library defines the system-wide collation
- Collation: Set of rules that describe how strings are compared and ordered
  - "A" < "B" < "C"
  - "1" < "2" < "3"
  - "10" < "2" OR "10" > "2"
  - "\\" < "/" OR "/" > "\\"
- Indexes
  - Need to be used with the collation they were created with!
  - If not, data corruption can occur!



# OS Upgrade - Simple Solution

- Some data types don't use collations and are unproblematic, e.g. INTEGER
- Rebuild all indexes (on strings) with the current collation
  - If this works for your use-case, great!
  - If you use the pg\_dumpall upgrade method you get it automatically
  - For [GitLab.com](#) this would take multiple days, longer than our upgrade window




# OS Upgrade - Optimized Approach

- Before we start the upgrade we automatically create a list of all indexes, where the new collation can lead to corruption. (Script based on [amcheck](#))
  - No need to recreate non-problematic types like INTEGER
  - No need to recreate indexes only containing non-problematic data
    - Example: md5 hashes (strings)
- We recreate all listed indexes on a test system, to measure the execution time
  - If it takes longer than acceptable, we can optimize beforehand
    - Replace indexes
      - Different type
      - Multiple partial indexes
    - If non-disruptive: lazily recreate after upgrade



# OS Upgrade - Optimized Approach

- Saturday: Upgrade
- Sunday: Switchover
- After the upgrade step we have between 12h to 24h before the Switchover to:
  - recreate all problematic indexes
  - run amcheck to verify no data corruption
  - run additional tests if necessary




Recreate Indexes  
Run amcheck

Switchover




# PostgreSQL and OS Upgrade

- 2025 we upgraded most of our database systems
  - PG16 ⇒ PG17
  - Ubuntu 20.04 ⇒ 22.04
- Let's walk through one of the last upgrades




# PostgreSQL and OS Upgrade


DDL



# PostgreSQL and OS Upgrade




# PostgreSQL and OS Upgrade - Preparation



Create Test clone

# PostgreSQL and OS Upgrade - Preparation




Test upgrade

Get execution times


Get list of corrupted indexes

# PostgreSQL and OS Upgrade - Preparation



Remove Test Cluster  
Create Target Cluster


# PostgreSQL and OS Upgrade - Saturday



Switch to logical replication  
(DDL would break it)

# PostgreSQL and OS Upgrade - Saturday


No  
DDL



Upgrade Target  
(no sync during upgrade)




# PostgreSQL and OS Upgrade - Saturday



Resync



# PostgreSQL and OS Upgrade - Saturday




Reindex  
Analyze (collect statistics)  
Corruption Check



# PostgreSQL and OS Upgrade - Sunday

No  
DDL




Switchover read-only queries partially  
Monitor performance



# PostgreSQL and OS Upgrade - Sunday

No  
DDL




Switchover all read-only queries  
Monitor performance



# PostgreSQL and OS Upgrade - Sunday

No  
DDL



Run full QA test suite


QA + live traffic

Monitor performance



No  
DDL

# PostgreSQL and OS Upgrade - Sunday




Switchover all load



No  
DDL


# PostgreSQL and OS Upgrade - Sunday



Reverse Replication



# PostgreSQL and OS Upgrade - Monday




Monitoring during peak hours  
(Fast Rollback possible)



No  
DDL

# PostgreSQL and OS Upgrade - Tuesday




Point of no return  
Remove PG16 cluster



# PostgreSQL and OS Upgrade - Tuesday

DDL



Normal operation  
Start planning next upgrade ;)



# Resources

- GitLab: [about.gitlab.com](https://about.gitlab.com)
- Our RDBMS: [about.gitlab.com/handbook/engineering/infrastructure/database](https://about.gitlab.com/handbook/engineering/infrastructure/database)
- Ansible Playbooks: [gitlab.com/gitlab-com/gl-infra/db-migration](https://gitlab.com/gitlab-com/gl-infra/db-migration)
- CR Template: [..../db-migration/.gitlab/issue\\_templates/pg\\_upgrade.md](https://..../db-migration/.gitlab/issue_templates/pg_upgrade.md)
- Extended Slide Deck with addition annotations:
  - [FOSDEM26 - fosdem.org/2026](#)
  - [FOSDEM PGDay 2026 - 2026.fosdempgday.org](#)
- Previous Talk
  - [How we execute PG major upgrades at GitLab, with zero downtime. \(PGConf.EU 2023\) youtube.com/watch?v=o08kJggkovq](#)
- Alexander Sosna
  - [sosna.de](http://sosna.de)



# Questions?

- During the event
- GitLab Stand at FOSDEM
- Later
- Now!



[sosna.de](http://sosna.de)

